



JÖNKÖPING UNIVERSITY

Exploring AI-Based Emotion Recognition in Swedish: Speech, Text, and Vocal Markers

Author(s): Sara LJUNG & Janna HAKKARAINEN
Main subject area: Computer Engineering
School: School of Engineering
Date: May 2025

This final thesis has been carried out at the School of Engineering at Jönköping University within Computer Engineering. The authors are responsible for the presented opinions, conclusions and results.

Examiner: Neziha AKALIN

Supervisor: Garrit SCHAAP

Scope: 15 Credits

Date: 2025-03-09

Contents

Abstract	4
1 Introduction	5
1.1 Background	5
1.2 Problem Description	9
1.3 Purpose and Research Questions	11
1.4 Scope and Limitations	12
1.4.1 Scope	12
1.4.2 Limitations	12
1.5 Disposition	13
2 Method and Implementation	14
2.1 Approach and design	14
2.2 Data Collection	14
2.2.1 Research Question 1	16
2.2.2 Research Question 2	16
2.2.3 Research Question 3	17
2.3 Data Analysis	17
2.3.1 Statistical Analysis Approach	17
2.3.2 RQ2 and RQ3: Speech-Based AI vs. Text-Based AI, AI-labels vs. Self-Assessed Emotions	18
2.3.3 Data Normalization	18
2.3.4 Visual Analysis	19
2.4 Model Configuration	20
2.4.1 NLP Cloud	20
2.4.2 Hume AI	20
2.5 Validity and Reliability	20
2.5.1 Validity	20
2.5.2 Reliability	21
2.6 Considerations	21
3 Theoretical Framework	22
3.1 Affective Computing	22
3.2 Natural Language Processing and Emotion Recognition	22
3.3 Speech-Based Emotion Recognition	23
3.3.1 Hume AI	23
3.3.2 Praat Parselmouth	25
3.4 Text-Based Emotion Recognition	26
3.4.1 NLP Cloud	26
3.5 Vocal Markers	28
3.6 The Experiment	29
3.6.1 Python Application	30
3.6.2 Interviews and Surveys	30
3.7 Statistical Analysis	31

4 Results	32
4.1 Presentation of Collected Data	32
4.1.1 Overview of Interviews	32
4.1.2 Data Collection for RQ1: Vocal Features & Speech	32
4.1.3 Data Collection for RQ2 and RQ3: Text, Speech and Self-Assessment	34
4.2 Data Analysis for RQ1: Vocal Features & Speech Emotion Recognition	36
4.2.1 Correlation Between Vocal Features and AI Emotion Scores (Hume AI)	36
4.2.2 Correlation with Praat-Based Emotion Scores	37
4.2.3 Correlation Custom Categorization and Hume AI Labels	38
4.2.4 Limitations of the Custom Vocal Emotion Categorization Method	40
4.2.5 ANOVA Tables of Vocal Features Across Emotions	41
4.2.6 Time-to-Time Analysis	42
4.2.7 Conclusion RQ1 Data Analysis	47
4.3 Data Analysis for RQ2: Text and Speech Based Emotion Recognition	48
4.3.1 Comparative Overview of Model Outputs	48
4.3.2 Statistical Analysis	49
4.3.3 Conclusion of RQ2 Data Analysis	53
4.4 Data Analysis for RQ3: AI and self-assessed emotion labels	53
4.4.1 Model Emotion Score and Self-Reports Comparison	53
4.4.2 Correlation and Visual Analysis	54
4.4.3 Statistical Analysis and Effect Sizes	58
4.4.4 Conclusion of RQ3 Data Analysis	60
5 Discussion	62
5.1 Result Discussion RQ1	62
5.1.1 Interpretation of Results	62
5.2 Result Discussion RQ2	65
5.2.1 Interpretation of Results	65
5.2.2 RQ2: Speech-based AI vs Text-based AI	65
5.3 Result Discussion RQ3	66
5.3.1 RQ3: AI vs Self-Assessed Emotions	66
5.4 Conclusions	68
5.4.1 Conclusion for RQ1	68
5.4.2 Conclusion for RQ2	69
5.4.3 Conclusion for RQ3	69
5.5 Method Discussion	69
5.5.1 RQ1 Methodological Considerations	69
5.5.2 RQ2 and RQ3 Methodological Considerations	70
5.5.3 Summary of Methodological Considerations	71
5.6 Limitations	71
5.6.1 Limitations RQ1	71
5.6.2 Limitations RQ2 and RQ3	72
6 Conclusion	73
6.1 Summary of Key Findings and Answering Research Questions	73
6.2 Contribution to the Field	73
6.3 Limitations of the Study	74
6.4 Future Research	74
6.5 Final Conclusion	74

Abstract

Keywords:

Introduction

This thesis aims to explore emotion recognition and its effectiveness in the Swedish language. With the rapid advancement of the technology industry and artificial intelligence, emotion recognition has started to play an increasingly important role in the enhancement of human-computer interactions. These areas hold potential to transform and develop several important fields, but there are still challenges in the field. Much of the research has been focused on specific languages, notably English. This research focuses on emotion recognition across two distinct modalities in Swedish, speech-based emotion recognition and text-based emotion recognition and aim to contribute to broadening the field of emotion recognition in a non-English language.

1.1 Background

According to Oatley et al. (2019), emotion recognition has attracted increasing attention with the rapid advancement of technology and artificial intelligence. Emotions are experienced by all humans but are difficult to define precisely. They are an internal experience that are foundational to our sense of identity, our relationships, and moral judgement. Scientists have faced challenges in the effort to characterize how emotions are communicated. Emotions are internal but also expressed externally through voice and movements of the body. They are not only communicated through the words we say, but also how we say them. Tones of the voice is a source of varied emotional expressions where its states may alter patterns in vocalizations. It is considered that various emotion-related physiological changes influence acoustic features such as pitch, tempo, pitch variability, and loudness in the speech autociteOatley2019. Beyond spoken signals, researchers have also developed a set of Natural Language Processing (NLP) techniques to infer emotional states and opinions directly from text, based on methods at the intersection of artificial intelligence, computer science, and linguistics (Kansara et al., 2020). With the development of Artificial Intelligence several techniques have accelerated in the recent years, including for NLP, even if its origins back to the 1950s when questions about whether a machine could learn and think to interact with humans raised. (Núñez et al., 2024). NLP has remained as a significant contributor of AI. Some of the active research areas in the NLP domain is Machine Translation, Chatbots, recognizing speech, text summarization, and sentiment analysis (S. Kusal et al., 2023). Figure 1.1 demonstrates the different subdomains of NLP.

Sentiment analysis is a computational branch in NLP that utilizes the detection and evaluation of people's emotions, opinions, and moods based on text, speech, facial expression, etc., without analyzing these feelings (Ermakova et al., 2023). The rise of sentiment analysis is associated with the growth of social media, which has generated vast amounts of digital option data recorded in digital forms. Since the early 2000's, the field has become one of the most researched parts in NLP (L. Zhang et al., 2018), expanding beyond computer science to fields like finance, marketing, political- and health science. Accordingly, sentiment analysis is valuable across different areas of society. Sentiment analysis is utilized in the popular index called the happy planet index (HappyPlanetIndex, n.d.), measuring sustainable well-being of different countries, even if it only can observe three feelings, positive, negative, or neutral. The happy

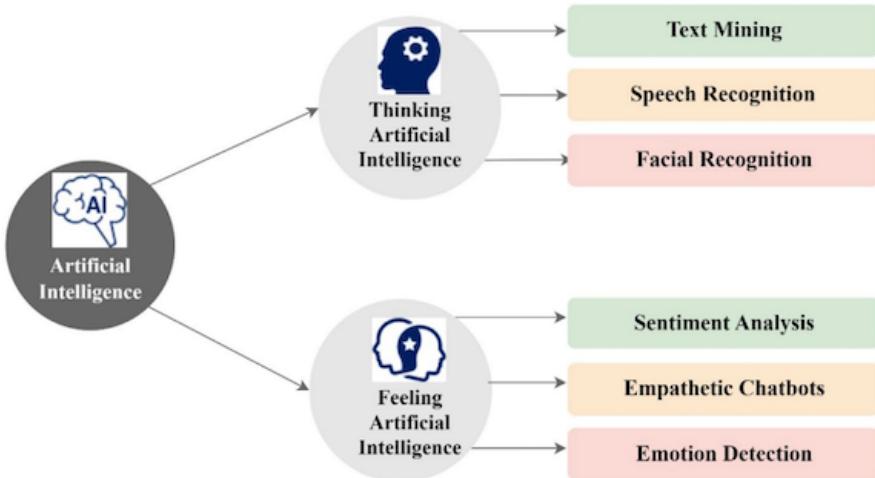


Figure 1.1: *Subdomains of NLP (S. Kusal et al., 2023)*.

planet index checks the happiness level calculated from a particular country, where emotion detection is used with sentiment analysis (Madhuri & Lakshmi, 2021). With the evolution of deep learning networks, emotion detection has advanced (Safari & Chalechale, 2023). Sentiment analysis identifying positive, negative, or neutral states have progressed into recognizing the six basic emotions; joy, sadness, anger, disgust, fear, and surprise in text. The emotions categorization fluctuate depending on the research. These basic six were determined by Paul Ekman (S. Kusal et al., 2023; Oatley et al., 2019) who determined that these six fundamental emotions is shared in people of different cultures, characterized by facial features. However, Ekman's classification was made over 20 years ago when no agreement about what emotions should be considered as existed. Today, the agreement about evidence for universal emotional signals and evidence for five emotions is robust: anger, disgust, sadness, happiness, and fear (Ekman, 2016).

Emotion recognition from textual data is important in various domains such as customer reviews, social media analysis, public monitoring, and conversational agents. A systematic review (S. Kusal et al., 2023) shows that Deep Learning models outperform traditional Machine Learning models due to their ability to capture contextual dependencies. The review further demonstrates the highest accuracy (76%) is shown by transformer-based models such as bidirectional encoder representations from transformers (BERT), highlights challenges such as small or imbalanced datasets that can affect the model reliability, and notes that multimodal approaches with text, speech, and images improve emotion recognition (Madhuri & Lakshmi, 2021). However, text-based emotion detection (TBED) has challenges with identifying hidden emotions, and adapting to diverse languages. Datasets based on different languages than English, as Arabic and Hindi, are tested in a study (Maruf et al., 2024) that identifies challenges as limited resources for non-English languages. The authors underscore the potential of TBED but notes limitations as it is no universal solution for challenges like sarcasm, dynamic emotions, and cultural variances.

Emotion detection research progressed with Speech Emotion Recognition (SER) (S. Kusal et al., 2023). It has shown that hearers can evaluate five emotions in speech-prosody, anger, happiness, sadness, fear, and tenderness, with 70 percent accuracy (Oatley et al., 2019). Speech emotion recognition focuses on how something is said rather than the words themselves. Acoustic features like amplitude, formants, and pitch help classify emotions. Those features offer invaluable insights into the subtle emotional expressions conveyed through speech, assisting the complicated process of emotion recognition (Lian et al., 2023).

Several studies distinguish different emotions through vocal features. Already in 2005, automatic recognition of positive and negative emotions in spoken dialogs was investigated (C. M. Lee & Narayanan, 2005). In that study, acoustic, lexical, and discourse information were combined to enhance emotion detection and move beyond traditional acoustic-only ways. The authors analyzed acoustic features, lexical features, and discourse features. Linear Discriminant Classifiers were used and resulted in good performance for acoustic and lexical information. A study by Bänziger et al. (2014) demonstrated that human listeners could reliably rate emotional expressions in acted voices. These human judgements had higher accuracy for detection of certain emotions, such as happiness, compared to technical analyses of acoustic measurements.

According to Khalil et al. (2019), acoustic features enable emotion recognition through speech using deep learning, which offers many advantages over traditional sentiment-analysis methods. Deep learning models have the capability to automatically detect complex patterns and varying features without requiring manual feature extraction. The goal of speech emotion recognition (SER) is identification of emotions in speech, unrelated to the semantic content (S. D. Kusal et al., 2024). Figure 1.2 represents a SER system.

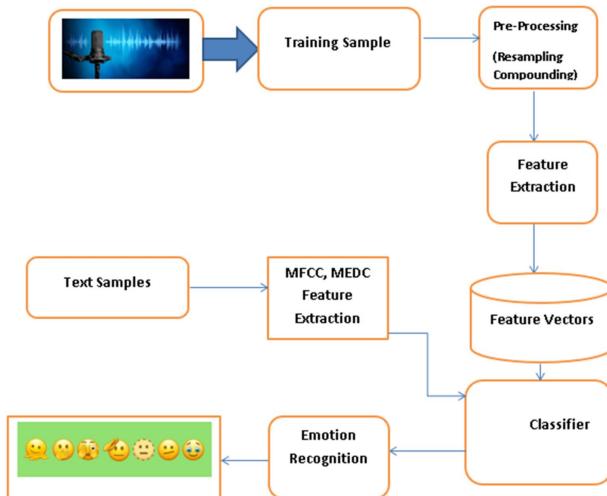


Figure 1.2: Block diagram of SER (Tyagi & Szénási, 2024).

In recent years, speech emotion recognition has emerged as a research area driven by its applications in human-computer interaction (S. Zhang et al., 2021). The advancements have led to the development of intelligent affective services in fields such as call centers, healthcare, surveillance, and affective computing. The accuracy of models tested in recent years have improved significantly (Adebiyi et al., 2024; Praseetha & Joby, 2022; Rahman et al., 2024). Several studies conducted in the last year's show emotion detection accuracy results over 90%.

Juslin et al. (2018) concluded in 2018 analyzed 1,877 voice clips from 23 datasets to compare spontaneous and posed emotions. Their findings highlighted key differences:

- Spontaneous expressions were rated as more genuine than posed ones, even when intensity was controlled.
- Posed expressions were more intense, but intensity alone did not fully explain perceived authenticity.
- Acoustic differences were small but present, mainly in pitch range, speech rate, and voice intensity.
- Highly intense spontaneous emotions conveyed emotions as clearly as posed ones, suggesting that emotion intensity plays a role in perception.

These findings underscore that posed and spontaneous emotions are not interchangeable and that SER datasets must distinguish between these sample types to build models that can generalize to real-world emotional speech accurately.

One recent review of SER corpora and features (Rathi & Tripathy, 2024) shows that most studies still target only six emotions—happiness, anger, sadness, surprise, fear, and neutrality—even though narrower sets (e.g. anger, fear, happiness, sadness) dominate earlier work (Scherer et al., 2018). In contrast, GoEmotions is a large, detailed text database for 27 distinct emotions, a study by Demszky et al. (2020) obtained an average F of 0.46 (0.86 for gratitude, 0.00 for grief) and 0.64 when reduced to six labels. While this GoEmotions-study is included in the research behind the commercial system Hume.ai, and highlights the value of fine-grained emotion categories (Hume, n.d.-a), which is important to acknowledge since because of potential biases.

Datasets drive both speech- and text-based emotion models. Speech emotion recognition datasets are gathered in three ways, acted by performers, induced in controlled settings, or from natural conversations, affecting how expressive and realistic the recordings are. Rathi and Tripathy (2024) analysed 93 research papers where IEMOCAP and RAVDESS are among the most widely used datasets, chosen by 35.83% and 21.50% of researchers, respectively. They further state that dataset choice, recording conditions, and selected features (e.g. MFCCs, pitch, intensity, prosody) impact SER accuracy significantly, and that natural speech is more difficult to classify due to its high variability and background noise.

The number of natural datasets is relatively limited (Cai et al., 2023), and many research papers test on acted datasets. For example, the empirical analysis Ahammed et al. (2024) demonstrates a high-accuracy SER system (100% accuracy, precision, and F1-score) on a combined RAVDESS, TESS, and SAVEE dataset. Each dataset includes posed or elicited emotions in English speech. Similarly, different models for SER achieved over 94% accuracy for these same acted datasets (Alroobaea, 2024). However, spontaneous speech is not validated in these studies and depends on acted data.

In contrast, Text-Based Emotion Detection (TBED) are driven by diverse text datasets, from six emotions to GoEmotions set with 27 emotions (S. Kusal et al., 2023). Researchers in TBED can use publicly available datasets with reliable annotation, for instance derived from stories, publications, news, social media texts, or reviews on movies. According to S. Kusal et al. (2023), many datasets are based on social media, including casual writing style which is a big challenge. The use of short messages and informal language has limited research. Human emotion expressions and the texts conveying them are ambiguous and subjective, additionally, emotions are multifaceted with varying expressions. Therefore, the authors claim that human mapping is important. Over 3.5 million self-labeled posts on Twitter was used to train a model in S. J. Lee et al. (2023), achieving up to 0.87 F1 on human-annotated sets and 0.79 F1 on self-reported hashtags. However, like SER, TBED is dominated by English and lacks large, naturalistic datasets in other languages.

The promising development of emotion recognition has been adapted in research for other areas than computer and machine learning science. SER is beneficial in translating languages, interactive courses and tutorials held online where the student's emotional state can be understood to help the machine make decisions on how to present the course (Abbaschian et al., 2021). It can be implemented in vehicles' safety structures to recognize the driver's emotional state and therefore prevent accidents.

Several studies (DeSouza et al., 2021; Drougkas et al., 2024; Simcock et al., 2020; Singh, 2023) demonstrate the potential benefit of AI-based emotion recognition in mental health, investigating it can assist psychiatrists diagnosing and identifying potential mental illnesses. DeSouza et al. (2021) showed how leveraging speech and text analysis with NLP can help detect late-life depression and predict its severity with 86-92% accuracy. Drougkas et al. (2024)

compared unimodal approaches, either text- or audio-based and combined audio-text models, resulting in text unimodal accuracy between 78% and 87% with F1 scores from 0.60 to 0.79, audio unimodal accuracy of 64%-72% with F1 values as low as 0.0 up to 0.46. Multimodal approaches, combining text and audio, showed similar accuracy (80% - 87%) and F1 scores (0.60-0.80) as text-unimodal approaches. The authors conclude that text model outperform the acoustic model in recognising mental health indicators, but that multimodal models can outperform unimodal techniques since positive F1 scores increase combining the models.

In summary, speech emotion recognition is proficient in capturing vocal cues, especially on acted datasets, while text-based emotion detection relies on transformer models trained on large text collections. However, most research is based on English and uses acted or social-media data, with few studies exploring natural, spontaneous speech or on other languages.

1.2 Problem Description

Despite significant progress in speech emotion recognition, there are limitations in current research. For instance, emotional voice samples are usually obtained from actors portraying emotions using scripted speech. These acted expressions tend to be more intense and exaggerated than naturally occurring emotions. However, this method risks overemphasizing obvious emotional cues while missing subtle ones. It is argued that such portrayals reflect social norms more than genuine physiological responses, although all public expressions may involve some degree of performance (Scherer et al., 2018).

The way emotional speech data is collected depends on the design and purpose of the SER system. As datasets shift from acted emotions to more spontaneous or real-life emotions, emotion recognition becomes more challenging. Many researchers prefer acted emotion datasets because they offer a wide range of emotions and large amounts of data (Rathi & Tripathy, 2024). Induced datasets are collected by constructing an artificial emotional situation, without the knowledge of the performer or speaker. This results in a more naturalistic database, but issues regarding ethics may apply, since the speaker should know they have been recorded for research (Khalil et al., 2019).

Estimation of emotions from spontaneous speech is a challenging task. Most studies test models on acted datasets (Ahammed et al., 2024; Alroobaea, 2024; Khalil et al., 2019; Praseetha & Joby, 2022). The primary reason for the concentration on acted SER tasks is that acted emotions can be easily performed in a controlled laboratory setting, often resulting in high SER accuracy. However, these emotions tend to be exaggerated and may not accurately reflect how emotions are expressed in real-world situations. Consequently, detecting spontaneous emotions in natural environments is significantly more complex and challenging compared to recognizing acted emotions (S. Zhang et al., 2021). Figure 1.3 demonstrates the difficulty level for varying settings.

Text-Based Emotion Detection (TBED) shows similar gaps regarding the data the majority of the researched models are trained and evaluated on. Although transformer models reach up to 76 % accuracy on English datasets (S. Kusal et al., 2023), they are heavily dependent on informal social-media or review texts. Moreover, TBED resources for other languages is limited, and challenges like sarcasm and cultural nuance affects the reliability of the models (S. J. Lee et al., 2023; Maruf et al., 2024).

The Swedish language is not widely spoken and therefore very limited research has been concluded on the Swedish speech. One study (Ekberg et al., 2023) investigated Swedish emotion recognition through feature extraction and concluded that emotions in Swedish speech have unique sound patterns. The emotion results showed that surprise is a very distinct emotion, but happiness and anger sound alike, which could confuse listeners.

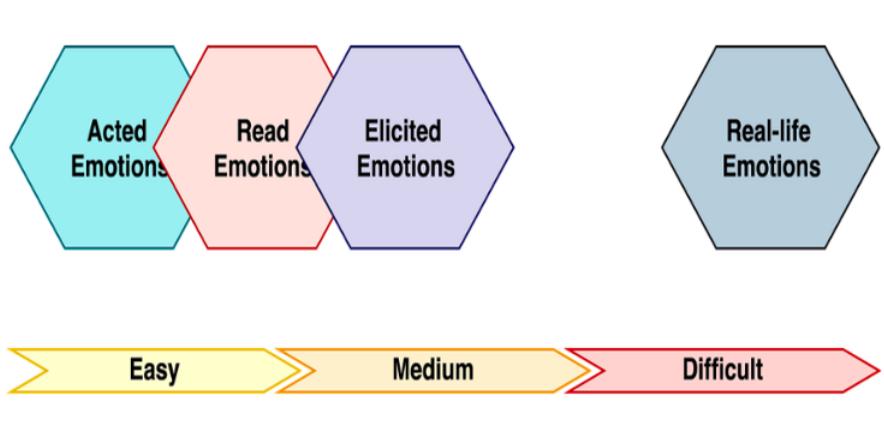


Figure 1.3: *Emotion recognition databases and their difficulty level (Khalil et al., 2019)*.

Limitations as overreliance on acted English data, lack of natural and non-English datasets, and modality-specific biases, are motivations for this thesis. We will compare SER and TBED in Swedish, using both speech and text from the same speakers and explore the alignment against their self-reported emotions to test real-world performance.

1.3 Purpose and Research Questions

The advancement of artificial intelligence (AI) has significantly improved the ability to recognize human emotions, both through speech and text. This offers transformative potential across domains such as mental health, education, and human-computer interaction. Speech Emotion Recognition (SER) and Text-based emotion detection (TBED) have become key areas within the field of Natural Language processing (NLP), leveraging deep learning to interpret different emotional cues with increasing accuracy. However, despite these advancements, significant challenges remain in ensuring that emotion recognition systems are robust, culturally inclusive, and reflective of real-world emotional expressions. Much of the existing research relies on acted datasets, which may underperform when it comes to subtle, spontaneous emotions in everyday contexts, and there is a notable gap in understanding how these models perform across diverse linguistic and cultural situations, such as the Swedish language. The number of studied languages is not that broad, and the studies on accuracy for a new language implies that more research on the generalizability to other languages is essential. Furthermore, while speech and text offer complementary perspectives on emotions, their alignment with individuals own perceptions of their emotions remains unexplored.

This study aims to address the dataset gap by investigating the performance of AI-driven emotion recognition systems in a specific context: Swedish speech and its transcribed textual content. By focusing on Swedish – a language with limited prior research in SER – this thesis seeks to contribute to a broader understanding of how linguistic and cultural factors can influence emotion recognition, which is applicable to multilingual understanding for emotion recognition for different languages. Additionally, the integration of speech and text analysis gives an opportunity to explore multimodal approaches. The alignment between AI-generated emotion labels and self-reported emotions is an overlooked area. Although emotions are inherently difficult to define and can be challenging for individuals to self-assess, it is valuable to examine the alignment between model outputs and people's own perceived emotions. Publicly available AI models and APIs, despite their use in real-world applications, are rarely compared agaisnt such subjective human data, making this comparative evaluation both novel and scientifically signficiant.

The purpose of this thesis is therefore to evaluate how the Llama-3 model from NLP Cloud and Hume AI recognizes emotions from Swedish speech, to assess whether its transcribed textual content can convey emotional states independently and compare these AI-generated labels with self-reported emotions from Swedish speakers. By addressing the specific challenge of emotion recognition in a less-studied language, the study contributes to the broader scientific discussion on emotion-recognitions generalizability. The study will provide insights into alignment between speech and text modalities, cultural emotional expression, and the alignment between AI outputs and human experience.

To explore speech emotion recognition for Swedish speech, vocal markers from Swedish speech recordings will be extracted and compared to a prior study (Ekberg et al., 2023). With the usage of this research, the performance of an AI model for Swedish can be compared, and therefore the first research question of this study is:

[1] How does an AI model for speech recognition compare to research on vocal markers for emotions in Swedish speech?

Text-based emotion recognition is a commonly used research field, but mostly for English text. To address this, it is interesting to assess whether transcribed Swedish speech can reveal emotions independently, which leads to the second research question:

[2] What similarities and differences emerge between emotions detected from audio features and from the textual transcripts of the same speech data?

The perception of emotions is a complex field, with few studies made on the alignment between machine-labeled emotions and human-perceived emotions. To undertake this, its comparison will be explored in the third research question:

[3] How do AI-generated emotion labels (speech & text-based) compare to self-reported emotions?

1.4 Scope and Limitations

The scope focuses on AI-based emotion recognition in Swedish speech and text, considering its constraints in design and resources. The study explores challenges like reliance on acted datasets, language differences, and the alignment between AI predictions compared to self-reported emotions. Since this is an exploratory thesis, some limitations are recognized but accepted for feasibility.

1.4.1 Scope

The study evaluates AI-driven emotion recognition in Swedish, a language with little prior research in this area. It analyzes emotions from about 15 Swedish-speaking participants through short interviews designed to evoke natural emotions. Participants, both male and female, have varying age from 20-78 years. The study includes:

- **Vocal Extraction:** With Praat Parselmouth, a Python library for Praat software used for feature extraction from audio recordings.
- **Speech-based analysis:** Using Hume.ai (Hume, n.d.-a), an API with AI-based emotion recognition in speech for AI-based speech emotion recognition.
- **Text-based analysis:** Using NLP Cloud (Cloud, n.d.), an API utilizing AI to transcribe speech and detect emotions from text.
- **Comparison with self-reports:** Participants rate their emotions on a scale of 1-6 (1 = very weak, 6 = very strong), compared to AI-generated labels.

To keep this study manageable, it focuses on two semantic orientations, one positive and one negative designed interview for each participant. Five emotions are derived from the audio and are reported by the participants.

The analysis relies on existing AI tools and API's (Cloud, n.d.; Hume, n.d.-a) and Praat software for voice feature extraction, without developing new models. A mixed method is used, combining AI outputs with qualitative insights.

1.4.2 Limitations

Several factors limit the study's depth and generalizability. With only 16 participants, the dataset is limited, and the results may not apply to all Swedish speakers as well as the findings may not apply beyond Swedish. The interviews are designed to elicit emotions and may not fully capture natural emotional responses, since they are partially induced, and the very nature of the interview setting cannot be directly applicable to real-world environments. The design of the emotion-eliciting scenarios may not be optimal because of deficient psychological expertise, even if the scenarios are based on prior research. By the same reason, the composition by the self-reports could be a limitation in combination with the subjectivity of participants' emotion reports, that may be influenced by personal biases or recall inaccuracies. The selected emotion

categories, commonly used in prior research, include more negative (anger, fear, sadness) than positive (joy) oriented emotions, leading to potential limitations on self-reports. Focusing on these emotions may exclude other relevant emotional states, as Hume AI's output in fact cover several more emotion labels. Pre-trained AI models are utilized without modifying their algorithms, which may introduce biases. The study does not include biometric data, which could provide additional insights. For vocal extraction, Praat Parselmouth is applied, which in our implementation, does not cover the full set of vocal features included in the Swedish research (Ekberg et al., 2023) used for comparison in RQ1.

These limitations are necessary compromises for feasibility within the study's timeframe and resource constraints. The study does not aim to develop new AI models or solve all SER challenges. Instead, it provides initial insights into Swedish emotion recognition, tests existing AI tools, and identifies areas for future research.

1.5 Disposition

From here, the report is structured as follows:

Method and Implementation: This section introduces explanatory mixed method, experimental approach used to answer the research question. It describes the experimental setup, data collection process, methods of analysis, and considerations regarding validity and reliability.

Theoretical Framework: This chapter explores the underlying theories relevant to this study. It provides an overview of Natural Language Processing (NLP), Speech Based Emotion Recognition (SER), Hume.ai, Praat Parselmouth, Text-Based Emotion Recognition (TBED), NLP Cloud, and theories behind vocal markers in speech. The experiment is explained with relevant research for the interviews used for this study.

Time Plan: This section outlines the remaining phases of this thesis, detailing the planned milestones and schedule.

Method and Implementation

This chapter outlines the work process for this study, designing a methodical approach to investigate emotions in Swedish speech using both AI-based analysis and self-reported data. The chapter describes the study's approach and design, justifies methodological decisions, provides details regarding data collection and analysis procedures, and addresses validity and reliability considerations.

2.1 Approach and design

This study adopts an explanatory sequential mixed method approach, which integrates both quantitative and qualitative approaches in a structured sequence. The study first collects and analyzes quantitative data, as AI-generated emotion labels and self-reported emotions, and then qualitatively interprets the results to explore alignment and divergence. This approach ensures a systematic, layered analysis rather than pure comparison (Creswell & Creswell, 2023).

The study follows a deductive research approach, as it builds upon existing theories of emotional expression in text and speech. The AI models will be tested and compared to established findings. Instead of developing new theories, the study aims to evaluate whether AI-based emotion recognition methods align with each other, prior research on vocal emotion markers and self-reported emotions for Swedish speech. This is classified as an experimental study, as it involves a controlled setting where participants are asked questions on predefined emotional recall scenarios. It does not manipulate independent variables in a traditional experimental way (Creswell & Creswell, 2023), instead observes and analyzes the natural emotional responses provoked through structured questions (Bryman et al., 2022).

The study evaluates AI-generated emotion labels from speech compared with existing research on vocal markers, text-based emotion recognition and self-reported emotions. The self-reports serve as a reference point and not a ground objective truth, to acknowledge the subjective nature of emotional perception.

2.2 Data Collection

The study involves participants for semi-structured interviews where they respond to predefined scenarios to provoke emotions. Ethical considerations, such as informed consent and anonymization, are followed firmly to ensure participant well-being. In the first phase, interviews are collected for analysis.

16 Swedish speakers, primarily acquaintances to the researchers, are recruited via invitations. Interviews include 2 scenarios, each scenario includes 5-7 questions designed to elicit either positive or negative emotions, the participants select one of these questions for each semantic area, to maintain freedom in the speech as well as avoiding asking questions the participant are not comfortable to answer. The participants are not pre-informed about the feeling that are aimed to be provoked during the interviews. The scenarios have been pilot-tested for effectiveness. The interviews are recorded in a quiet room. Each scenario last between 2 and 4 minutes with breaks between them, the order of the scenarios varies to minimize affecting the

results because of the possibility of the order influencing the different emotions. Each recording have been edited to delete our questions and any longer silent pauses.

Participants are asked open-ended questions designed to bring out previously lived through personal experiences of anger and happiness. The questions about anger are focused on previous experiences of unfair treatment and frustration regarding their everyday lives or society, while the questions related to happiness explore moments of pride and unexpected joy through memories. The semi-structured format allows for follow-up questions based on participant responses, to bring out as much emotion as possible. The follow-up questions include "How did that make you feel?", "Can you evaluate on that specific situation?", and "What feelings did you experience?".

Two example questions that can be selected in each scenario are stated below: **Negative:**

- Think about a situation when something did not go as planned. What was it? What feelings did you experience?
- Is it something about society that makes you upset? What? Can you elaborate?

Positive:

- Describe a situation when something unexpected raised immense happiness in you. How did you feel in that moment?
- Think about a moment when you were very proud of yourself. What did you do and how did you feel?

The study adopts a mixed-method interview approach (Bryman et al., 2022), where the qualitative data from speech is transformed into quantitative AI-generated emotion labels, to enable comparison in a structured way (Creswell & Creswell, 2023). Audio is recorded and pre-processed to reduce background noise. Vocal markers from each recording are extracted using Parselmouth, a Python library for vocal feature extraction, to answer RQ1. The audio is analysed simultaneously using two emotion recognition models: Hume.ai to extract speech-based emotional labels, and NLP Cloud to transcribe the speech and then analyse the textual content in terms of emotion scores. The same dataset is used for all research questions to ensure consistency.

The diagram in Figure 2.1 visualizes the multi-modal pipeline used in this study. The interview audio files are processed through three primary channels: speech-to-text transcription via NLP Cloud, Speech emotion recognition via Hume AI and acoustic feature extraction via Praat. These channels represent two main pipelines. The entities presented in yellow are prevalent in both pipelines, where the audio recordings are analyzed with Hume AI, the output is filtered to the 6 emotions analyzed in this study. The pipeline illustrated in green represents the analysis to answer research question 1. The vocal features extracting utilizing Praat Parselmouth are chosen are based on previous research, see 3.5, Theoretical Framework - Vocal Markers, where pitch, intensity/loudness, formant frequencies (F1, F2, F3), HNR, jitter and shimmer have distinguished values for certain emotions. To compare the extracted data with Hume AI, these values are clustered into emotion groups. Data from speech analysis and vocal markers are combined to statistically analyze the results for RQ1. The pipeline used to answer the second and third research question is presented in orange. Interview audio is processed the same way as for RQ1 but extended with normalization for the Hume values to enable comparison with outputs from NLP Cloud and self-assessment. For text-based emotion recognition, the recording is transcribed before text-analysis is composed. Results from speech and text prediction are combined with the self-assessment scores to answer RQ2 and RQ3.

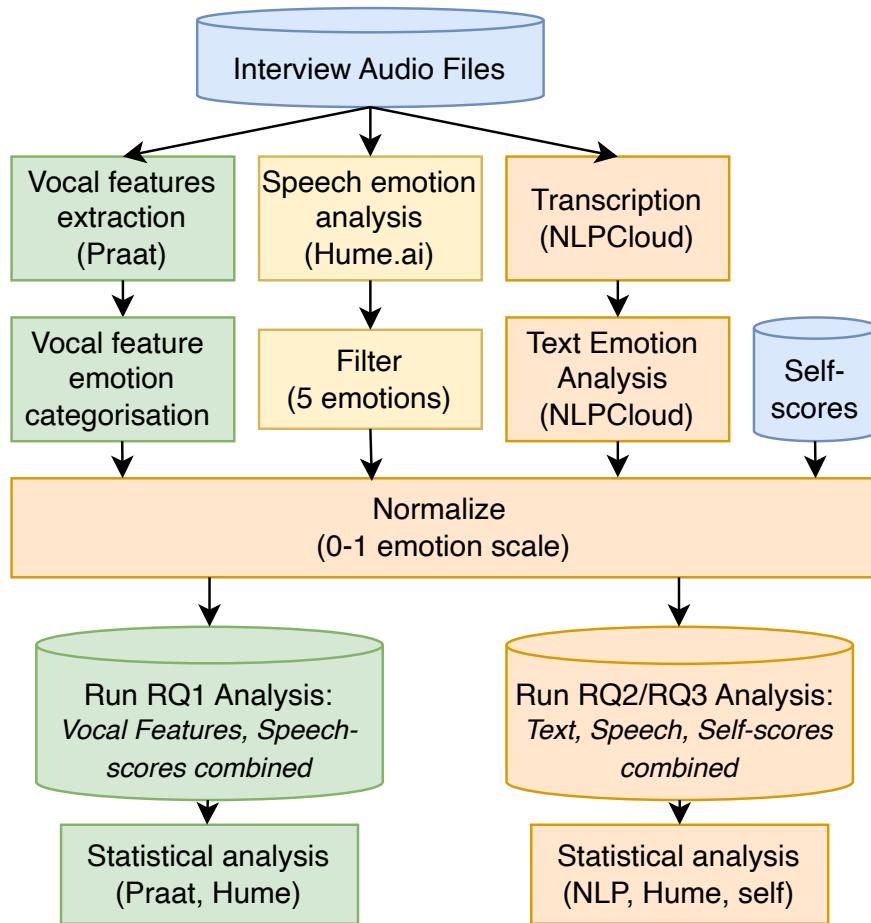


Figure 2.1: The multi-modal pipeline used in this study.

2.2.1 Research Question 1

How does AI-model for speech emotion recognition compare to research on vocal markers for emotions in Swedish speech?

To answer this question, speech recordings are collected from participants as they describe emotionally charged experiences. AI-based emotion recognition using Hume.ai, are used for AI-based Speech Emotion Recognition. Voice feature extraction from the recordings is made, to compare to AI-labeled emotions with known vocal markers from existing Swedish emotion research (Ekberg et al., 2023).

2.2.2 Research Question 2

What similarities and differences emerge between emotions detected from audio features and from the textual transcripts of the same speech data?

To answer the second question, the recorded speech is transcribed and analyzed for emotion recognition using NLP Cloud's emotion recognition to assess the emotional content of speech transcripts. The text-based AI labels are compared with speech-based AI labels to determine whether emotion is preserved in textual content alone.

2.2.3 Research Question 3

How do AI-generated emotion labels (speech & text-based) compare to self-reported emotions?

For the third question, participants complete a self-assessment survey after each interview segment, where they rate their emotional state on a 1-6 scale (1 = very weak, 6 = very strong) for relevant emotions. The self-reported emotions are compared with AI-generated labels from both speech and text models to analyze agreement and divergence. The results are clustered as agreements, partially agreements, and disagreements across methods.

2.3 Data Analysis

To systematically evaluate the agreement between different emotion detection methods, a combination of statistical analyses and visualizations was applied for speech-based AI, text-based AI, vocal markers, and self-reported emotions. The analysis aimed to assess the alignment with established vocal marker research and subjective human perception, where identification and categorization of differences were analyzed.

2.3.1 Statistical Analysis Approach

RQ1: Vocal Markers vs. Speech-Based Emotion Recognition

The first step involved applying a custom emotion categorization function based on standardized distances from the Swedish research on vocal markers (Ekberg et al., 2023), to cluster vocal features into emotional categories. This approach adapted absolute standardized deviations across key acoustic features, as detailed in section 3.5 and 3.7.

Following the initial clustering, correlations were calculated between:

- Individual vocal features and Hume AI emotion scores.
- The categorized vocal emotion scores and corresponding vocal features.
- The categorized vocal emotion scores and Hume AI emotion scores.

Limitations of the customized categorization method were identified during this phase and are discussed in the results.

For further exploration of how vocal features varied across AI-detected emotions, one-way ANOVA tests were conducted, followed by Tukey's HSD analysis to evaluate specific feature differences, see 3.7 Statistic Analysis for details.

To track changes over time, each audio recording were segmented into timeframes of around 5s including Hume AI emotion scoring and vocal features for that specific segment. These were analyzed by tracking Z-score variations, see 3.7, in key vocal features (pitch, intensity, jitter, and shimmer) throughout each recording. This enabled observation of how vocal features shifted within a single clip and to see how these changes matched Hume's emotion predictions.

Custom Emotion Categorization Method

To complement the correlation analysis with Hume AI, a custom rule-based categorization method was integrated to group emotion probabilities from extracted vocal features. The approach was based on the Swedish research on vocal markers in emotions (Ekberg et al., 2023), where statistical patterns were identified in vocal markers across five emotions: anger,

joy, sadness, fear, and surprise. The Swedish research did not state their used baseline for neutral speech, therefore the baseline for this study is the average vocal features of all clips.

The method functions as following:

- 1. Predefined means and standard deviations for each vocal features identified for each emotion were retrieved from the Swedish research (Ekberg et al., 2023). These features are stored in JSON format.
- 2. For every feature included in a recording, the function calculates the standardized distance between the measured vocal value and the mean for each emotion:

$$d_{\text{emo}} += \frac{|x - \mu|}{\sigma}$$

where x is the observed value, and μ, σ are the mean and standard deviation for the pair of feature and emotion.

- 3. To increase the functionality, distances are inverted so smaller distances can result in higher emotion scores. This is calculated with the function below, where ϵ is a small constant to avoid division with zero.

$$\text{score}_{\text{emo}} = \frac{1/(d_{\text{emo}} + \epsilon)}{\sum_e 1/(d_e + \epsilon)}$$

- 4. The output is a normalized probability value that is distributed across all five emotions.

The method allowed to map vocal data to emotion estimates in an interpretable way to enable structured comparison with speech-based emotion labels. However, the performance of this method for our dataset had limitations and gave uniform emotion labels.

2.3.2 RQ2 and RQ3: Speech-Based AI vs. Text-Based AI, AI-labels vs. Self-Assessed Emotions

For both RQ2 and RQ3, statistical analyses were used to evaluate the alignment between AI-generated emotion scores and self-reported emotions. The following methods were applied across both research questions:

- Pearson correlation coefficients and associated p-values for measurement both for single clips and the full dataset.
- T-tests and calculation of Cohen's d to evaluate statistical significance and effect sizes.

For all statistical tests, a p-value below 0.05 was considered as statistical significant, implying that the observed correlations or differences were unlikely to have occurred by chance. These statistical methods were applied to the entire dataset, and separately for negatively and positively oriented interview scenarios to identify potential contextual differences. For RQ2, comparisons focused on speech-based vs. text-based AI results, and RQ3 extended the comparison to include self-reported emotions as a subjective component.

2.3.3 Data Normalization

To enable direct comparison across different sources, all emotion scores were normalized to sum up to 1. The normalization included the following steps:

- 1. Surprise combination: Hume AI predicted two separate labels for Surprise, one positive and one negative. These were merged into a single "surprise" by calculating the average.
- 2. Filtering and formatting: Filtering to only include the five target emotions (anger, joy, sadness, fear, surprise), since Hume predicted around 30 different emotions. All emotion labels were converted to lowercase.
- 3. Normalization: The emotion scores were normalized so the sum of all five target emotion values equals 1. This was done by dividing each score by the total sum of the emotion values. If the total sum was zero (no emotion detected), all scores were set to 0.

Normalization ensured consistent comparability between the sources, for both AI models and self-assessments, regardless of scale differences in raw scores.

2.3.4 Visual Analysis

RQ1:

Visualizations included:

- Heatmaps showing correlation matrices.
- Line plot diagrams for the full dataset.
- Bar chart diagrams for single clips.
- Composite correlation diagrams.
- Over-time diagrams comparing vocal features and Hume labels.

To evaluate the performance of the custom vocal feature-based emotion categorization method, line plots and bar charts were implemented to visualize differences between the generated scores and Hume AI's predictions. The line plot summarizes average emotion scores across the full dataset, while bar charts presented detailed comparisons for individual audio recordings. These diagrams emphasize deviations and alignments between the categorized vocal marker method and AI-based emotion prediction.

Composite correlation diagrams were used to explore associations between single vocal features and AI-generated emotion scores. For these diagrams, Pearson correlation coefficients were calculated for each emotion and its relation to pitch and intensity, the results are visualized as grouped bar charts to easily identify positive or negative tendencies.

RQ2 and RQ3:

Visual methods included:

- Difference bar charts to illustrate the average difference in Hume AI and NLP Cloud's emotion scores for each emotion category, to provide a clear view of systematic deviations.
- Grouped bar charts to present average emotion scores from both AI systems, separated by interview sentiment (positive vs. negative), to explore variations in emotion detection depending on the interview context.
- Scatter plots (RQ3) to explore detailed relationships between AI-based emotion scores and self-reported emotions for single emotions and recordings.

These visualizations were integrated to support the identification of alignment patterns between the AI systems, and contributed with insights into how the modality of the AI influences emotion recognition results. Python have been utilized to create all visualizations.

Given the limited dataset size and timeframe, a combination of statistical methods and visual analysis, was utilized to balance quantitative data with qualitative interpretation and support the exploratory nature of this study.

2.4 Model Configuration

2.4.1 NLP Cloud

The text-based emotion recognition is classified with NLP Clouds finetuned-llama-3-70b model through prompting, which allows a more flexible approach. Each text input uses the following prompt:

Listing 2.1: NLP Cloud configuration prompt.

```
prompt = (
    "You are an emotion analysis system.
    Given a Swedish text, respond only with a JSON
        object using these emotion labels:
    joy, surprise, fear, anger, sadness.
    Each value must be a float between 0.0 and 1.0.
    Respond with the JSON directly and nothing else.
    f'{transcription}'"
)
```

The prompt that is used returns a JSON response with float values ranging from 0.0 to 1.0 for each of the emotions with the labels “joy”, “surprise”, “fear”, “anger”, “disgust” and “sadness”. This approach was chosen to ensure these specific emotions being analyzed due to them being the feelings of the basic six, which are the feelings used in the research about acoustic features in swedish speech done by Ekberg (Ekberg et al., 2023).

2.4.2 Hume AI

To ensure consistency across the different models used in this research, some changes have been made to adjust the output from the Hume AI model to better match the format used in NLP Cloud. Additionally, NLP Cloud has the feeling surprise while Hume has two different feelings for surprise, the two being positive surprise and negative surprise. Therefore, the scores of the two feelings of surprise from the Hume model have been combined in this research to give just one number that creates the average of the two to match the format.

2.5 Validity and Reliability

2.5.1 Validity

To ensure validity, the interview scenarios are pilot tested to ensure they provoke intended emotions (Bryman et al., 2022). The use of multiple AI models (speech- and text-based) allows for cross-validation of results. Standardized interview prompts ensure consistency across participants. Participant self-assessment serves as a secondary reference to evaluate AI-labeled

emotions. Triangulation across AI, vocal markers, text analysis, and self-assessments enhance convergent validity (Creswell & Creswell, 2023).

2.5.2 Reliability

To ensure reliability, standardized equipment and scenario are used to ensure replicability. Hume.ai, NLP Cloud, and Praat provide consistent measures. The AI models used in the study (Hume.ai and NLP Cloud) are pre-trained and validated emotion recognition systems. Correlation will be determined and are used to quantify the reliability of AI models in detecting emotions. The same data is not analysed multiple times to check if the results are different. The prompt used for NLP Cloud is therefore zero-shot. The study has a replicable experimental setup, with documentation supporting replication to allow researchers to reproduce similar evaluations.

Triangulation is achieved in the study through comparison of speech AI, text AI, and self-reports which improves creditability. Any discreteness will be analyzed qualitatively to contextualize potential biases rather than assuming errors. Reliability is ensured through standardization in data collection. All interviews are preprocessed to reduce background noise and normalize volume levels. The online tool Auphonic (Auphonic, n.d.) is used for this, due to its simple usability for noise reduction, ability to cut out pauses and limit loudness. The same data processing steps are applied consistently for all recordings, ensuring equality in analysis. The study has a replicable experimental setup, with usage of pre-trained, publicly available APIs, and documentation supporting replication to allow researchers to reproduce similar analyses. These measures ensure that our study is generalizable within the scope of automated emotion recognition for stress analysis.

2.6 Considerations

To consider the implications of this study, several factors must be recognized. To address ethical and privacy concerns, all participant data is anonymized and securely stored to ensure privacy. The participants provide informed consent before engaging in this study. The emotion-provoking scenarios are designed to minimize distress, focusing on natural, everyday emotions rather than triggering events. The participants will have scenarios to choose from, see 2.2 Data Collection.

Scientific considerations extend to emotion research to Swedish speech and AI tools. Findings in the study can inform future human-interaction research in emotion-based applications. Societal considerations include that the insights could enhance AI-driven mental health tools and future research, especially for Swedish language and real-world interviews.

Theoretical Framework

The following chapter will introduce the relevant theories and key concepts related to emotion recognition, such as speech-based and text-based models and technologies. This chapter will explain how vocal features and speech prosody can help to identify different emotions in spoken languages, using different AI tools and software. This study aims to compare accuracy and effectiveness of different approaches by conducting interviews and collecting data, which will be analyzed using a Python application. The elements of the Python application for analyzing the data from the interviews will be comprehensively explained in this chapter.

3.1 Affective Computing

Affective computing was introduced by Professor Rosalind Picard in the mid-1990s to early 2000s (Tian et al., 2022). By exploring the ways in which human emotions are recognized, understood and expressed through different forms of behaviors and communication, the domain of affective computing is a field that merges the principles of artificial intelligence with insights from social and behavioral science (Tian et al., 2022).

3.2 Natural Language Processing and Emotion Recognition

The first English language lexical database was created in 1998 for Natural Language Processing (NLP) tasks, the term sentiment and emotional analysis came to practice in 2001 to predict the stock market (S. Kusal et al., 2023). In 2005, the first article was written on emotion and opinion detection from text. Concept-level sentiment analysis resources were publicly available in 2009. Word embedding is the term to represent words for NLP text analysis and was developed in 2013, the same year as neural network first was adopted in NLP tasks. The field had a massive upwelling when the transformers concept was published in 2017, followed by the evolution of BERT, a pretrained model that automated text analysis and classification in 2018 (S. Kusal et al., 2023).

Traditional approaches for sentiment analysis classification have been used since the past few decades, which rely on rule-based methods such as “bag of words” method to process text (Kansara et al., 2020). The method represents text based on word frequency without consideration of word order. It can identify sentence structure, negation, emphasis, subjectivity and irony. Recent models leverage deep learning algorithms that process raw text by first cleaning and pre-processing it, including punctuation, stop words and markups, as well as applying stemming (the process of reducing words to their root form by removing prefixes or suffixes to simplify text analysis in NLP).

Deep learning applicate artificial neural networks (ANN) to learn tasks using multiple layers of network. In traditional models only one or two layers could be used, but in deep learning much more learning power of artificial neural networks is exploited (L. Zhang et al., 2018).

Studies have shown consistently higher accuracy for sentiment analysis using deep learning algorithms compared to traditional machine learning algorithms (Kansara et al., 2020).

3.3 Speech-Based Emotion Recognition

Studies about speech-based emotion recognition (SER) have been ongoing since 1978 (Sönmez & Varol, 2024). SER identifies how something is being said without the context of the words spoken. These systems are used in many different areas, most often in areas of interactions between humans and machines (Zhang, 2025). A typical SER system contains of three components: signal preprocessing, feature extraction, and classification (Sahoo et al., 2023).

Although there is a wide range of SER-algorithms, with some approaches using more complex setups that involve Convolutional Neural Networks (CNN) based SER algorithms among others (Ri et al., 2023), the process of a SER algorithm could look like figure 3.1, involving several steps as feature extraction, selection and classification.

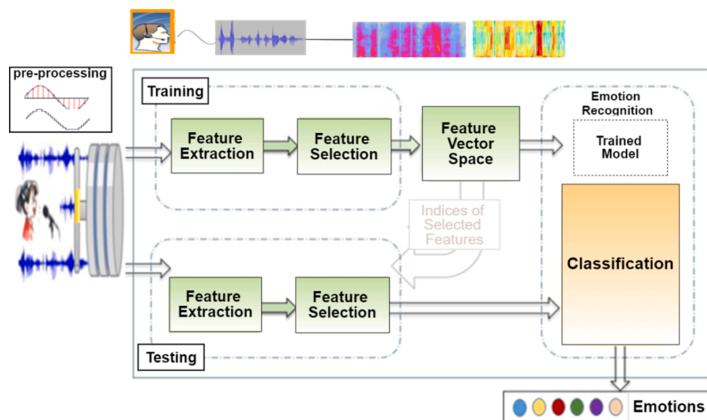


Figure 3.1: *An overview of the stages in SER to analyze speech data for emotion detection* (Sönmez & Varol, 2024)

While emotions can be recognized via many channels such as speech, facial expressions and text, speech signals are rapid and natural which makes vocal audio fitting for emotion recognition. According to (Sönmez & Varol, 2024) there are several key benefits with SER, such as a limited amount of hardware needed for the capturing the vocal data which simplifies the process of the vocal data collection. Another benefit is that vocal data being less demanding in terms of storage, compared to video footage for example, and participants in SER experiments may feel more comfortable with vocal analysis than face analysis in terms of confidentiality, resulting in datasets reflecting real emotions more accurately.

3.3.1 Hume AI

Hume is a technology company dedicated to advancing the field of emotion recognition. Having conducted extensive psychology studies to explore human emotions and the way these emotions are expressed, Hume AI has used the research to develop advanced machine learning models (Hume, n.d.-b) as well as using deep learning for the research and development (Brooks et al., 2023).

The official website of Hume AI outlines several influences on their emotional mapping. Drawing influence from key figures such as David Hume, Charles Darwin and Paul Ekman, Hume AI's research is grounded in these foundational theories of emotions. Paul Ekman's "The Basic 6" is mentioned (Hume, n.d.-a) and remains relevant throughout this research.

One of the measurements used to recognize emotions in vocals with Hume AI in this research is speech prosody.

Speech prosody gives crucial insights into a speaker's purpose in their communication. Particular emotions and the intensity of those emotions are indicated with intonation, rhythm and pitch of the speaking voice (Thompson et al., 2004; Tomasello et al., 2022). It simply refers to the patterns and tone in the speech that are not related to the actual words being spoken (Cowen et al., 2019).

Happiness and sadness show the opposite characteristics of each other, where happiness is linked to quicker tempo and higher pitch while sadness has the opposite, a slower tempo and lower pitch. The clear difference between the characteristics serves as the difference in the speech prosody of the two emotions (Thompson et al., 2004).

In figure 3.2, a visual presentation of Hume AI's speech prosody model is visible (AI, n.d.-a). Emotions are clustered with other similar emotions, one example being amusement and joy, or distress and anxiety.

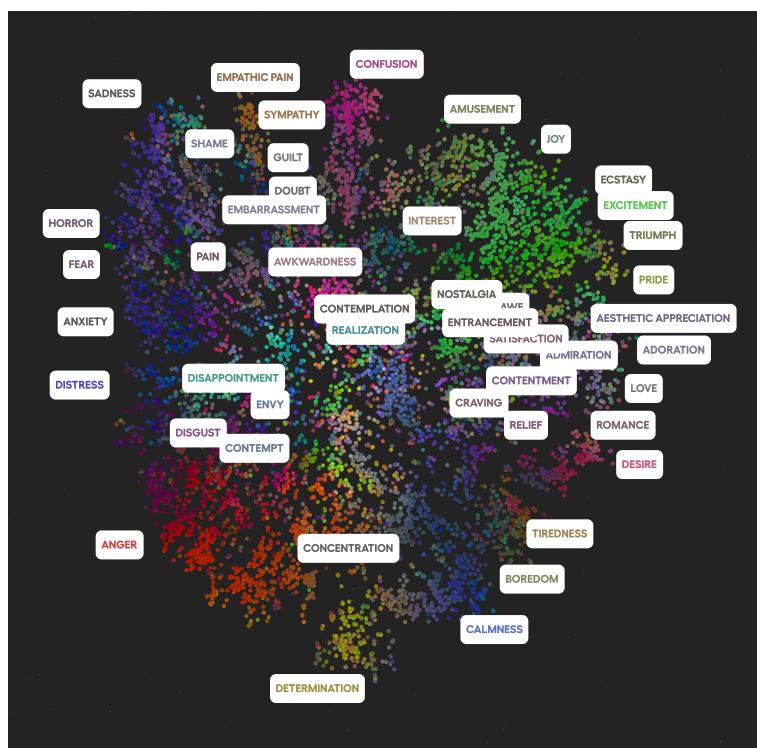


Figure 3.2: *Visual representation of Hume AI's speech prosody model* (AI, n.d.-a)

To ensure a broader range of emotion recognition with a more comprehensive analysis of human voices in this research, speech prosody is used in combination with another measurement, vocal bursts.

Vocal bursts play a key role in social communication between humans. They are short emotional sounds which occur naturally, some examples being laughs, sighs or cries (Brooks et al., 2023).

Vocal burst and voice have received less attention in the fields of machine learning and affective computing due to more focus being held on facial expressions. Even if speech prosody has been studied more extensively, there has been newer research showing that vocal bursts convey more than ten different emotions with consistency, being mostly consistent across different cultures as well (Baird et al., 2022).

Figure 3.3 shows Hume AI's mapping of non-verbal communication, vocal bursts (AI, n.d.-b). Emotions are shown and as well as in Hume AI's speech prosody model, the emotions are clustered indicating some emotions are more associated with each other.

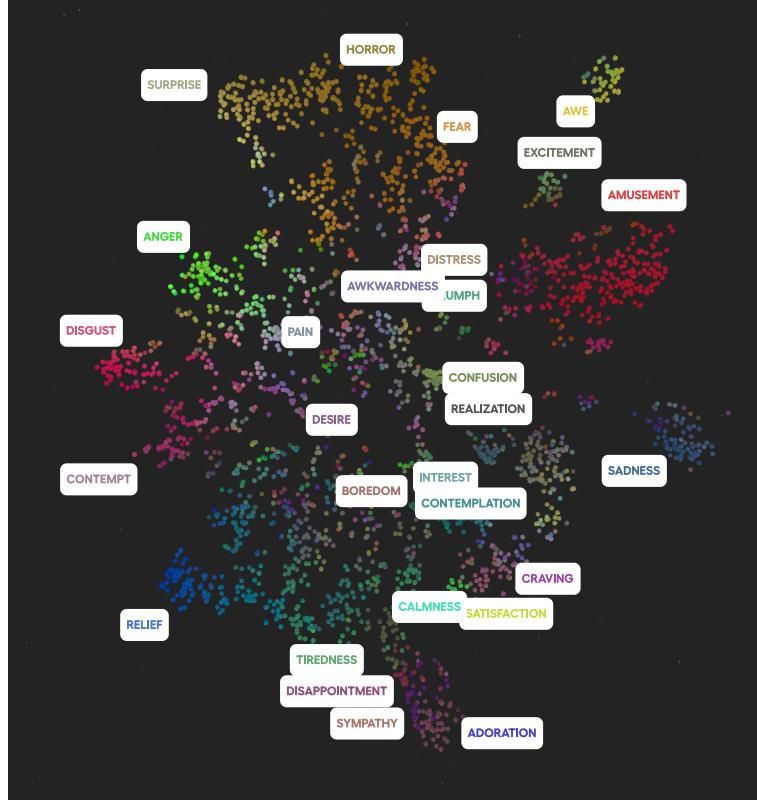


Figure 3.3: *Visual representation of Hume AI's vocal burst model* (AI, n.d.-b)

While there are other tools for emotion recognition, Hume AI is among the few that are specifically designed for Voice AI while being able to recognize emotions through specifically speech prosody and vocal burst with no need to finetune it yourself. Using models needing finetuning would not fit the scope of this thesis given the limited timeframe, and while there are other models, like OpenAI whisper, which is an extensively trained model on hundreds of thousands of hours on data, their main focus leans toward transcribing speech (OpenAI, 2022). This ultimately led to the decision to use Hume AI.

3.3.2 Praat Parselmouth

In the field of software for linguistic analysis, Praat is a well-established tool to analyze different elements in speech. Being able to estimate elements such as fundamental frequency and intensity among others, Praat holds a broad spectrum of algorithms in acoustics, being a successful tool for analyzing acoustics (Jadoul et al., 2024).

Designed to provide efficient access to the core functionalities of Praat in Python for programmers, Parselmouth is an open-source Python library (Jadoul et al., 2018).

Python is widely used for data analysis, but it had been noted that there were challenges with analyzing acoustics in Praat, this due to the functionality often being missing or scattered across multiple incompatible libraries.

Parselmouth streamlines and optimizes workflows in a single programming environment by enabling a deeper integration of the capabilities of Praat in combination with other libraries (Jadoul et al., 2024). Not designed to replace Praat, but rather a way to enable users to access the functionality of Praat directly in Python, there some main objectives in Parselmouth according to (Jadoul et al., 2018). One objective is to enable users already experienced with Praat to effectively incorporate its functionality with Python's scientific tools, being tools that are not obtainable in Praat. Providing Python users with the ability to utilize the functionality

of Praat, even if they are not experienced users of it, is also an important aspect, as well as enhancing the optimal aspect of workflow for users preferring to conduct their work within a single programming language.

The benefits of Parselmouth both in terms of the usage for completing this thesis and overall, are it being open source and compatible with Python as Python is widely utilized and backed by a vast community of researchers and engineers, among others. Parselmouth integrates the different strengths of different approaches to provide a library following the principles of Python and behaving consistently with other well-known Python libraries.

Parselmouth directly utilizes the official C/C++ source code of Praat instead of having to reconstruct its algorithms. This simplifies the process since it guarantees full consistency with Praat without the requirement of learning its scripting language (Jadoul et al., 2018).

There are other similar tools that essentially could accomplish the same task, like Librosa although it is more tailored for both audio and music analysis. It does have feature extraction (Babu et al., 2021), but the decision on which software to use for linguistic analyzation still falls on Praat Parselmouth due to it being more fitting for the purpose of this thesis.

3.4 Text-Based Emotion Recognition

In the field of NLP, the comprehension of the context behind words in text-form has gone from only being able to determine the tone in text to actually identifying the emotions behind them (Esfahani & Adda, 2024), recognizing these capabilities has valuable practical applications in enhancing different domains within human-computer interaction (Shelke et al., 2022).

Text-based emotion detection relies on four main approaches, according to S. Kusal et al. (2023). The first approach is keyword-based, which matches words in a text with predefined emotion keywords from resources like WordNet, adjusting for intensity and negation. The second approach is rule-based, and this approach uses linguistic rules and probabilistic affinity to classify emotions after preprocessing. The third is machine learning-based which applies supervised or unsupervised models to classify emotions, extracting key features from preprocessed text. The fourth and last approach is deep learning-based and it utilizes neural networks to learn complex patterns from tokenized and embedded text data for emotion classification.

Machine learning classifiers are significantly used in text-based classification, since they use labeled datasets and are therefore data driven. Machine learning models are trained on large number of datasets and learn from experience, with classifiers that contain labels for input and desired output. Transformer-based models, such as BERT, are based on machine learning models which are trained on vast amounts of data and can be fine-tuned for specific tasks. BERT is a deep learning model based on attention processing. It gains a thorough text-understanding through considering left and right contexts equally. The model solves NLP issues and is used to train general language models on large datasets (S. D. Kusal et al., 2024).

3.4.1 NLP Cloud

There are limited publicly available APIs for text-based emotion detection. The decision to use NLP Cloud has several reasons. Compared to other available TBED APIs, that do not require pre-training, NLP Cloud has comprehensive documentation for both API and the models it is based on. The company has support as well as information about Data Privacy and Security (Cloud, n.d.), which also other APIs for TBED has. For example, Vern AI (AI, u.d.) provides customer support but has very limited information about the API or documentation that is easily accessible. TwinWord (TwinWord, n.d.) was another choice, also providing contact support and privacy information, but was deficient in comprehensive documentation and limited information about the API.

NLP Cloud was the obvious choice for several reasons. In addition to what is mentioned above, the company has solid customers like Zoom and collaboration with Nvidia. They provide information about what models they use for their API which can be downloaded and fine-tuned if desired. In contrast to the other APIs, NLP Cloud provide APIs for Speech-to-Text transcription, and they have an emotional analysis model supporting Swedish Language. Therefore, no translation is required beforehand. Their speech-to-text API is based on OpenAI's Whisper model (Cloud, n.d.). OpenAI provides a research report on the model, which is a speech recognition system designed to process and transcribe audio with remarkable robustness and generalization. Contrasting traditional models that heavily rely on unsupervised pre-training or dataset-specific fine-tuning, Whisper leverages large-scale weakly supervised training from over 600,000 hours of multilingual audio data. This includes 96 languages beyond English. Whisper handles several tasks, for instance speech recognition, language identification, and translation (Radford et al., 2022).

For sentiment and emotion analysis, NLP Cloud provides several options. Two of them are equally researched with relatively comparable results. DistilBERT Base Uncased Emotion is one option, with studies on DistilBERT, and Transformer-model which require finetuning and require less computational resources than traditional BERT (Bidirectional Encoder Representations from Transformers) models. The model was developed by Google AI Language researchers in October 2018. BERT has challenges regarding fixed input length and computational complexity, reasons that led to DistilBERT's development in October 2019. This pre-trained model uses technology that accomplishes to reduce a BERT model by 40% while preserving 97% of language understanding capabilities 60% faster. The model accuracy for sentiment analysis ranges from 95.7% to 96.6% on Yelp Open Dataset, which has been demonstrated as a valuable resource of sentiment-labeled review-data (Areshey & Mathkour, 2024). NLP Cloud provides a pre-fine-tuned version of DistilBERT. However, it does not support Swedish language and translation beforehand is necessary for this study. A fine-tuned version of Llama 3 is an option on NLP Cloud that support several languages including Swedish (Cloud, n.d.). Opposed as to DistilBERT, some studies on emotion detection have been conducted on Llama 3. Both models have predominantly been evaluated for sentiment analysis. Emotion identification using Llama 3 showed an F1 score of 0.48 as average for all tested emotions: anger, joy, sadness, surprise, fear, and love (J. Zhang & M, 2024). A fine-tuned version of Llama 3 was tested for sentiment analysis, which resulted in an accuracy increase from 0.333 to 0.923 and F1 score from 0.50 to 0.91. The authors of this study imply that these results are superior performance against other models that are included for comparison, including DistilBERT (Kumar & Singh, 2024). In a text classification study, mainly examines speed performance, DistilBERT had an accuracy between 0.94-0.96 for the Amazon Alexa Reviews dataset but 0.35-0.41 for the Brexit Blog Corpus dataset (Barbon & Akabane, 2022). As for many models for emotion recognition mentioned in this report, the dataset affects the results heavily.

Regarding multiple language performance, a study examined LLaMA 3 vs. State-of-the-Art Large Language Models on their ability to detect fake news. Two datasets were tested, one Romanian and one English. Their proposed Llama 3 model accomplished higher precision and accuracy across several metrics in fake new detection. For the English dataset, the fine-tuned Llama 3 model had lower accuracy compared to ChatGPT 4 and Gemini. Yet, it outperformed these models for the Romanian dataset. The study also explored their fine-tuned Llama model compared to its base version. The fine-tuned model outperformed earlier models in distinguishing nuanced categories, particularly for the Romanian dataset where it achieved a remarkably high accuracy of 68% in one category (Repede & Brad, 2024).

Comparing these two alternatives for text-based emotion detection in Swedish, the fine-tuned Llama 3 model emerges as the most suitable choice. Although the exact fine-tuned version of the model available on NLP cloud has not been publicly researched, its built-in compatibility

with Swedish, combined with research on a Romanian dataset, makes it a stronger candidate than DistilBERT. Both models have achieved high accuracy for TBED. Nevertheless, given this study is aimed to focus on emotion recognition models for Swedish speech, the Llama model without need for prior translation is a more valuable choice.

3.5 Vocal Markers

States of emotions are determined by many different factors, but in speech, different vocal markers such as prosody, pitch and loudness are among the telling features when it comes to emotions (Ekberg et al., 2023), where frequency is perceived as pitch, while amplitude is heard as loudness (Fröhholz & Belin, 2019).

In this research vocal markers are crucial to understand the difference in certain emotions in voice.

Figure 3.4 (Ekberg et al., 2023) shows a table of comparisons with certain parameters of vocal markers measured for five different emotions in the Swedish language and some explanations for the different parameters. This is used in this thesis as a comparison for the vocal markers collected from the interview data in the conducted experiment, although this thesis does not use all acoustic features to measure emotions.

TABLE 2.
Comparisons of Acoustic Parameters Between Emotions

Acoustic features (parameters)	Anger	Happiness	Fear	Sadness	Surprise	Comparisons			
	M (SD)	M (SD)	M (SD)	M (SD)	M (SD)	F	P	pEta2	Post hoc-tests (Bonferroni adjusted)
Frequency-related:									
pitch jitter	5.00 (5.39) -0.13 (0.38)	7.18 (6.25) 0.58 (0.38)	5.81 (2.31) -0.98 (0.41)	3.99 (5.36) 0.32 (0.39)	3.56 (4.14) 2.14 (0.39)	2.77 8.24	<0.001	.178	Happiness > Surprise ($P=0.039$) Surprise > Anger ($P<0.001$) Surprise > Fear ($P<0.001$) Surprise > Sadness ($P=0.014$) Happiness > Sadness ($P=0.012$) Happiness > Sadness ($P=0.008$) Fear > Sadness ($P=0.038$) Happiness > Sadness ($P=0.008$)
F1Frequency F2Frequency	0.78 (0.34) 1.20 (0.35)	1.76 (0.34) 1.94 (0.35)	1.47 (0.37) 1.75 (0.37)	0.12 (0.35) 0.23 (0.36)	0.57 (0.35) 1.03 (0.36)	3.60 3.53	0.008 0.009	.086 .085	Happiness > Sadness ($P=0.012$) Happiness > Sadness ($P=0.008$) Happiness > Sadness ($P=0.008$) Happiness > Sadness ($P=0.008$)
F3Frequency F1Bandwidth	0.80 (0.34) -1.05 (1.29)	1.59 (0.34) -0.96 (0.95)	0.88 (0.37) -0.44 (0.94)	-0.10 (0.35) -0.88 (1.35)	0.72 (0.35) -0.82 (0.88)	2.95 1.38	0.022 0.244	.072	Happiness > Sadness ($P=0.008$) Happiness > Sadness ($P=0.008$)
Amplitude-related:									
shimmer	-1.03 (0.21)	-1.02 (0.21)	-1.43 (0.23)	-1.02 (0.22)	0.13 (0.22)	7.12	<0.001	.158	Surprise > Anger ($P=0.002$) Surprise > Fear ($P<0.001$) Surprise > Happiness ($P=0.002$) Surprise > Sadness ($P=0.002$) Happiness > Sadness ($P=0.003$) Happiness > Surprise ($P=0.001$) Anger > Sadness ($P<0.001$) Anger > Surprise ($P<0.001$) Fear > Surprise ($P=0.001$) Happiness > Sadness ($P=0.003$) Happiness > Surprise ($P=0.001$)
loudness	7.16 (0.66)	6.49 (0.66)	5.09 (0.71)	2.96 (0.68)	1.24 (0.68)	13.36	<0.001	.260	Anger > Sadness ($P<0.001$) Anger > Surprise ($P<0.001$) Fear > Surprise ($P=0.001$) Happiness > Sadness ($P=0.003$) Happiness > Surprise ($P=0.001$) Anger > Anger ($P=0.014$) Anger > Sadness ($P=0.007$) Anger > Surprise ($P=0.001$) Happiness > Surprise ($P=0.004$) Happiness > Surprise ($P=0.043$) Surprise > Anger ($P=0.032$) Surprise > Anger ($P=0.002$) Fear > Happiness ($P=0.006$) Fear > Sadness ($P=0.010$) Fear > Surprise ($P=0.001$) Anger > Surprise ($P=0.042$) Happiness > Surprise ($P=0.019$)
HNR	2.36 (0.52)	3.99 (0.52)	4.83 (0.55)	2.16 (0.54)	1.31 (0.54)	7.09	<0.001	.157	Fear > Anger ($P=0.014$) Fear > Sadness ($P=0.007$) Fear > Surprise ($P=0.001$) Happiness > Surprise ($P=0.004$) Happiness > Surprise ($P=0.043$) Surprise > Anger ($P=0.032$) Surprise > Anger ($P=0.002$) Fear > Happiness ($P=0.006$) Fear > Sadness ($P=0.010$) Fear > Surprise ($P=0.001$) Anger > Surprise ($P=0.042$) Happiness > Surprise ($P=0.019$)
alphaRatio	2.52 (0.40)	2.15 (0.40)	1.14 (0.43)	1.95 (0.41)	0.48 (0.41)	4.05	0.004	.096	Happiness > Surprise ($P=0.005$) Happiness > Surprise ($P=0.043$) Surprise > Anger ($P=0.032$) Surprise > Anger ($P=0.002$) Fear > Happiness ($P=0.006$) Fear > Sadness ($P=0.010$) Fear > Surprise ($P=0.001$) Anger > Surprise ($P=0.042$) Happiness > Surprise ($P=0.019$)
Hammarberg slopevOV500	-1.57 (0.28) 2.53 (0.43)	-1.19 (0.28) 2.68 (0.43)	-0.74 (0.30) 4.90 (0.46)	-1.4 (0.29) 2.76 (0.44)	-0.35 (0.29) 1.81 (0.44)	2.97 6.54	0.022 <0.001	.072 .147	Fear > Anger ($P=0.002$) Fear > Happiness ($P=0.006$) Fear > Sadness ($P=0.010$) Fear > Surprise ($P=0.001$) Anger > Surprise ($P=0.042$) Happiness > Surprise ($P=0.019$) Happiness > Surprise ($P=0.043$) Surprise > Anger ($P=0.032$) Surprise > Anger ($P=0.002$) Fear > Happiness ($P=0.006$) Fear > Sadness ($P=0.010$) Fear > Surprise ($P=0.001$) Anger > Surprise ($P=0.042$) Happiness > Surprise ($P=0.019$)
slopev500V1500	1.45 (0.32)	1.57 (0.32)	1.28 (0.34)	0.35 (0.33)	0.12 (0.33)	4.21	0.003	.100	Happiness > Surprise ($P=0.005$) Happiness > Surprise ($P=0.043$) Surprise > Anger ($P=0.032$) Surprise > Anger ($P=0.002$) Fear > Happiness ($P=0.006$) Fear > Sadness ($P=0.010$) Fear > Surprise ($P=0.001$) Anger > Surprise ($P=0.042$) Happiness > Surprise ($P=0.019$)
F1Amplitude F2Amplitude	-0.3 (0.22) 0.32 (0.20)	-0.31 (0.22) 0.43 (0.20)	-0.19 (0.24) 0.21 (0.21)	-0.49 (0.23) 0.10 (0.20)	-0.85 (0.23) -0.56 (0.20)	1.30 3.68	0.274 0.007	.088	Anger > Surprise ($P=0.024$) Happiness > Surprise ($P=0.007$) Surprise < Anger ($P=0.030$) Happiness > Surprise ($P=0.008$) Happiness > Surprise ($P=0.012$)
F3Amplitude	0.34 (0.20)	0.46 (0.20)	0.24 (0.21)	0.14 (0.21)	-0.52 (0.21)	3.54	0.009	.085	Surprise < Anger ($P=0.030$) Happiness > Surprise ($P=0.008$) Happiness > Surprise ($P=0.012$)
H1H2 H1A3	1.44 (0.24) -0.91 (0.29)	1.61 (0.24) -1.19 (0.29)	0.66 (0.25) -1.61 (0.30)	0.48 (0.25) -1.40 (0.29)	1.13 (0.25) -0.83 (0.29)	4.00 1.23	0.004 0.301	.095	Happiness > Sadness ($P=0.012$)
Temporal-related:									
loudnesspeaksRate	-1.79 (0.27)	-1.35 (0.27)	-0.71 (0.28)	-1.30 (0.27)	-0.13 (0.27)	5.65	<0.001	.129	Surprise > Anger ($P<0.001$) Surprise > Happiness ($P=0.016$) Surprise > Sadness ($P=0.029$)
voicedLength unvoicedLength	0.28 (0.19) 0.15 (0.27)	0.31 (0.19) -0.05 (0.27)	0.17 (0.20) -0.17 (0.28)	0.35 (0.19) 0.42 (0.28)	-0.40 (0.19) 0.22 (0.28)	2.68 0.69	0.034 0.598	.066	
pseudovsyllableRate	-0.34 (0.19) -0.24 (0.19)	-0.22 (0.19) -0.26 (0.20)	-0.20 (0.19)	-0.38 (0.19)	0.44 (0.19)	3.00	0.020	.073	Surprise > Anger ($P=0.046$) Surprise > Sadness ($P=0.034$)

Note: F1Frequency = Frequency- formant 1, F2Frequency = Frequency-formant 2, F3Frequency = Frequency-formant 3, F1Bandwidth = Formant 1 bandwidth, HNR = Harmonics-to Noise ratio, AlphaRatio = Alpha ratio, Hammarberg index, v0v500 = Spectral Slope V 0-500 Hz, v500v1500 = Spectral slope V 500-1500 Hz, F1Amp = Formant 1 relative energy, F2Amp = Formant 2 relative energy, F3Amp = Formant 3 relative energy, H1H2 = Harmonic difference H1-H2, H1A3 = Harmonic difference H1-A3, LoudPeak = Rate of loudness peaks, Voice = Length of continuously voiced regions, Unvoiced = The length of unvoiced regions, Pseudo = Pseudo syllable rate.

Figure 3.4: Table comparing acoustic parameters between emotions (Ekberg et al., 2023)

Figure 3.5 shows a table of variations in different emotions measuring the acoustic parameters pitch, intensity, speaking rate and voice quality which are often used to identify emotions (Khalil et al., 2019). Although figure 3.4 displays a wider variety of more specific acoustic features, figure 3.5 provides a foundational understanding of different acoustic features connected to the different emotions.

Emotions	Pitch	Intensity	Speaking rate	Voice quality
Anger	abrupt on stress	much higher	marginally faster	breathy, chest
Disgust	wide, downward inflections	lower	very much faster	grumble, chest tone
Fear	wide, normal	lower	much faster	irregular voicing
Happiness	much wider, upward inflections	higher	faster/slower	breathy, blaring tone
Joy	high mean, wide range	higher	faster	breathy; blaring timbre
Sadness	slightly narrower	downward inflections	lower	resonant

Figure 3.5: Acoustic variations in different emotions (Khalil et al., 2019).

3.6 The Experiment

An experiment will be conducted and will consist of short interviews with voluntary people. These interviews will be recorded for data collection and will be used to extract emotions with a Python application.

3.6.1 Python Application

The Python application serves as the central system for processing and analyzing emotions in speech and will integrate several tools and frameworks to extract emotions.

The steps for the process consist of:

1. Audio recording
2. Feature extraction
 - Prosody and vocal bursts will be analyzed with Hume AI to detect emotional cues from pitch, intonation and vocal bursts using Hume AI's models.
 - Extracting vocal features such as jitter, shimmer and frequency with Praat Parselmouth.
 - Convert speech to text and analyze emotions in text with NLP Cloud.

3.6.2 Interviews and Surveys

The interviews involve voluntary participants engaging in short audio-recorded interviews, designed to draw out natural emotional responses. The participants will be asked questions to prompt them to recall and reflect on past experiences which encourages them to revisit emotions they felt at that time.

For ethical purposes the participants will be given a selection of topics to choose from, minimizing the risk of discomfort or distress. The audio recordings will be anonymous and recorded in a controlled acoustic environment to ensure minimal noise interference.

Questions asked during the interview follow one of many emotion induction techniques, “autobiographical recall”. This is a method used to facilitate the re-experiencing of emotions felt in a previous moment (Siedlecka & Denson, 2019), which is what is intended for the interviews to be able to collect emotional data from vocal recordings. By letting the participant think and speak about a memory from the past, emotions felt in that moment reflect in their voice.

After the interview is done, the participants will answer a survey doing a self-assessment of their emotions felt during the interview. This will enable a comparison between emotion detection and the participants reported emotional experiences.

There are many different methods for self-assessment, and emotional self-assessment is linked to many different theories. Many are connected to emotional intelligence (EI), trait emotional intelligence (trait EI) and Core Self Evaluation (CSE) (Montasem et al., 2013), but rather than conducting a comprehensive exploration of different psychological theories in self-assessment, this research will use simplified surveys at a basic level for the purpose of fulfilling the technical objectives of the work.

The theories behind the interviews are stated in the possibility of detecting emotions in voices. While there are a lot of recognized emotions that can be detected in different AI models and software tools, the interview will focus on bringing out two different basic emotions to maintain a manageable scope, while ensuring ethical feasibility.

Research has stated that there are different levels of unique universal signs for different affective states and while there are evidence supporting the universality for certain emotions such as anger, fear, surprise, sadness, happiness and more, there are also emotions that do not include all characteristics that distinguish them from other mental states, two examples being guilt and shame (Ekman & Cordaro, 2011). Research of this nature supports the rationale for having the focus solely on two of the basic emotions for this thesis. The questions in the interviews will focus on bringing out two separate emotions, one on the positive spectrum, happiness, and one on the negative spectrum, anger.

3.7 Statistical Analysis

Pearson Correlation Coefficient

Pearson's r is a measurement of the strength and direction of a linear relationship between two variables. The value range is from -1 to 1, where positive values implies a positive correlation and negative values the opposite. Values close to +1 indicate a strong correlation, values between +-0.30 and +-0.49 a moderate correlation, and values below +-0.29 are seen as a weak correlation. Values around 0 implies no linear correlation (Bruce & Bruce, 2017).

P-Value

A p-value indicates if the observed results have a probability of occurrence by chance. A widely accepted threshold for statistical significance is $p < 0.05$, which means there is less than a 5% possibility that the observed effect is random (Bruce & Bruce, 2017).

Z-score Standardization

Acoustic features such as pitch, intensity, jitter and shimmer can have great variation. To ensure comparability in statistical analyses, features are often standardized using Z-score standardization (Ekberg et al., 2023). This method transforms data to have a mean of zero and a standard deviation of one, ensuring meaningful comparisons across features. By this, a variable does not have an overly influence due to a scale of the measurement. The measurements are described as "standard deviations away from the mean". (Bruce & Bruce, 2017).

Standardized Distance for Emotion Categorization

Emotion categorization based on vocal features can be operated through standardized distance methods, where deviations from the baseline of acoustic profiles are quantified. Using standardized differences allow an interpretable measure of how vocal features aligns with expected patterns for each emotion (Ekberg et al., 2023) (Bruce & Bruce, 2017). The categorization method used in this study is a custom method inspired by this standard practice.

ANOVA Tests

ANOVA (Analysis of Variance) is a standard statistical method used to determine if there are any significant differences in means across multiple groups (Bruce, 2017). It is used to categorize grouping factors and are one method in the Swedish research for vocal markers (Ekberg, 2023).

Tukey's HSD

When ANOVA presents significant differences between group means, Tukey's HSD test is incorporated as a post analysis to identify which groups are divergent from each other. This method controls for errors when making multiple comparisons (Bruce & Bruce, 2017).

T-Tests and Cohen's d

Paired T-tests are used to compare the means between two groups to determine statistical significance, while Cohen's d provides a standardized measure of the effect size which indicates the magnitude of the observed differences (Cohen, 1977) (Bruce & Bruce, 2017).

Results

4.1 Presentation of Collected Data

4.1.1 Overview of Interviews

We conducted semi-structured interviews with 16 native Swedish speakers (10 M/6F, age 23-78), each lasting 1-3 minutes. Each participant was interviewed for two different scenarios, resulting in 32 different recordings. All interviews were audio-recorded in a quiet room and elicited two target emotions – anger and happiness – via open-ended prompts (e.g. “Is there anything in society that makes you upset? What? How does that make you feel?”; “Can you remember one time you felt really proud of yourself?”). The participants rated their perceived emotions on a 1-6 scale immediately after each scenario. The rated emotions covered the basic 5 emotions mentioned in this report: anger, joy, sadness, fear, and surprise.

Table 4.1 presents the participants ID, gender, age, and self-assessed scores for their perceived emotions for respective interview scenario.

Participant			Negative					Positive				
ID	M/F	Age	A	J	Sad	F	Sur	A	J	Sad	F	Sur
1	M	23	5	1	3	1	1	1	6	1	1	4
2	M	26	6	1	3	4	1	1	6	1	2	1
3	F	27	4	1	6	1	2	1	6	1	1	3
4	M	29	2	1	3	2	1	1	4	2	2	2
5	F	28	4	1	4	1	2	1	5	1	1	5
6	M	25	2	2	1	1	1	1	3	1	1	1
8	M	27	3	1	2	1	2	1	5	1	1	1
9	F	26	3	1	3	1	1	1	5	1	1	1
10	F	78	5	1	3	2	4	1	6	4	1	1
11	F	27	3	3	2	1	1	1	6	1	1	1
12	M	58	1	3	1	2	1	1	6	1	1	3
13	F	54	4	1	4	3	1	1	6	1	1	1
14	M	20	1	3	1	2	2	1	4	1	1	3
15	M	30	3	2	2	3	1	2	5	1	1	1
16	M	25	4	1	2	1	1	1	6	1	1	1

Table 4.1: Participant table. A: Anger, J: Joy, Sad: Sadness, F: Fear, Sur: Surprise.

4.1.2 Data Collection for RQ1: Vocal Features & Speech

The collected audio recordings from the 32 interviews were processed for research questions 1 to specifically focus on vocal features and speech-based emotion recognition.

Vocal Feature Extraction (Praat Parselmouth)

Audio recordings were processed with Praat Parselmouth. Vocal parameters were extracted from each recording, which have been validated by Swedish emotion research on vocal markers (Ekberg et al., 2023).

- Pitch: mean pitch in Hz and semitones (ST).
- Intensity: mean intensity measured in decibels (dB).
- Voice Quality Metrics: Harmonic-to-Noise Ratio (HNR), jitter (local frequency perturbation), shimmer (local amplitude perturbation).

These acoustic features were then categorized into discrete emotional labels (anger, joy, sadness, fear, surprise) based on thresholds and criteria defined in prior Swedish research (Ekberg et al., 2023).

Speech-Based Emotion Recognition (Hume AI)

The same audio were analyzed with Hume AI, that provided probability distributions across several emotions, where the five targeted emotions was filtered from. The results were normalized according to ADD THAT PART INTO METHOD FROM PDF REFER TO HERE.

Data Structure

The extracted vocal features, custom emotion categorizations of vocal features, and Hume AI outputs were combined for each recording and stored in JSON format, to enable direct comparison and further analysis. All data is normalized to sum up to 1 before loaded into the JSON files. Table 4.2 includes the mean values and standard deviation for the analysed features (pitch, intensity, HNR, jitter, shimmer) separated by sentiment. Filtered Hume probabilities are presented in Table 4.3 as mean values with standardized deviation for positive and negative recordings.

The custom categorization of emotion based on vocal features are presented as mean values including standard deviations for positive and negative sentiments in Table 4.4.

Negative Clips			Positive Clips		
Feature	Mean	Std	Feature	Mean	Std
mean_pitch_st	-0,1606	5,9616	mean_pitch_st	-0,6688	6,0049
mean_pitch_hz	157,1044	53,2213	mean_pitch_hz	152,8906	54,6353
mean_intensity_db	62,9767	2,3435	mean_intensity_db	63,7063	3,0511
mean_hnr_db	5,7883	5,8332	mean_hnr_db	5,3269	5,8034
jitter_local	0,0246	0,0032	jitter_local	0,0253	0,0046
shimmer_local	0,1219	0,0267	shimmer_local	0,1286	0,0218

Table 4.2: Summary Statistics: Vocal Features by Sentiment

Positive Clips			Negative Clips		
Metric	Mean	Std	Metric	Mean	Std
hume_anger	0,2262	0,0584	hume_anger	0,2832	0,0751
hume_fear	0,1485	0,0513	hume_fear	0,1523	0,0372
hume_joy	0,3318	0,1279	hume_joy	0,2758	0,0951
hume_sadness	0,1673	0,0644	hume_sadness	0,1771	0,0691
hume_surprise	0,1264	0,0209	hume_surprise	0,1117	0,0205

Table 4.3: Summary Statistics: Hume AI Probabilities by Sentiment

Positive Clips			Negative Clips		
Metric	Mean	Std	Metric	Mean	Std
custom_anger	0,2066	0,0031	custom_anger	0,2064	0,0034
custom_joy	0,2036	0,0039	custom_joy	0,2043	0,0041
custom_fear	0,2028	0,0079	custom_fear	0,2042	0,0087
custom_sadness	0,1913	0,0027	custom_sadness	0,1906	0,0032
custom_surprise	0,1956	0,0062	custom_surprise	0,1946	0,0063

Table 4.4: Summary Statistics: Custom Emotion Categorization Scores by Sentiment

Segment-Level Data

Certain analyses in RQ1 rely on time-segmented data. For each recording, Hume AI returns emotion probabilities at regular time segments, an example of this is presented in Table 4.5. In the data analysis, we extract the same set of acoustic features from matching time intervals with Praat, enabling time-to-time comparisons across modalities.

time (s)	anger	fear	joy	sadness	surprise
1,47	0,2332	0,1590	0,4244	0,1214	0,0620
5,15	0,1469	0,0342	0,6693	0,0110	0,1387
8,27	0,0993	0,0259	0,7804	0,0184	0,0759
...
43,2342	0,1216	0,0837	0,5861	0,0500	0,1586

Table 4.5: Segment-Level Hume Probabilities for clip: id_001_neg

4.1.3 Data Collection for RQ2 and RQ3: Text, Speech and Self-Assessment

The data collection for RQ2 and RQ3 is based on the same audio recordings as for RQ1. Each recording was transcribed and analyzed with NLP Cloud (text-based), to extract emotion probabilities from the transcription. The same audio was analyzed using Hume AI for speech-based emotion detection, resulting in paired emotion probability scores alongside self-reported emotion ratings. All scores were normalized for comparison.

The data was structured in JSON format, each audio object consists of five emotion labels from each data type (Hume, NLP, Self).

Table 4.6 summarize the average emotion scores and standard deviations for both speech-based (Hume AI) and text-based (NLP Cloud) models across all clips in the dataset.

All Recordings

Emotion	Self Mean	Hume Mean	NLP Mean	Self Std	Hume Std	NLP Std
Anger	0,210	0,260	0,200	0,124	0,072	0,223
Joy	0,312	0,302	0,396	0,200	0,117	0,351
Sadness	0,190	0,167	0,181	0,105	0,065	0,138
Fear	0,136	0,150	0,093	0,061	0,045	0,092
Surprise	0,149	0,118	0,129	0,082	0,022	0,089

Table 4.6: Means and standard deviations of self-reported, Hume, and NLP emotion intensities.

The interviews were conducted with either a positive or negative orientation. Each recording was analyzed individually, and the data structure distinguishes between negative and positive audio files. The corresponding emotion scores from Hume AI and NLP Cloud are presented in Table 4.7 for positively oriented interviews, and in Table 4.8 for negatively oriented interviews. Each table displays the mean and the standard deviation for the respective AI model’s emotion probability.

Positive Recordings

Emotion	Self Mean	Hume Mean	NLP Mean	Self Std	Hume Std	NLP Std
Anger	0,103	0,227	0,015	0,032	0,060	0,057
Joy	0,497	0,333	0,708	0,081	0,132	0,169
Sadness	0,117	0,163	0,067	0,059	0,065	0,062
Fear	0,108	0,148	0,040	0,033	0,052	0,067
Surprise	0,173	0,126	0,171	0,098	0,022	0,096

Table 4.7: Means and standard deviations of self-reported, Hume, and NLP emotion intensities for positive recordings.

Negative Recordings

Emotion	Self Mean	Hume Mean	NLP Mean	Self Std	Hume Std	NLP Std
Anger	0,305	0,289	0,363	0,092	0,072	0,183
Joy	0,148	0,275	0,121	0,106	0,098	0,205
Sadness	0,254	0,171	0,282	0,095	0,066	0,103
Fear	0,161	0,152	0,141	0,069	0,038	0,087
Surprise	0,128	0,112	0,092	0,059	0,021	0,064

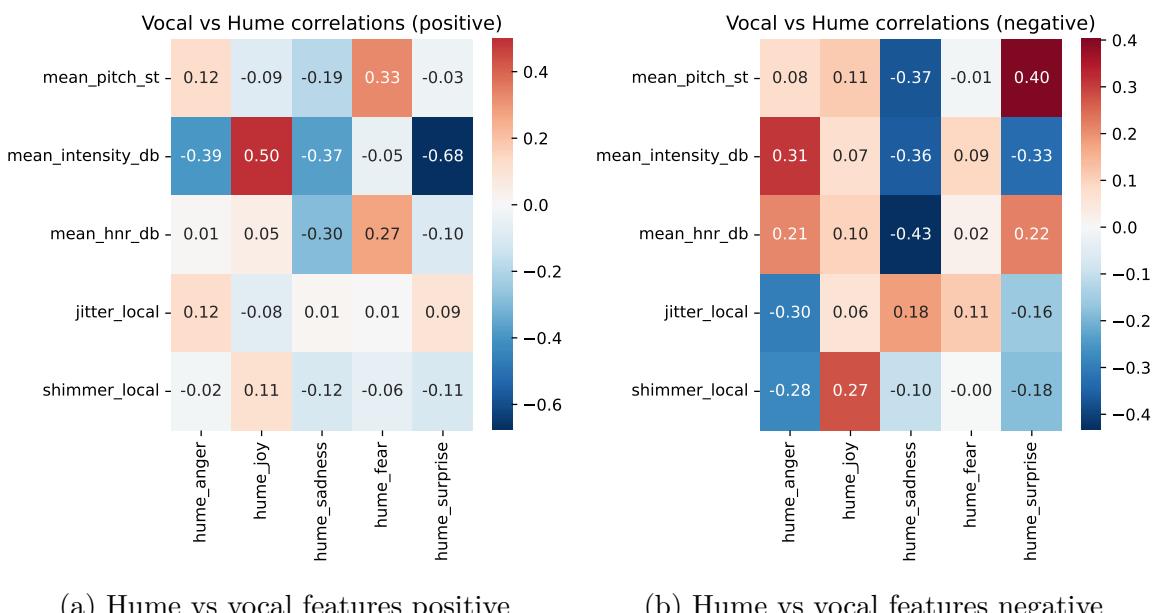
Table 4.8: Means and standard deviations of self-reported, Hume, and NLP emotion intensities for negative recordings.

4.2 Data Analysis for RQ1: Vocal Features & Speech Emotion Recognition

The first research question explores how vocal features correlates with AI-based emotion detection in conversational Swedish speech. To analyse this, acoustic features such as pitch, intensity, jitter, shimmer and HNR were extracted using Praat Parselmouth and compared with emotion scores from the speech-based model Hume AI. A custom categorization method based on Swedish vocal emotion research (Ekberg et al., 2023) were tested for comparison. The goal with this analysis was to explore if these vocal markers could explain or predict how speech-based AI systems interpret emotional expressions in semi-structured, spontaneous speech in an interview setting.

4.2.1 Correlation Between Vocal Features and AI Emotion Scores (Hume AI)

Figure 4.1 demonstrates heatmaps of the Pearson correlation coefficients between selected vocal features and Hume AI emotion labels across positive clips in Figure 4.1a and negative clips in Figure 4.1b. The results show generally weak correlations, with slightly stronger correlations for the negative recordings.



(a) Hume vs vocal features positive

(b) Hume vs vocal features negative

Figure 4.1: Correlation Heatmaps between Hume AI and vocal features.

Correlation values for positive interviews in Figure 4.1a ranging roughly between -0.7 and 0.5, most values suggest generally weaker correlations than the highest values. Stronger positive correlations suggests that certain feature is higher when Hume rates the correlating emotion. Negative correlations implies the opposite, low value of a certain feature for the correlated emotion. Mean intensity stands out from other vocal features with a moderate positive correlation with Hume Joy ($r = 0.50$), a moderate negative correlation with Hume Anger ($r = -0.39$), a moderate negative relationship with Hume sadness ($r = -0.37$) and a strong negative correlation with Hume Surprise ($r = -0.68$). Mean pitch shows strongest effect on Hume Fear ($r = 0.30$) and a moderate negative link with Hume Sadness ($r = -0.21$). Mean HNR has a moderate negative correlation with Hume Sadness ($r = -0.30$) as well it is moderately positively correlated with Hume Fear ($r = 0.27$). Jitter and shimmer remain near zero for most emotions

in the positive clips, with none exceeding $|r| = 0.13$. This suggests that, in more positively oriented interviews, variation in pitch, intensity, and HNR capture the core emotion-related cues moderately, while jitter and shimmer have small predictive influence in semi-spontaneous speech during interview conversations.

Figure 4.1b illustrates Pearson correlation values for negative clips in the dataset, again presenting generally weak effects even if slightly stronger correlations occur, approximately ranging between -0.5 and 0.5. The strongest relationship appears for sadness, where mean HNR shows a strong negative correlation ($r = -0.43$) and moderately negative correlated with pitch ($r = -0.37$) and intensity ($r = -0.36$). Anger had the highest positive correlation with intensity ($r = 0.31$) followed by a weak positive link with HNR ($r = 0.21$) and weak to moderate negative correlations with jitter ($r = -0.30$) and shimmer ($r = -0.28$). Correlations between vocal features and joy are all perceived weak, shimmer emerges as strongest ($r = 0.27$) compared to all other features ($r < 0.11$). Hume predicted fear presents coefficients indicating no linear correlation for all vocal features ($r < 0.11$). Surprise presents a moderate positive correlation with pitch ($r = 0.36$) and negative relationship with intensity ($r = -0.33$), other features are weakly correlated to the emotion. Jitter and shimmer show similar relationships to Hume Fear ($r = 0.2$). Shimmer has strongest correlation with Hume Anger, and a moderate relationship with Hume Joy. Jitter is moderately correlated with Anger as well and have a moderate correlation with Sadness where shimmer has almost no correlation.

Overall, the negative and positive diversion suggests that intensity consistently reflect Hume AI's anger, sadness, and surprise predictions, and more prominent for joy in positive recordings and near zero for joy rated in negative contexts. Correlations between pitch and emotions predicted by Hume is generally weak and varying between the sentiment categories. Jitter and shimmer remain minor indicators for positive conditions, while having moderate correlations with Hume anger and joy in negative recordings.

4.2.2 Correlation with Praat-Based Emotion Scores

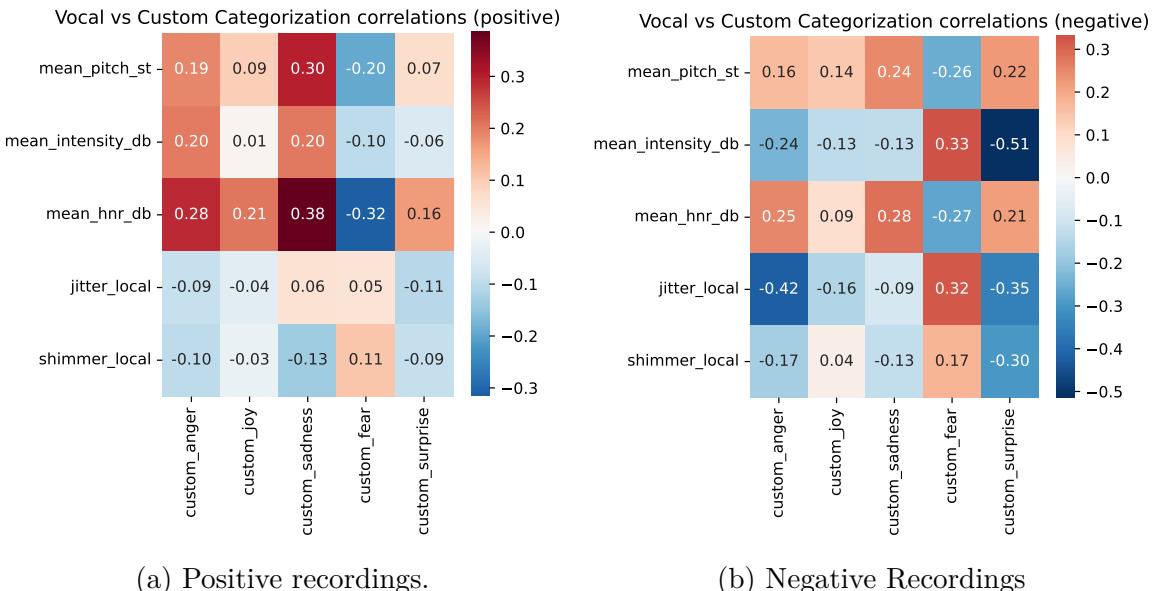


Figure 4.2: Correlation heatmaps for custom categorized emotion scores vs. vocal features.

Figure 4.2 presents heatmaps of Pearson's correlation r between selected vocal features and the emotion scores obtained from the custom emotion categorization function for Figure 4.2a

positive oriented clips and Figure 4.2b negative clips. In contrast to the correlation results between Hume AI and vocal markers, these heatmaps present similar values and orientations for both negative and positive recordings, with values roughly ranging from -0.5 to 0.4. In positive recordings, pitch and HNR correlation coefficients show a similar pattern where custom categorized sadness have the most distinct relationship (pitch $r = 0.30$, HNR $r = 0.38$). Pitch is weakly correlated with other custom categorized emotions, with a slightly heightened positive correlation with anger ($r = 0.19$) and negative with fear ($r = -0.20$). HNR demonstrates the strongest correlations of all features, approaching a positive moderate relationship with anger ($r = 0.28$), while negatively correlated with fear ($r = -0.32$). Intensity shows a slight increase in correlation, yet below the threshold for moderate strength, for anger and sadness ($r = 0.20$). Jitter and shimmer have minor correlations with all custom categorized emotion labels ($|r| < 0.13$).

Correlations for the negative subset have lower value distribution than the positive subset. Intensity and custom surprise present the singular strong negative correlation ($r = -0.51$). Subsequently, fear and intensity are moderately correlated ($r = 0.33$). As for positive recordings, pitch and HNR demonstrate similar behaviors for all emotions - both have a similar weak positive correlation with anger ($r = 0.16$, $r = 0.25$ respectively), minor correlation with joy ($r < 0.14$), slightly stronger positive relationship with sadness and surprise ($r = 0.21-0.28$), and a negative link with fear ($r = -0.26$). Jitter has the strongest association with anger ($r = -0.42$), and a moderate positive link with fear ($r = 0.32$). Surprise is the single emotion presenting a moderate correlation with shimmer ($r = -0.30$), other relationships remain weak.

In general, the heatmaps reveal differences in correlation values and directions across positive and negative clips, while some consistent patterns occur. HNR correlates most strongly with sadness in both subsets ($r = 0.38$, $r = 0.28$), but weaker in negative contexts. Pitch has strongest relationship with sadness in the positive set ($r = 0.30$), but is more evenly distributed in the negative set (highest for surprise, $r = 0.22$). Jitter and shimmer are almost negligible in positive contexts correlations, but in negatives jitter is a stronger indicator for three of five emotions, and shimmer for surprise individually. Overall, HNR and pitch remain the most stable vocal markers, while intensity, jitter and shimmer has substantial variability dependent on the sentiment context.

4.2.3 Correlation Custom Categorization and Hume AI Labels

To explore the alignment of top-labelled emotion by Hume AI and customised categorization based on vocal features, Figure 4.3 illustrates the correlations in a confusion matrix, treating Hume labels as ground truth and comparing these to emotions grouped by the custom function. For positive recordings, Hume and the custom function agree on anger for one clip, while Hume rates joy highest in four cases where the custom function selects anger. Both methods rate joy as the top emotion for four clips, although it is one case where the custom function picks joy and Hume anger. Sadness is the only other emotion with agreement, otherwise divergences occur in fear-anger, joy-surprise, and anger-surprise pairs. In negative contexts, the methods agree on anger as the top emotion for six recordings and agree on joy for one clip. Beside these, the top-rated emotion is discrepant where Hume labels two clips as joy that the custom function rates as anger, and two joy-rated clips by Hume are rated as surprise by the customized categorization. Other divergences occur for sadness-anger, anger-joy, anger-sadness, and anger-surprise in single-pairs.

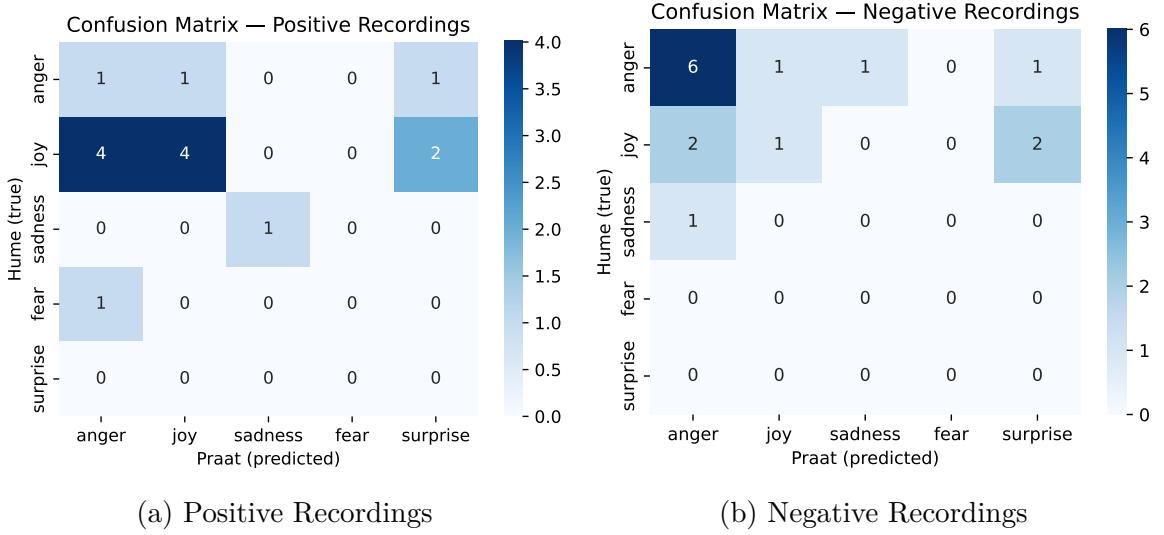


Figure 4.3: Confusion matrix diagram for top label emotion between Hume and Custom Categorization.

The mean rated scores for Hume versus custom categorized emotion labels for positive recordings are illustrated in Figure 4.4, negative in Figure 4.5. As shown, some ratings follow the same pattern while other emotions diverge, as Hume joy are significantly higher for both sentiment categories while custom labelled surprise is overestimated compared to the Hume ranking.

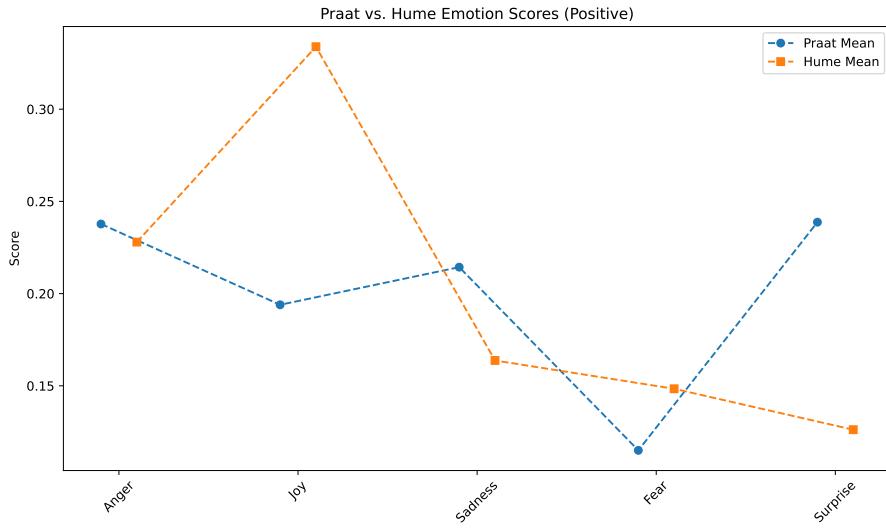


Figure 4.4: Scatter Plot of Mean Hume and Custom Emotion Scores, positive recordings.

Custom labelled joy, sadness, and fear are ranked similarly in both positive and negative contexts. This pattern aligns with Hume probabilities that are more equal regarding surprise than the custom function for separated sentiments. Anger is the most divergent emotion label across sentiments, with higher ratings in negative clips. Both sources rates anger, sadness, and fear relatively corresponding, Hume predicts anger slightly higher in negative contexts, while sadness is ranked lower than by the custom function for both sentiments.

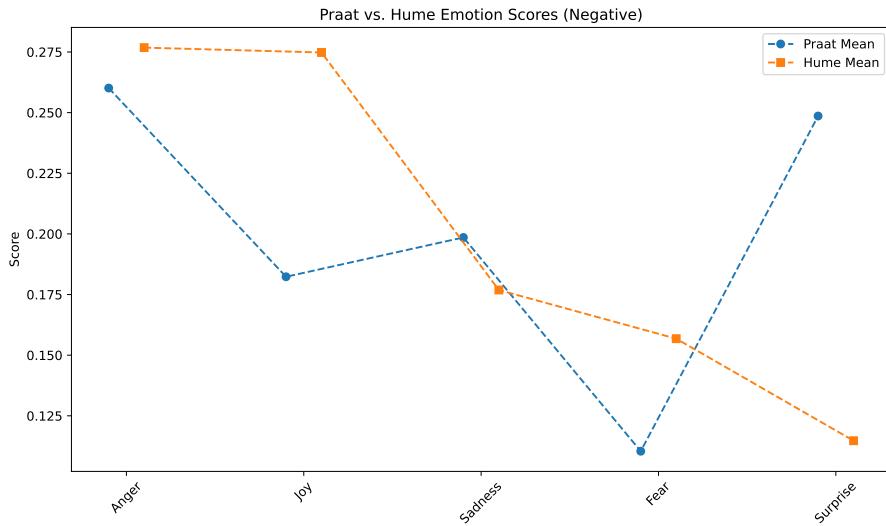


Figure 4.5: Scatter Plot of Mean Hume and Custom Emotion Scores, negative recordings.

Pearson's correlation coefficients between Hume and custom categorized emotion labels are presented in Table 4.9 separated by sentiment. Sadness in positive recordings is the single significant correlation ($r = 0.614$, $p = 0.0149$), other relationships remain weak. Correlations in the negative subset are slightly higher regarding the negative association with anger ($r = -0.295$, $p = 0.2854$) and a positive link with fear ($r = 0.269$, $p = 0.3329$), yet both of these correlations are weak without significance. All remaining labels present near negligible correlation coefficients ($r < 0.09$).

Positive Recordings				Negative Recordings			
Emotion	Pearson's r	p-value	Sign.	Emotion	Pearson's r	p-value	Sign.
Anger	-0.091	0.7474	No	Anger	-0.295	0.2854	No
Joy	0.049	0.8613	No	Joy	-0.029	0.9183	No
Sadness	0.614	0.0149	Yes	Sadness	-0.060	0.8330	No
Fear	-0.044	0.8772	No	Fear	0.269	0.3329	No
Surprise	0.022	0.9375	No	Surprise	-0.081	0.7741	No

Table 4.9: Pearson's r, p-values, and significance for positive vs. negative recordings.

4.2.4 Limitations of the Custom Vocal Emotion Categorization Method

The rule-based function, customised to group emotions by vocal features as presented in the Swedish study (Ekberg et al., 2023), often assigns equal top scores to multiple emotions for the same clip. The number of top rated emotion for separate recordings are summarised in Table 4.10, showing that 60% of the total clips have more than one top label for the custom function. By contrast, Hume has one singular top label for 29 of 30 clips.

(a) Custom-categorisation			(b) Hume AI		
#Top Emotions	Clips	Percentage	# Top Emotions	Clips	Percentage
1	12	40%	1	29	97%
2	9	30%	2	1	3%
3	5	17%			
4	4	13%			

Table 4.10: Distribution of recordings by number of top emotions

To further demonstrate the occurring similarities in the emotion scoring by the rule-based function, bar charts for single clips is presented in Figure 4.6 for two separate positive clips and Figure 4.7 illustrates two single negative interviews. Concurrence occurs for clip 4.6b, for both Hume and the custom function except for joy that is predicted at high levels by Hume. Diagram 4.6a and 4.7a reveals reasonable approximation for certain emotions, as joy and fear in 4.6a where Hume rates fear substantially higher. Similar pattern appears in 4.7a for joy and sadness, in comparison to high predicted Hume-joy. However, discrepancies are presented in all single clips with varying dispersed.

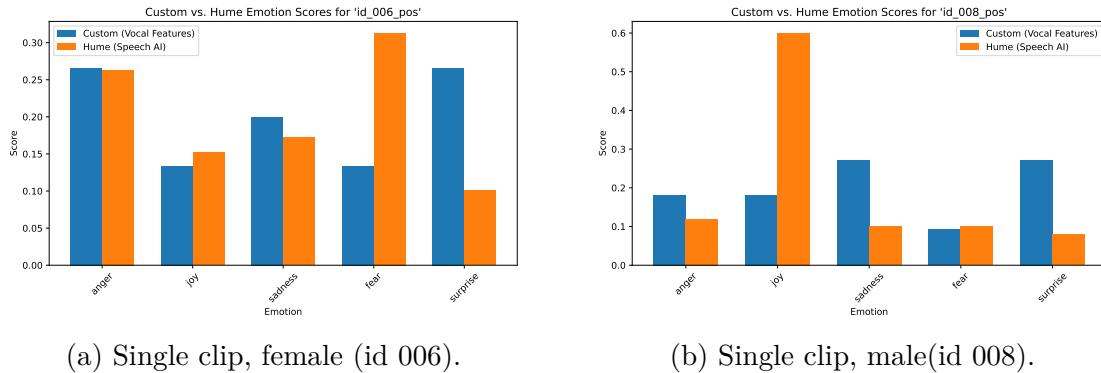


Figure 4.6: Custom emotion categorization vs. Hume for two separate positive clips.

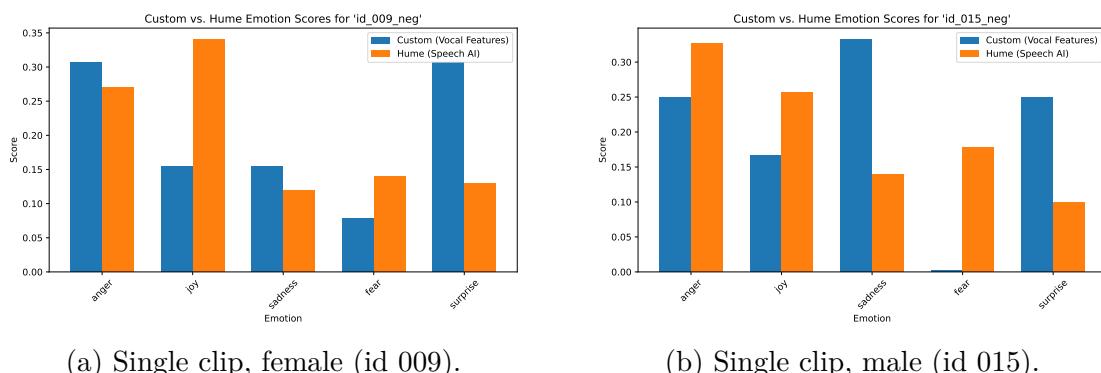


Figure 4.7: Custom emotion categorization vs. Hume for two separate negative clips.

4.2.5 ANOVA Tables of Vocal Features Across Emotions

An ANOVA was implemented to further examine whether essential vocal features varied across Hume-labelled emotions. This was conducted on mean values of full recordings for pitch, intensity, HNR, jitter, and shimmer. The results are summarised in Table 4.11a for positive

recordings and 4.11b for negative recordings, showing that none of the features showed statistically significant differences between the five Hume emotion categories (all p-values > 0.13 , ranging up to > 0.90).

Positive recordings				Negative recordings			
Feature	ANOVA	P-value	Sign.	Feature	ANOVA	P-value	Sign.
Pitch		0.7595	No	Pitch		0.5393	No
Intensity		0.8627	No	Intensity		0.1307	No
HNR		0.6149	No	HNR		0.5142	No
Jitter		0.9564	No	Jitter		0.9066	No
Shimmer		0.7828	No	Shimmer		0.6863	No

(a) ANOVA: Positive recordings.

(b) ANOVA: Negative recordings.

Table 4.11: ANOVA for vocal features across emotions.

These results imply that within our dataset of spontaneous speech during interviews, the average values of the acoustic features did not systematically vary according to AI-labeled emotions.

4.2.6 Time-to-Time Analysis

Time-to-Time Analysis: Full Dataset

To understand both if acoustic cues correlates with emotion probability by Hume and when they produce clear categorical shifts, two analyses were conducted at time segmened level. These effects were observed further with sentiment seperations of all, negative, and positive contexts to see if the emotion-acoustic feature relationship are dependent on positive vs. negative sentiment. Table 4.12 presents the significant correlations between z-scored acoustic features and Hume emotions.

Feature	Emotion	Pearson's r	p-value	Significant
pitch	joy	0.065	0.0448	Yes
pitch	sadness	-0.230	0.0000	Yes
pitch	fear	0.082	0.0110	Yes
intensity	joy	0.164	0.0000	Yes
intensity	sadness	-0.142	0.0000	Yes
intensity	fear	-0.173	0.0000	Yes
intensity	surprise	-0.110	0.0006	Yes
HNR	sadness	-0.253	0.0000	Yes
HNR	fear	0.112	0.0005	Yes
jitter	sadness	0.084	0.0091	Yes

Table 4.12: Significant Pearson correlations for the full dataset.

Only significant correlation is included (all $p < 0.05$). 9 correlations out of 24 indicated significance, the full analysis is presented in Appendix [FIGURE REF]. All correlations are perceived as weak even if there is statistical significance. Pitch correlated positively with joy ($r = 0.065$), negatively with sadness ($r = -0.230$) and positive with fear ($r = 0.082$). Intensity had a increased relationship with joy ($r = 0.164$), but decreased with sadness ($r = -0.142$), fear ($r =$

-0.173) and surprise ($r = -0.110$). HNR showed a weak negative correlation with sadness ($r = -0.253$) and positive with fear ($r = 0.112$). Jitter had a single significant, yet weak correlation with sadness ($r = 0.084$). Shimmer showed no significant correlations.

Table 4.13 compares the top 30% and bottom 70% of emotion-probability time-segments, and tests whether the mean z-scored feature differs between the high vs. low groups. A large t-statistic value indicate a reliable shift in that feature when Hume rates that emotion high.

Feature	Emotion	t-statistic	p-value	Significant
pitch X	anger	2.529	0.0116	Yes
pitch	joy	2.293	0.0221	Yes
pitch	sadness	-7.769	0.0000	Yes
pitch	fear	5.185	0.0000	Yes
intensity X	anger	1.975	0.0485	Yes
intensity	joy	4.602	0.0000	Yes
intensity	sadness	-3.981	0.0001	Yes
intensity	fear	-3.567	0.0004	Yes
intensity	surprise	-2.949	0.0033	Yes
HNR	anger	2.709	0.0069	Yes
HNR	sadness	-7.506	0.0000	Yes
HNR	fear	4.914	0.0000	Yes
HNR X	surprise	2.287	0.0224	Yes
jitter	sadness	1.811	0.0705	No

Table 4.13: High-vs-low t-test results for significant acoustic features (full dataset)

High-anger predictions of segmented recordings had higher pitch ($t = 2.529$), intensity ($t = 1.975$) and HNR ($t = 2.709$), none of these associations was significant in the data for Table 4.12. Joy segments had increased pitch ($t = 2.293$) and intensity ($t = 4.602$), while fear segments have lower intensity ($t = -3.567$) and increased pitch ($t = 5.185$) and HNR ($t = 4.914$). High-sadness segments showed notable decreased pitch ($t = -7.769$), intensity ($t = -3.981$), HNR ($t = -7.506$), and a small, not significant increase in jitter ($t = 1.811$, $p = 0.0705$). Segments with high surprise predictions has the moderately low intensity ($t = -2.949$) and a modest positive differentiation for HNR ($t = 2.287$).

Time-to-Time Analysis by Sentiment

(a) Pitch and Anger (Pearson r)				(b) Pitch and Anger (t-test)			
Sentiment	r	p	Sign.	Sentiment	t	p	Sign.
All	0.032	0.3284	No	All	2.529	0.0116	Yes
Positive	-0.042	0.3772	No	Positive	1.392	0.1647	No
Negative	0.101	0.0209	Yes	Negative	2.642	0.0085	Yes

Table 4.14: (a) Pearson correlations and (b) high-vs-low t-test results for pitch vs. anger by sentiment.

Table 4.14 presents an absent correlation between pitch and anger for the full dataset ($r = 0.032$, $p = 0.3284$), and positive clips ($r = -0.042$, $p = 0.3772$) but a significant, yet weak

correlation in the negative subset ($r = 0.101$, $p = 0.0209$). The t-tests in 4.14b confirms this where high-anger segments have moderate higher pitch in the negative set ($t = 2.642$, $p = 0.0085$) and in for all clips ($t = 2.529$, $p = 0.0116$), but not in positive contexts. This implies that pitch is a considerable signal of anger when the overall context is negative.

(a) HNR and Anger (r)				(b) HNR and Anger (t-test)			
Sentiment	<i>r</i>	<i>p</i>	Sign.	Sentiment	<i>t</i>	<i>p</i>	Sign.
All	0.059	0.0695	No	All	2.709	0.0069	Yes
Positive	-0.001	0.9750	No	Positive	1.941	0.0529	No
Negative	0.111	0.0109	Yes	Negative	3.123	0.0019	Yes

Table 4.15: HNR–Anger correlations and t-test results by sentiment.

Table 4.15 presents harmonic-to-noise correlation with anger for the full dataset ($r = 0.059$, $p = 0.695$), implying an unsignificant weak correlation. T-tests in 4.15b shows that HNR slightly distinguish high vs low anger for both the full and negative dataset ($t = 2.709$, $t = 3.123$ respectively). These correlations is significantly weak, in negative contexts ($r = 0.111$, $p = 0.0109$) but non-significant in positive recordings ($r = -0.001$, $p = 0.9750$), which also has a non-significant distinction in the t-tests. Together with pitch, HNR appears to be the most prevalent features correlated with anger in negative circumstances, even if the shifts in vocal features for this emotion is fairly weak.

(a) Pitch and Joy (r)				(b) Pitch and Joy (t-test)			
Sentiment	<i>r</i>	<i>p</i>	Sign.	Sentiment	<i>t</i>	<i>p</i>	Sign.
All	0.065	0.0448	Yes	All	2.293	0.0221	Yes
Positive	0.110	0.0210	Yes	Positive	1.816	0.0701	No
Negative	0.014	0.7567	No	Negative	0.617	0.5374	No

Table 4.16: Pitch–Joy correlations and t-test results by sentiment.

Pitch and joy correlations is presented in Table 4.16, showing unlike anger, a weak correlation for the full dataset ($r = 0.065$, $p = 0.0448$). Pitch reaches a weak correlation with joy in positive recordings ($r = 0.110$, $p = 0.0210$) and no significance in the negative subset. The t-tests is only significant in the full dataset ($t = 2.293$, $p = 0.0221$), not for the positive clips ($t = 1.861$, $p = 0.0701$), implying a more context-specific and non-linear affect.

(a) Intensity and Joy (r)				(b) Intensity and Joy (t-test)			
Sentiment	<i>r</i>	<i>p</i>	Sign.	Sentiment	<i>t</i>	<i>p</i>	Sign.
All	0.164	0.0000	Yes	All	4.602	0.0000	Yes
Positive	0.175	0.0002	Yes	Positive	3.343	0.0009	Yes
Negative	0.152	0.0005	Yes	Negative	2.375	0.0179	Yes

Table 4.17: Intensity–Joy correlations and t-test results by sentiment.

Table 4.17 demonstrates the correlation between intensity and joy, showing stronger associations compared to joy and pitch. All contexts show a positive relationship, strongest for positive clips ($r = 0.175$, $p = 0.0002$) with significant mean differences where the full dataset

has highest significance ($t = 4.602$, $p < 0.001$). These results suggests that intensity is a reasonably stable cue to joy regardless of the overall sentiment context, even if higher correlation occurs for the full and positive sets than the negative subset ($r = 0.152$ and $t = 2.375$).

(a) Pitch and Sadness (r)				(b) Pitch and Sadness (t-test)			
Sentiment	<i>r</i>	<i>p</i>	Sign.	Sentiment	<i>t</i>	<i>p</i>	Sign.
All	-0.230	0.0000	Yes	All	-7.769	0.0000	Yes
Positive	-0.275	0.0000	Yes	Positive	-6.332	0.0000	Yes
Negative	-0.187	0.0000	Yes	Negative	-4.980	0.0000	Yes

Table 4.18: Pitch–Sadness correlations and t-test results by sentiment.

Table 4.18 display seperated sentiment correlations for pitch and Hume predicted sadness, with significant difference between high and low groups indicating prominent shifts in pitch when Hume rates sadness high. Correlations are weak, but significant for all sentiment contexts. Positive recordings show the largest negative correlation ($r = -0.275$, $p < 0.001$), although the full dataset reveals greater diverge in variations ($t = -7.769$, $p < 0.001$). The negative subset accommodate lowest correlation and shifts in pitch for sadness.

(a) HNR and Sadness (r)				(b) HNR and Sadness (t-test)			
Sentiment	<i>r</i>	<i>p</i>	Sign.	Sentiment	<i>t</i>	<i>p</i>	Sign.
All	-0.253	0.0000	Yes	All	-7.506	0.0000	Yes
Positive	-0.243	0.0000	Yes	Positive	-5.442	0.0000	Yes
Negative	-0.262	0.0000	Yes	Negative	-5.049	0.0000	Yes

Table 4.19: HNR-Sadness correlations and t-test results by sentiment.

HNR showed to be an indicator of sadness in Table 4.13 and 4.12, presented further in Table 4.19. The patterns between sadness and HNR is similar to pitch, such that significant correlation occur across all sentiments with large differentions in HNR fluctuation. Negative conditions show strongest negative correlation ($r = -0.262$, $p < 0.001$), followed by the full dataset ($r = -0.253$, $p < 0.001$) that exhibit the largest variations of vocal variance regarding sadness ($t = -7.506$, $p < 0.001$).

These large t-statistics suggests that reduced pitch and HNR are indicators of sadness independent from sentiment orientation.

Case Examples

For a more concrete illustration of the prior tendencies , three interview recordings were analyzed in detail. The purpose was to examine whether emotional shifts become more apparent when evaluating shorter time segments within individual speakers, compared to the weaker correlations observed at the dataset level.

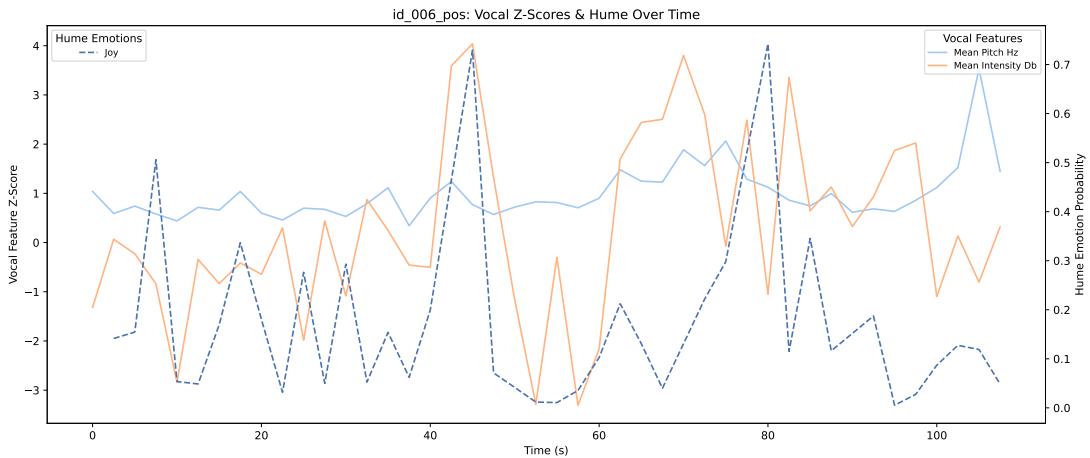


Figure 4.8: Pitch and Intensity vs. Hume Joy for single positive clip.

(a) Pearson Correlation (Clip id_006_pos)

Feature	Emotion	<i>r</i>	<i>p</i>	Sign
pitch	joy	0.188	0.1965	No
intensity	joy	0.351	0.0134	Yes
intensity	fear	-0.413	0.0032	Yes
jitter	fear	-0.291	0.0426	Yes

(b) t-test (Clip id_006_pos)

Feature	Emotion	<i>t</i>	<i>p</i>	Sign
pitch	joy	0.331	0.7419	No
intensity	joy	2.718	0.0092	Yes

Table 4.20: Statistical results for clip id_006_pos

24 rows for all!!!

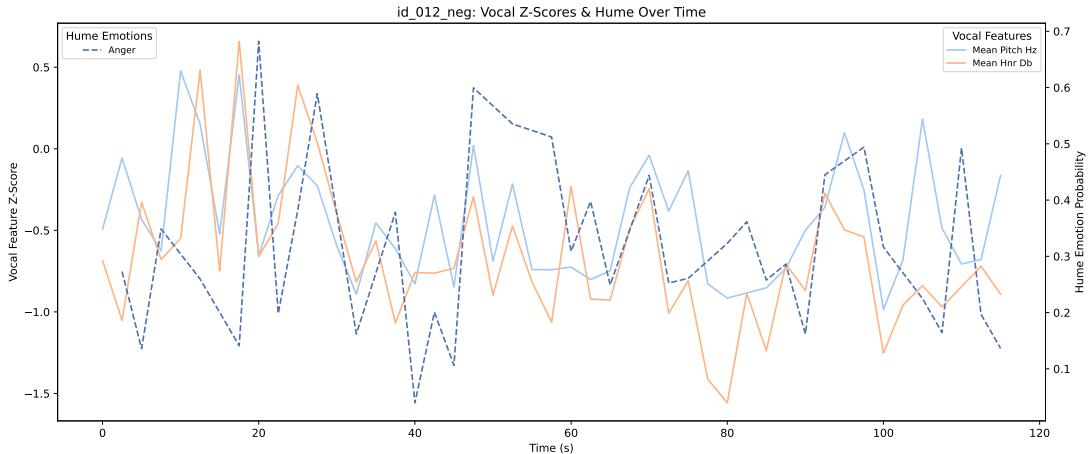


Figure 4.9: Pitch and HNR vs. Hume Anger for single negative clip.

(a) Clip id_012_neg – Pearson Correlation				
Feature	Emotion	r	p	Sign
pitch	fear	0.393	0.0122	Yes
intensity	fear	-0.465	0.0025	Yes
hnrr	anger	-0.114	0.4846	No
hnrr	fear	0.437	0.0048	Yes
shimmer	sadness	-0.354	0.0248	Yes
shimmer	fear	-0.356	0.0240	Yes

(b) Clip id_012_neg – t-test				
Feature	Emotion	t	p	Sign
pitch	anger	-0.353	0.7259	No
pitch	fear	2.700	0.0103	Yes
intensity	anger	2.255	0.0300	Yes
hnrr	anger	0.880	0.3844	No

Table 4.21: Statistical results for clip id_012_neg

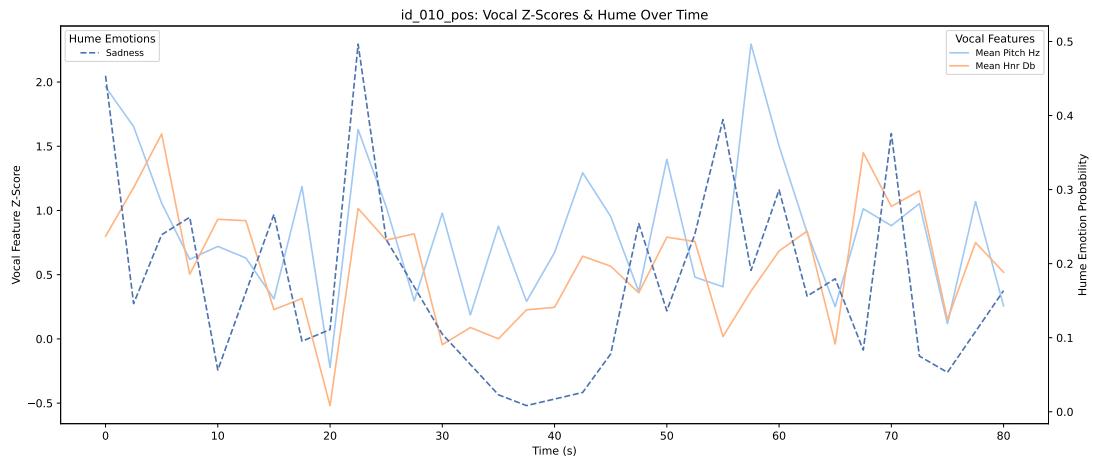


Figure 4.10: Pitch and HNR vs. Hume Sadness for single positive clip.

(a) Clip id_010_pos – Pearson Correlation				
Feature	Emotion	r	p	Sign
mean_pitch_hz	sadness	-0.032	0.8632	
mean_pitch_hz	fear	0.399	0.0260	
mean_pitch_hz	surprise	0.383	0.0332	
mean_hnr_db	sadness	-0.149	0.4225	
mean_hnr_db	fear	0.389	0.0306	
mean_hnr_db	surprise	0.406	0.0233	
jitter_local	anger	-0.359	0.0476	
shimmer_local	joy	0.370	0.0405	

(b) Clip id_010_pos – t-test				
Feature	Emotion	t	p	Sign
mean_pitch_hz	sadness	-0.876	0.3885	No
mean_hnr_db	sadness	-0.378	0.7082	No
jitter_local	anger	-2.047	0.0499	Yes
shimmer_local	sadness	-2.663	0.0125	Yes

Table 4.22: Statistical results for clip id_010_pos

4.2.7 Conclusion RQ1 Data Analysis

The results revealed only weak to moderate correlations for the analysis between individual vocal features and how Hume AI predicted emotions, where intensity and pitch showed most patterns consistently. The custom vocal categorization method did not function well in this context and resulted in very uniform results. This method was built on a basic group of vocal features which may overlooked important indicators for certain emotions. ANOVA tests found

no significant differences in vocal features across AI-labeled emotions. However, examining pitch and intensity fluctuations over time segments in individual clips gave more promising results. This implies that dynamic changes in vocal features can offer more insights than static averages when analysing conversational, yet spontaneous speech during interviews.

4.3 Data Analysis for RQ2: Text and Speech Based Emotion Recognition

Research Question 2 explores the degree to which two modalities for AI-based emotion recognition systems - speech-based (Hume AI) and text-based (NLP Cloud) - agree or diverge when labelling emotional expressions in semi-structured interviews. We examine five target emotions (anger, joy, sadness, fear, surprise) across the full dataset, as well as positive and negative interviews separately. To acquire a detailed picture of how the models align, we compare their average emotion scores, measure Pearson correlations and paired t-tests with Cohen's d. This multimethod approach supports a comprehensive understanding of how the two modalities responds to the same emotional input, to find mutual strengths and diverse tendencies in how they classify emotions.

4.3.1 Comparative Overview of Model Outputs

As presented in Table 4.6, Table 4.7, and Table 4.8 (4.1.3 Data Collection), the mean emotion scores and standard deviations differ between the two models across the full dataset, including patterns within positive and negative interviews.

Figure 4.11 visualises these differences for positive and negatives recordings separately. As presented, anger in positive interviews was detected as significantly higher levels by Hume compared to NLP, that rated anger near zero. For the negative interviews, the rating was more aligned where NLP rated anger slightly higher. Joy is rates substantially high by NLP in the positive interviews, compared to both other emotions and Hume's probability. In contrast, Hume rates joy higher than NLP for negative recordings. Sadness and fear are both rated higher by Hume than NLP in positive contexts, while NLP rates sadness higher in negative contexts where fear has more aligned scoring by the systems. Surprise was detected at similar, low levels by both models for both sentiment categories. Highest contrast for surprise is found in positive interviews where NLP rated it slightly higher.

The differences in the average emotion scoring are presented further in Figure 4.12. Positive values indicate that Hume AI assigned higher scores for respective emotions, while negative values imply higher scores from NLP Cloud. As explained for Figure 4.11, the most evident difference was shown for joy in both sentiment contexts, where NLP rates it significantly higher in positive settings and Hume higher in negative settings. Differences for sadness and surprise were insignificant in negative interviews, aligned with surprise in positive interviews. In Figure 4.12, the divergence in rating of fear in negative contexts is obvious where NLP rated the emotion more frequent.

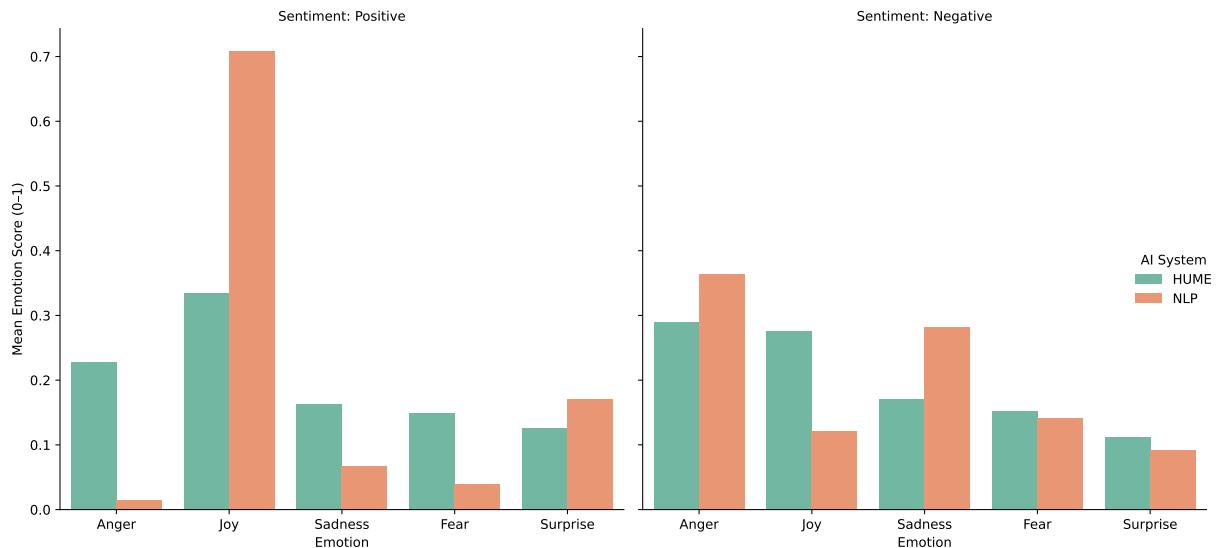


Figure 4.11: Average emotion score for Hume AI and NLP Cloud, separated by positive and negative recordings.

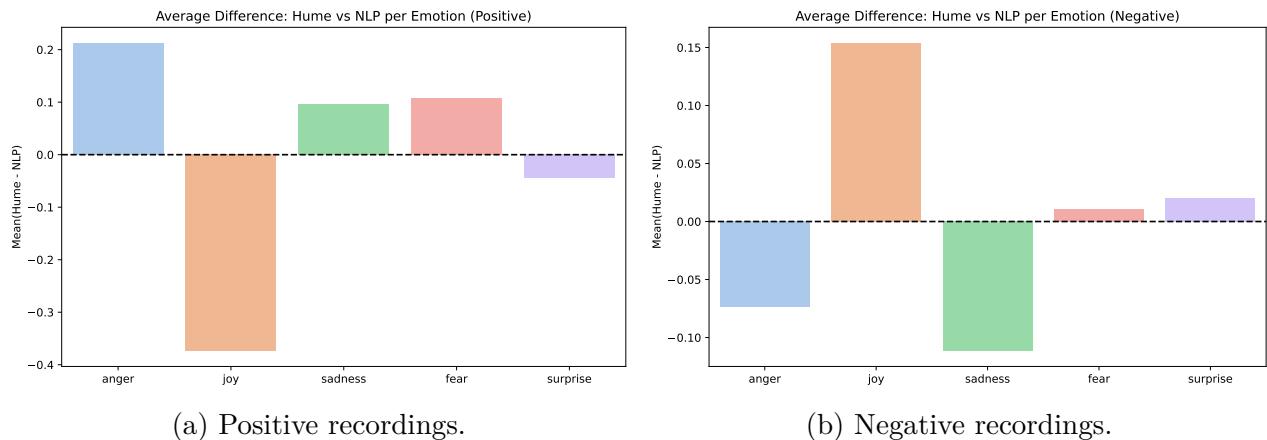


Figure 4.12: Average difference in emotions scores between Hume and NLP.

4.3.2 Statistical Analysis

Correlation Analysis

To evaluate how text-based (NLP Cloud) and speech-based (Hume AI) emotion recognition aligns, Pearson correlation coefficients (r) were calculated for each emotion across all interview recordings. Table 4.23 includes the full dataset (positive and negative recordings), presenting the correlation values as well as corresponding p-values to examine the statistical significance.

All recordings

Emotion	Pearson r	p-value	Significant
Anger	0.466	0.007	Yes
Joy	0.521	0.002	Yes
Sadness	0.167	0.362	No
Fear	0.171	0.348	No
Surprise	0.197	0.281	No

Table 4.23: Pearson Correlations Between NLP and Hume Emotion Scores (Full dataset)

This data demonstrates a reasonable positive correlation for Anger ($r = 0.466$, $p=0.007$) and Joy ($r=0.521$, $p=0.002$), implying that these emotions are relatively consistent identified throughout the AI systems on the full dataset. The p-values ($p<0.05$) show a statistical significance and highlights a relevant relationship in how Anger and Joy are detected through different processes. Sadness, Fear, and Surprise show contrasted results with weak correlations ($r<0.20$) where the p-values indicate no significance with low agreement between the AI models for these emotions when analysing the full dataset. Overall, some alignment for the more distinct emotions as Anger and Joy are declared through the correlation analysis, but some difficulties with consistent agreement are prominent for more nuances emotions as Sadness, Fear, and Surprise.

Positive Recordings

Emotion	Pearson r	p-value	Significant
Anger	0.160	0.568	No
Joy	0.682	0.005	Yes
Sadness	0.546	0.035	Yes
Fear	0.098	0.729	No
Surprise	-0.050	0.860	No

Table 4.24: Pearson Correlations Between NLP and Hume Emotion Scores (Positive)

Table 4.24 presents the same data as Table 4.23, but for positive recordings separately. As for the full dataset, Joy shows a significant correlation ($r = 0.682$, $p = 0.005$) between the model's detection. In contrast, Anger has a lower correlation ($r = 0.160$, $p = 0.568$) in positive contexts and Sadness presents a significant correlation ($r = 0.549$, 0.035) distinct from the full dataset. Fear and Surprise has even lower correlations for positive recordings than the dataset combined, with p-values close to 1.

Negative Recordings

Emotion	Pearson r	p-value	Significant
Anger	0.260	0.313	No
Joy	0.556	0.020	Yes
Sadness	0.028	0.914	No
Fear	0.270	0.294	No
Surprise	0.209	0.422	No

Table 4.25: Pearson Correlations Between NLP and Hume Emotion Scores (Negative)

Table 4.25 summarizes the correlation coefficients for the negative recordings. Consistent with the full dataset and positive subset, Joy again demonstrated a significant correlation in negative contexts ($r = 0.556$, $p = 0.020$). All other emotions failed to reach significance, with values that markedly diverged from their corresponding values in the positive recordings: Sadness resulted $r = 0.028$ ($p = 0.914$) versus $r = 0.546$ ($p = 0.035$) for positives, and Fear showed $r = 0.270$ ($p = 0.294$) compared to $r = 0.098$ ($p = 0.729$) in the positive context.

Paired t-Tests and Effect Sizes

To further explore alignment and differences between speech-based (Hume AI) and text-based (NLP Cloud) emotion recognition, paired t-tests and Cohen's d were conducted. Table 4.26 shows the t-statistics, p-values, and Cohen's d for each emotion across the full dataset. Positive t-values implies that Hume rated that emotion more frequent than NLP, negative t-values suggest the opposite.

Full Dataset

Emotion	t-statistic	p-value	Significant	Cohen's d
Anger	1.717	0.096	No	0.303
Joy	-1.726	0.094	No	-0.305
Sadness	-0.548	0.588	No	-0.097
Fear	3.341	0.002	Yes	0.591
Surprise	-0.657	0.516	No	-0.116

Table 4.26: t-statistics, p-value with significance, and Cohen's d for all clips.

Across all interviews, only Fear had statistically significant difference between the AI-models ($t = 3.341$, $p = 0.0022$), and had a medium effect size (Cohen's d = 0.591). Hume AI rated fear consistently higher than NLP Cloud, suggesting a systematic modality difference for this emotion. Although Anger, Joy, Sadness, and Surprise had some mean-score differences, none reached statistical significance (all $p > 0.05$) and their effect sizes were small ($|d| < 0.03$). Apart from Fear, the two models demonstrated close agreement in recognizing these emotional expressions.

Positive Recordings

Emotion	t-statistic	p-value	Significant	Cohen's d
Anger	10.903	0	Yes	2.815
Joy	-11.665	0	Yes	-3.012
Sadness	6.177	0	Yes	1.595
Fear	5.125	0	Yes	1.323
Surprise	-1.723	0.107	No	-0.445

Table 4.27: t-statistics, p-value with significance, and Cohen's d for positive interviews.

Table 4.27 demonstrates t-tests and Cohen's d for positive oriented interviews, where all emotions except for surprise ($t = -1.723$, $p = 0.107$, $d = -0.445$) shows significant differences ($p < 0.001$) with certainly large effect sizes. Negative T-value and Cohen's d for Joy ($t = -11.665$, $d = -3.012$) indicates that NLP have the aspects of overestimating this emotion compared to Hume with large effect sizes, where Hume in contrast tends to overestimate Anger ($t = 10.903$,

$d = 2.815$) in positive contexts. Hume rates Sadness and Fear more prominent than NLP, and Surprise remain inconsistent as previous results with no significant difference ($t = -1.723$, $p = 0.107$).

Negative Recordings

Emotion	t-statistic	p-value	Significant	Cohen's d
Anger	-1.702	0.108	No	-0.413
Joy	3.720	0.002	Yes	0.902
Sadness	-3.796	0.002	Yes	-0.921
Fear	0.536	0.599	No	0.130
Surprise	1.311	0.208	No	0.318

Table 4.28: t-statistics, p-value with significance, and Cohen's d for negative interviews.

Table 4.28 presents conducted t-tests and Cohen's d in negative interviews, with significant differences for Joy ($t = 3.720$, $p = 0.002$), where Hume rates it significantly higher than NLP. In contrast, NLP has clear higher scoring for Sadness with large effect size ($t = -3.796$, $d = -0.921$). However, the effect sizes are not as big as for the positive recordings. For example, the effect sizes for Joy ($t = 0.902$) are lower than Joy in positive contexts ($t = -3.012$) where NLP overestimated the emotion compared to Hume. Anger has a moderate difference, even if it is not statistically significant. No notable differences are detected for either Fear or Surprise. This implies that the AI systems strongly disagrees on Joy and Sadness detection in the negative contexts of the dataset.

Conclusion Statistical Analysis

Comparison of speech-based (Hume AI) and text-based (NLP Cloud) with statistical analysis demonstrates correlation particularly for clear expressed emotions as Anger and Joy when analysing the full dataset. However, anger shows no correlation between the models for either positive or negative recordings when separated. Joy shows a significant correlation throughout all sentiment contexts, where t-tests confirmed that NLP had higher predictions for joy in positive contexts and Hume in negative. Emotions that are more subtle like Sadness, Fear, and Surprise, revealed low correlations for all sentiment contexts except positive that showed a strong correlation for sadness, indicating modality-specific distinctions. Paired t-tests strengthened this observation regarding the full dataset and negative subset, pointing out Fear as the only emotion with statistically significant divergence in the full dataset where speech-based analysis assigned higher scores consistently. However, negative recordings showed significant difference for joy and sadness, while positively oriented clips showed significant difference for all emotions except surprise.

Conclusion Sentiment-Based Analysis

In conclusion, Hume AI and NLP Cloud show moderate to strong agreement on anger ($r = 0.47$) and joy ($r = 0.52$) across the full dataset, but weak correlations on sadness, fear, and surprise. Paired t-tests showed that fear is the single emotion that exhibits a significant mean difference across the full interview set (Hume > NLP, $d = 0.59$), while joy and anger showed modality divergencies in positive and negative subsets where NLP overestimated joy in positive interviews ($d = 3.01$) and Hume overestimated anger ($d = 2.82$). These results suggest that text and speech modalities agree on certain emotions particularly when considering the full dataset. However, divergencies occur for sentiment-specific analyses, especially for positive interviews.

Case Example

single clip comparison

briefly illustrate how speech vs text differ in practice

4.3.3 Conclusion of RQ2 Data Analysis

The results of this research question show that even if Hume AI and NLP Cloud partially aligns in detecting emotions, certainly for clearly expressed emotions such as Anger and Joy, they diverge significantly in their predictions of more nuanced emotions such as Fear, Sadness, and Surprise. Statistical tests confirmed a significant difference for Fear. Sentiment-based analysis showed that emotional context have an impact on the results, when analysing five basic emotions, where positive scenarios had a larger model divergence. As discussed above, the interview setting and overall data collection may have different impacts on the results. Still, the findings highlights how speech- and text-based models are complementary, each with their own strengths to capture different aspects of emotion expression, and indicate that relying on a single modality could have limitations for comprehensive emotion detection in speech.

4.4 Data Analysis for RQ3: AI and self-assessed emotion labels

The third research question explores how AI-generated emotion labels - from both speech-based (Hume AI) and text-based (NLP Cloud) – aligns with participants' own emotion ratings, to evaluate the agreement and divergence in different interview sentiments (positive and negative). This section includes average emotion scores from self-reports, Hume AI, and NLP Cloud across all recordings and for each sentiment category. Linear agreements are quantified with Pearson correlations and mean-level differences are analysed with paired t-tests and Cohen's d for assess effect sizes. This approach allows to see the overall alignment between AI-models and participants own assessment as well as how it depends on the sentiment context.

4.4.1 Model Emotion Score and Self-Reports Comparison

An initial overview is summarised in Table 4.6, Table 4.7, and Table 4.8 (4.1.3 Data Collection), with average emotion scores across all 30 interview recordings for each emotion category (anger, joy, sadness, fear, surprise). The table presents mean values and standard deviation for self-reported scores aside both AI-systems. Table 4.7 includes the same values for positive recordings and Table 4.8 presents the data from negative recordings.

These differences are visualized in Figure 4.13, illustrating a bar chart that compares the average emotion scores defined by Hume AI, NLP Cloud, and participants self-assessment seperated by positive and negative oriented interviews.

For the positive recordings, Joy consistently had higher self-reported scores than other emotions. NLP rates Joy higher than the participants while Hume rates it lower. In contrast, Joy was markedly rated higher by Hume in negative contexts than NLP and self-reporting which rated the emotion equally. The negative related emotions (Anger, Sadness, Fear) were assessed at lower levels by participants in the positive interviews. Self-assessed scores generally matched Hume's higher detection of Sadness, Fear and Surprise than NLP's minor predictions. Anger had higher rating by Hume than both self-reports and NLP in positive contexts. Surprise had similar average score across all sources, slightly lower detection rate by Hume.

For negative recordings, Anger had similar rating across all sources, with slightly higher rating by NLP. Joy has markedly higher average score by Hume compared to the other sources,

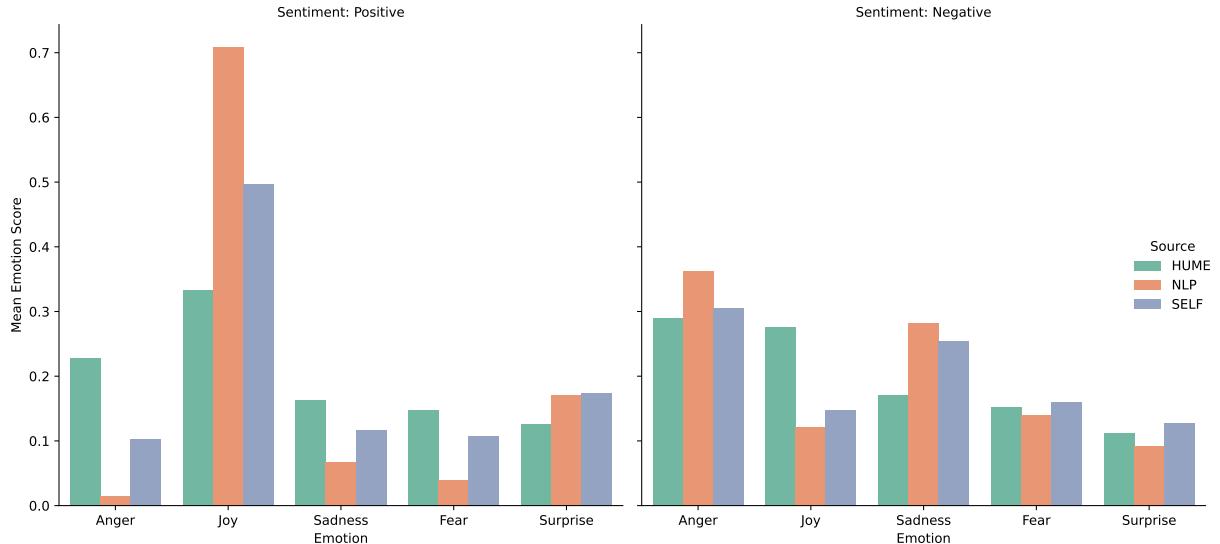


Figure 4.13: Comparison of emotional labels for Hume, NLP, and self-assessed.

while the speech-model rates Sadness lower than the text-model and participants. Fear and surprise had aligned rating by all sources, with a similar pattern where both emotions are rated slightly higher by the participants, closely followed by Hume and lowest rating by NLP.

This comparison suggests that emotional ranking are more aligned in for the negative recordings, where Joy is the emotion most distinct in the rating by Hume. The positive oriented interviews have more varying results between the sources, Joy are significantly rated higher by NLP while the text-based model rates Anger close to zero compared to Hume that rates this emotion higher than all other emotions except for Joy.

The sentiment-based comparison clearly presents that emotional expression and self-awareness have a significant variance between modalities and emotional contexts. Explicit emotions articulated in words are closely aligned between self-assessed rating and text-based analysis. While implicit or subtle emotions expressed through vocal tone have a notable divergence.

4.4.2 Correlation and Visual Analysis

To evaluate the alignment between AI-generated emotion scores and participants self-reported emotions, Pearson correlation analyses were conducted across the five emotion categories for both speech-based (Hume AI) and text-based (NLP Cloud) compared to self-reporting. With these measurements the relationship's strength and direction and the statistical significance can be reviewed.

Hume AI vs Self-Reported Emotions

All Hume

Emotion	Pearson's r	p-value	Significant
Anger	0,359	0,043	Yes
Joy	0,334	0,062	No
Sadness	0,050	0,784	No
Fear	-0,007	0,969	No
Surprise	0,088	0,631	No

Table 4.29: Pearson's r, p-values, and significance for all Hume recordings.

The correlation results for Hume AI predictions on all recordings in the dataset is demonstrated in Table 4.29, and indicate generally weak correlations across the majority of emotions. Anger is the only emotion showing a statistic significant correlation ($r = 0.359$, $p = 0.043$), which indicates a moderate alignment between Hume AI's speech based emotion detection and participants own perception for this emotion. Joy shows a moderate correlation but without statistical significance ($r = 0.334$, $p = 0.062$), other emotions, such as Fear ($r = 0.007$, $p = 0.969$), presents no relevant correlation.

(a) Positive Recordings (Hume)				(b) Negative Recordings (Hume)			
Emotion	r	p	Sign.	Emotion	r	p	Sign.
Anger	0.404	0.136	No	Anger	-0.105	0.690	No
Joy	0.401	0.138	No	Joy	0.127	0.627	No
Sadness	0.320	0.244	No	Sadness	-0.146	0.576	No
Fear	-0.027	0.924	No	Fear	-0.036	0.891	No
Surprise	0.091	0.748	No	Surprise	-0.143	0.585	No

Table 4.30: Pearson's r, p-values, and significance for Hume AI vs. self (positive and negative).

Table 4.30a presents correlation coefficients for positive recordings with no significant agreement occurs between Hume predictions and self-reported emotions. Anger, sadness, and sadness show moderate correlations ($r = 0.320-0.404$) with no statistical significance ($p = 0.136-0.244$). Weak correlation appears for both fear and surprise with high p-values suggesting no convincing evidence for these correlations.

Negatively oriented interviews, Table 4.30b show similar results as for positive interviews where no correlations of significance are found ($r = -0.146-0.127$, $p = 0.576-0.0.891$). Four out of five emotion correlations are negative while joy has a weak positive relationship. Each correlation is considered weak without statistical significance, implying that Hume predicted emotions distinct from participants own evaluation.

NLP Cloud vs Self-Reported Emotions

All NLP

Emotion	Pearson's r	p-value	Significant
Anger	0,739	0,000	Yes
Joy	0,863	0,000	Yes
Sadness	0,710	0,000	Yes
Fear	0,669	0,000	Yes
Surprise	0,092	0,616	No

Table 4.31: Pearson's r, p-values, and significance for all NLP recordings.

Table 4.31 presents correlation coefficients between self-reported and NLP-predicted emotions for the full dataset. When analysing the full dataset, NLP Cloud demonstrated strong and statistically significant correlations with self-reporting for four of five emotions. Joy showed the strongest correlation ($r = 0.863, p < 0.001$), followed by Anger ($r = 0.739, p < 0.001$) and Sadness ($r = 0.710, p < 0.001$). Fear had a moderately strong correlation with high statistical significance ($r = 0.669, p < 0.001$). Surprise was the single emotion showing weak correlation with no statistical significance ($r = 0.092, p = 616$).

(a) Positive Recordings (NLP)				(b) Negative Recordings (NLP)			
Emotion	r	p	Sign.	Emotion	r	p	Sign.
Anger	-0.199	0.477	No	Anger	0.286	0.266	No
Joy	0.622	0.013	Yes	Joy	0.366	0.149	No
Sadness	0.363	0.183	No	Sadness	0.429	0.086	No
Fear	0.527	0.043	Yes	Fear	0.599	0.011	Yes
Surprise	0.011	0.969	No	Surprise	-0.146	0.575	No

Table 4.32: Pearson's r, p-values, and significance for NLP Cloud vs. self (positive and negative)

Table 4.32a presents correlation data between self-reports and NLP Cloud for positive interviews, where lower alignments between self-reports and NLP is found compared to the full dataset. Only correlations for Joy ($r = 0.622, p = 0.013$) and Fear ($r = 0.527, p = 0.043$) are statistically significant. Sadness had a moderate correlation without statistical significance, Anger and Surprise presented weak correlations.

Correlation coefficients for negative interviews are presented in Table 4.32b, with similar results as for the positive interviews with weaker correlations compared to the full dataset. The single strong correlation with statistical significance is Fear ($r = 0.599, p = 0.011$). Moderate correlation is found for Joy ($r = 0.366, p = 0.149$) and Sadness ($r = 0.429, p = 0.086$), both with no statistical significance. As for the positive recordings, both Anger and Surprise had weak correlations between NLP and self-reporting.

Visual Correlation

Figure 4.14 illustrates the correlation between self-reported Anger scores and AI-labelled predictions for positive oriented interviews, while Figure 4.15 illustrates the correlated data for negative interviews. As shown, Hume shows a moderate positive correlation with no statistical

significance ($r = 0.40$, $p = 0.136$) where the data points have some spreading around the trend line. NLP Cloud shows a weaker correlation with self-reports for anger ($r = -0.20$, $p = 0.477$) than Hume in positive recordings, as demonstrated in the Figure 4.14 for NLP vs Self where data points are spread out vertically in line with the 0.0 axis.

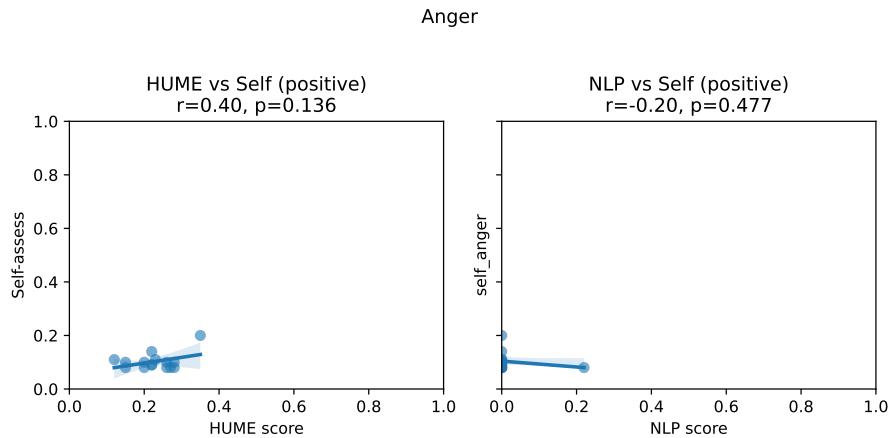


Figure 4.14: Scatter plot, Hume, NLP vs. Self for Anger.

The correlation coefficients remain low in Figure 4.15, where NLP presents a moderate positive correlation ($r = 0.29$, $p = 0.266$) with self-assessed anger in negative contexts, while the relationship with Hume is weaker ($r = -0.10$, $p = 0.690$) than in the positive interviews. The dispersed data points around the trend line visualises the divergence between the AI-systems and participants own judgement.

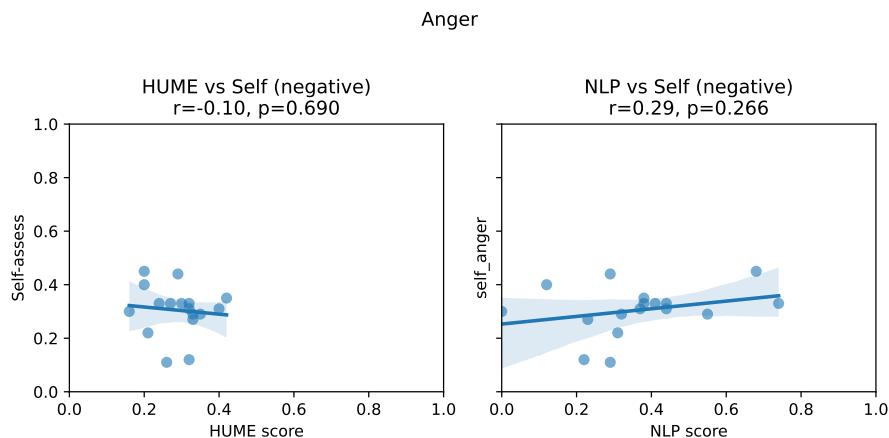


Figure 4.15: Scatter plot, Hume, NLP vs. Self for Anger.

Correlations for Joy is presented further in Figure 4.16 for positive recordings. Both Hume and NLP show a moderate to strong correlation with self-reported joy, NLP with the strongest correlation with statistical significance ($r = 0.62$, $p = 0.013$). This relationship is clearly presented with the data points being relatively close to the trend line for NLP, while the Hume diagram has more dispersed data points.

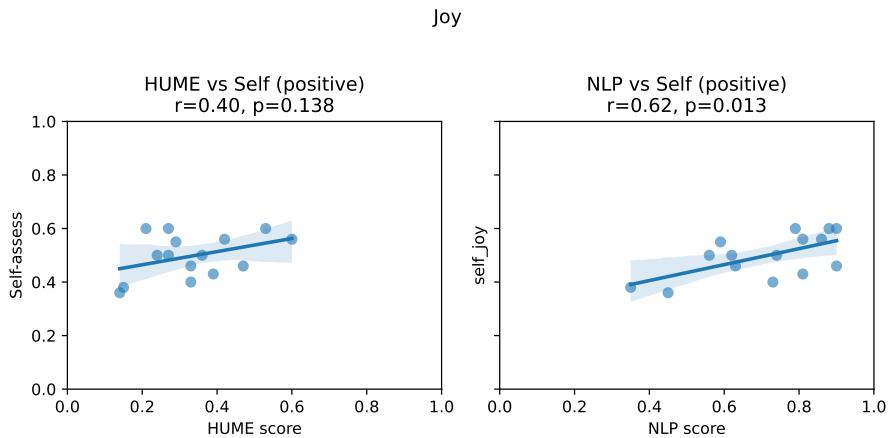


Figure 4.16: Scatter plot, Hume, NLP vs. Self for Joy.

Agreement between self-reported and AI-predicted joy is demonstrated for negative interviews in Figure 4.17. The trend where NLP has a higher correlation ($r = 0.37$, $p = 0.149$) remain, however the moderate relationship has no statistical significance. Hume shows a weak correlation, in contrast with the positive oriented interviews. Data points are more widespread for both Hume and NLP correlation with self-reported joy, suggesting varying rating of this emotion in negative contexts.

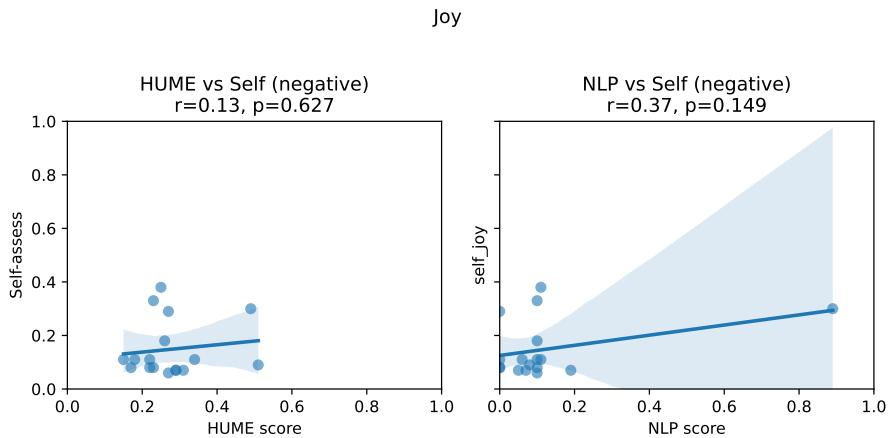


Figure 4.17: Scatter plot, Hume, NLP vs. Self for Joy.

Conclusion Correlation and Visual Analysis

MAYBE INCLUDE TEXT OR DELETE

4.4.3 Statistical Analysis and Effect Sizes

To explore if AI-generated emotion scores has a significant difference from self-reported emotions, paired t-tests were conducted for both Hume AI and NLP Cloud across each emotion for each sentiment. To evaluate the effect size of these differences, Cohen's d were calculated.

Hume vs Self-Reports

Table 4.33 presents the results on paired t-tests with Cohen's d to compare Hume AI's speech-based emotion scores to participants reports across the full dataset. As shown, Hume AI ratings

on anger are higher than self-reports ($t = 2.399$, $p = 0.023$, $d = 0.424$), suggesting a moderate tendency for Hume overestimating anger compared to participants own perception. In contrast, Hume underestimate Surprise relatively to self-reports ($t = -2.109$, $p = 0.043$, $t = -0.373$). No significant differences are found for sadness, fear, and surprise ($p > 0.05$, $|d| < 0.20$).

All Recordings

Emotion	t-statistic	p-value	Significant	Cohen's d
Anger	2.399	0.023	Yes	0.424
Joy	-0.271	0.788	No	-0.048
Sadness	-1.069	0.293	No	-0.189
Fear	1.052	0.301	No	0.186
Surprise	-2.109	0.043	Yes	-0.373

Table 4.33: Hume AI vs. Self—All Recordings (paired t-test & Cohen's d)

Table 4.34 seperates these comparisons by sentiment. In positive interviews, four of five emotion comparisons show significant differences. Hume tends to remarkably overestimate anger ($t = 8.776$, $p < 0.001$, $d = 2.266$) and underestimate joy ($t = -5.112$, $p < 0.001$, $d = -1.320$), and assigning notable higher scores for sadness and fear compared to participants own evaluation. Surprise is the single emotion that remains unsignificant. In negative interviews, significant differences appears for joy ($t = 3.878$, $p = 0.001$, $d = 0.941$) where Hume predicts higher levels than self-reported scores, and for sadness ($t = -2.890$, $p = 0.013$, $d = -0.677$) that is rated higher by participants than Hume. Other emotions show no reliable difference.

(a) Positive Recordings					(b) Negative Recordings				
Emotion	<i>t</i>	<i>p</i>	Sign.	<i>d</i>	Emotion	<i>t</i>	<i>p</i>	Sign.	<i>d</i>
Anger	8.776	0.000	Yes	2.266	Anger	-0.517	0.612	No	-0.125
Joy	-5.112	0.000	Yes	-1.320	Joy	3.878	0.001	Yes	0.941
Sadness	2.451	0.028	Yes	0.633	Sadness	-2.790	0.013	Yes	-0.677
Fear	2.463	0.027	Yes	0.636	Fear	-0.454	0.656	No	-0.110
Surprise	-1.855	0.085	No	-0.479	Surprise	-1.033	0.317	No	-0.250

Table 4.34: Paired t-test and Cohen's d for Hume AI vs. Self in positive and negative interviews.

These results suggests that Hume AI's speech-based assessments only have weak agreements with participant's self-reports, with varying alignment depending on emotion and sentiment context. In positive interviews, the model remarkably over- or underestimates anger and joy compared to self-reported emotions, while in negative interviews the only significant differences occur for joy and sadness.

NLP Cloud vs Self-Reports

Paired t-tests with Cohen's d for comparison of NLP Cloud's text-based emotion scores and participants' self-reports for all recordings are presented in Table 4.35. Significant differences are found for joy, where NLP rates it higher than self-reports ($t = 2.331$, $p = 0.026$, $d = 0.412$). In contrast, NLP tends to underestimate fear comparing to self-reports ($t = -3.496$, $p = 0.001$, $d = 0.618$). Anger, sadness, and surprise show now significant difference ($p > 0.05$, $|d| < 0.20$).

All Recordings

Emotion	t-statistic	p-value	Significant	Cohen's d
Anger	-0.373	0.711	No	-0.066
Joy	2.331	0.026	Yes	0.412
Sadness	-0.525	0.603	No	-0.093
Fear	-3.496	0.001	Yes	-0.618
Surprise	-1.011	0.320	No	-0.179

Table 4.35: Paired t-tests and Cohen's d for NLP Cloud vs. Self. All Recordings

Table 4.36 separates the comparisons by sentiment. In positive interviews, NLP rates anger at significantly lower levels than participants ($t = -4.853$, $p < 0.001$, $d = -1.253$), while rating joy higher than self-reports ($t = 6.066$, $p < 0.001$, $d = 1.566$). Both sadness and fear have a significant difference where NLP tends to underestimate these emotions compared to self-reports. Surprise remains without significance. In negative oriented recordings, no significant difference between NLP and self-reports are found ($p > 0.05$, $|d| < 0.389$), suggesting closer alignment during negative contexts.

(a) Positive Recordings					(b) Negative Recordings				
Emotion	t	p	Sign.	d	Emotion	t	p	Sign.	d
Anger	-4.853	0.000	Yes	-1.253	Anger	1.335	0.200	No	0.324
Joy	6.066	0.000	Yes	1.566	Joy	-0.578	0.571	No	-0.140
Sadness	-2.852	0.013	Yes	-0.736	Sadness	1.073	0.299	No	0.260
Fear	-4.603	0.000	Yes	-1.188	Fear	-1.149	0.267	No	-0.279
Surprise	-0.075	0.941	No	-0.019	Surprise	-1.604	0.128	No	-0.389

Table 4.36: Paired t-test and Cohen's d for NLP Cloud vs. Self in positive and negative interviews.

Overall, these results implies that NLP Cloud have higher agreement with self-reports in negative contexts, but in positive interviews notable divergencies are found for certain emotions, most appearing for anger, joy and fear.

4.4.4 Conclusion of RQ3 Data Analysis

The result for the full dataset shows that Hume AI only has two small significant divergencies from self-reports, where anger is overestimated ($d = 0.42$) and surprise underestimated ($d = -0.37$) compared to ratings by the participants, while all other emotions show no significance in t-tests and weak correlations. When separated by sentiment, Hume tends to overestimate anger and underestimate joy relatively to self-reports in positive contexts, in negative interviews it diverges on joy and sadness. In contrast, NLP Cloud are closely aligned with self-reports in negative contexts, with no significant differences, but in positive interviews NLP remarkably underestimates anger ($d = -1.25$), rates fear at lower levels, and joy at higher levels compared to self-reports. Correlation analyses strengthen these patterns, where NLP correlations with self-assessed scores are strong for anger, joy, sadness, and fear in the full set, and Hume only shows a moderate correlation for anger and weaker correlations for other emotions. However, the correlations are not as strong when separating the interviews by sentiment. Overall, text-based emotion detection with NLP Cloud has higher agreement with

participants' self-assessments, especially in negative interviews, while speech-based detection with Hume have more variations between positive and negative contexts. These results shows that each modality captures distinct features when comparing with human-labelled rating of their own emotions, and the alignment is fluctuating depending on the interview sentiment.

Discussion

5.1 Result Discussion RQ1

For the first research question, this thesis investigated how AI models for speech recognition compare to existing research on vocal markers. More specifically, the goal is to assess whether the AI models align with the findings of the Swedish research on vocal markers done by Ekberg (Ekberg et al., 2023).

5.1.1 Interpretation of Results

Vocal Features and Hume AI Emotion Scores

Correlations between different vocal markers (pitch, intensity, HNR, jitter and shimmer) and Hume AI emotion labels, were visualized on a heatmap using Pearson correlation coefficients. Overall, low to very low correlations were found here, with only a few instances of what could be considered moderate correlations. Where values closer to 1.0 for positive linear relationships or -1.0 for negative linear relationships would suggest a strong correlation, the heatmap showing the correlation between the vocal features and AI emotion scores from Hume AI primarily showed low numbers ranging from -0.44 as the lowest, to 0.34 as the highest. Some numbers appeared to have an extremely low correlation. For example, the correlation between sadness and shimmer was as low as 0.00, with other correlations hovering around being 0.01-0.04. This suggests a weak linear relationship between the vocal features mean pitch, mean intensity, mean HNR, jitter and shimmer and the emotion labels from Hume AI, indicating that the AI model may not fully have captured the complex details from the vocal markers as well as the findings from Ekberg (Ekberg et al., 2023).

Despite this, some of the results align with the findings of Ekberg. For example, in this analysis mean pitch with anger shows the value $r = 0.26$ while mean pitch with sadness gave the value $r = -0.37$. Ekberg's research showed anger to have elevated pitch and while sadness had lower pitch than anger and happiness, so both studies show anger to be associated with elevated pitch and sadness to be associated with lower pitch. Looking at intensity, this analysis showed the value of joy to be $r = 0.34$, sadness $r = -0.26$, and surprise $r = -0.39$. In Ekberg's study, happiness (joy) showed higher intensity while sadness and surprise showed patterns of lower intensity. Although these emotions show matching patterns, other emotions show mismatches. Ekberg's study reported higher intensity for anger and fear, while the heatmap analysis in this study presents fear to have a minimal correlation ($r = 0.01$), while anger has a negative correlation ($r = -0.18$) with intensity, suggesting that intensity decreasing as anger increases.

The findings for HNR only partially matched with the findings of Ekberg's study which reported that fear and happiness were linked with higher HNR while sadness is associated to lower HNR. In this study sadness showed a moderate negative correlation with HNR ($r = -0.44$), which aligns with the Ekberg study. However, correlations for fear ($r = 0.08$) and joy ($r = 0.06$) were very weak, differing from the associations observed by Ekberg. Jitter showed no moderate or strong correlations with any emotions in this study, and shimmer proved to be slightly positive for joy ($r = 0.19$) and moderately negative for anger ($r = -0.33$) but not for

surprise which was the only emotion in Ekberg's study which had higher shimmer.

Although some Pearson correlations presented in the heatmap showed consisting pattern with the Ekberg research and others diverged, it is important to note that there are methodological differences between this analysis and Ekberg's research. Ekberg employed a different statistical approach using both simple and multiple logistic regression models to predict the emotions from speech. While Pearson correlation which was used for this analysis is useful for detecting linear associations, it might not have captured the complex non-linear interactions that the logistic regression models are able to capture. The weak correlations found in the heatmap showing the correlation between the vocal features and AI emotion scores from Hume AI might also suggest that the vocal features used in this analysis were insufficient in predicting the emotions in spontaneous speech. Where the research done by Ekberg uses an acted dataset with repeated sentences, emotional expression in natural speech (e.g. interviews, which were used in this analysis) tends to be more subtle.

Praat-Based Emotion Scores

As for the second heatmap presented in the result, we see the correlations between vocal features and the emotion scores from the custom Praat-based categorization function. While the Pearson correlation values in this heatmap are seemingly stronger than with the Hume AI labels, these results may be misleading as they reflect an over-reliance on pitch and HNR, rather than authentic emotional differences. For instance, examining the highest correlations, pitch correlated strongly with fear ($r = 0.93$) and surprise ($r = -0.90$). HNR also demonstrated high correlations with fear ($r = 0.92$) and surprise ($r = -0.90$).

While the correlations here suggest a strong relationship between the emotional labels and the vocal features, this specific pattern indicates methodological limitations. The Praat-based categorization function appears to prioritize a narrow range of features (notably pitch and HNR), leading to inflated correlations that do not necessarily capture the complexity of emotional expression.

The failure of accurately categorizing vocal emotions is likely due to the spontaneous nature of the interviews which unlike acted datasets where emotions are exaggerated, result in subtle emotional expression. The limited number of vocal features may also be a contributing factor as to why full emotional complexity was not captured, together with variability in recordings which potentially diluted the emotional markers across time. Contrary to theoretical expectations, the categorization failed to distinguish between the emotional states in a meaningful way as indicated by the lack of clear emotional differentiations. These findings suggest that although some acoustic features were captured effectively, the Praat-based function did not categorize the emotions accurately, potentially due to the speech being in a spontaneous interview format instead of, for example, an acted dataset.

Custom Vocal Emotion Categorization Method

To assess the effectiveness of the custom emotion categorization function developed for this study, the outputs were compared to the Hume AI generated emotion scores.

Despite individual vocal features having high correlations, the Praat-based emotion categorization function reveals significant limitations. With an average score of 0.2, the Praat based function presented minimal variability, suggesting that the function was unable to differentiate emotional expressions within the spontaneous speech obtained from the interviews.

The Hume model demonstrated wider variations for the emotion scores, reflecting on a more nuanced detection of the different emotions. Several factors likely contributed to this outcome as these results highlight the challenge of categorization in spontaneous speech. In real world

emotional expressions, as opposed to acted ones, emotional signals are often more subtle and dependent on context.

The limited set of acoustic features may as well have excluded some important clues, further restricting the sensitivity of the function. All factors point to spontaneous speech potentially being too complex for this function.

ANOVA Vocal Features

To further investigate whether fundamental vocal features used in the previous analyses varied systematically across the five different AI-labeled emotions, an analysis of variance (ANOVA) was conducted. None of the results revealed any significant statistical differences for pitch, intensity, HNR, jitter or shimmer. This was further confirmed by Tukey HSD tests, suggesting that within the context of spontaneous speech obtained from the interviews, emotional states may not reliably be differentiated by average values of the core vocal features.

Clear differences in acoustic features between emotions were reported by Ekberg (Ekberg et al., 2023), which contrasts with these findings, possibly due to the controlled nature of Ekberg's acted dataset. The lack of variance for the present study likely reflects the subtle nature of emotional expressions in an interview format, being a natural conversational form of expression with some emotions possibly being interwoven with context and which may also have more variations from time to time.

Correlation Between Vocal Features and Hume AI Emotion Scores

Following the limitations identified with the rule-based emotion categorization, the analysis shifted to examining direct correlations between raw acoustic features and Hume AI's emotion scores. With a composite correlation analysis visualized, showing the two acoustic features pitch and intensity and the correlation with Pearson correlation coefficients, generally weak correlations were found across all emotions.

However, some patterns aligned with established findings of Ekberg (Ekberg et al., 2023), where intensity showed a negative correlation with fear, sadness and surprise which bears some resemblance with the results of Ekberg's study. Although pitch presented minor positive correlations with anger and fear, and happiness, these results alone suggest that average pitch and intensity alone are insufficient in capturing the complex nature of emotional expression made in spontaneous speech, as it remains too subtle. Although some expected relationships have been observed, emotions fluctuate dynamically within a clip, contributing to the difficulty of detecting consistent patterns through static averages of single features.

Observations from individual interviews

Analyses of individual interviews revealed results can give more clarity by analyzing a voice recording in its entirety combined with looking at peaks in specific moments. Where static measures often result in an average of emotions when analyzing an entire voice recording, segment analysis done from time to time shows more detailed emotional shifts. Although correlations are not entirely consistent throughout the recordings, the positive interview revealed increasing intensity corresponding with higher joy, whereas one of the negative voice recordings revealed peaks in pitch aligned with elevated anger scores. There are several possible factors that can account for this result, one being that the emotional expression in natural conversations such as the interview format for this study is context driven with emotions fluctuating throughout the entire conversation. Static measurements take an average of all the peaks and fluctuations, likely resulting in a somewhat neutral end result, where datasets with acted recordings

might not differ as much in one recording. These findings highlight the importance of analyzing emotions on a segment level rather than relying exclusively on an average for one whole clip. Looking at one clip at its entirety for a natural conversational clip gives the opportunity to inspect all emotional peaks and distinguish between different emotional states throughout the recording without emotions being averaged and neutralized. Furthermore some variability across individual speakers was observed demonstrating different vocal patterns suggesting that some personal vocal traits can impact the level of detectability of emotions. Some of these patterns could be due to factors such as gender or speaking style.

5.2 Result Discussion RQ2

For the second research question in this thesis the aim was to investigate whether we could understand the emotions from textual content of the speech, with the same data as in RQ1. This was achieved by transcribing the vocal recordings and analyzing them using NLP Cloud's emotion recognition to detect emotions in the textual content of the speech.

5.2.1 Interpretation of Results

5.2.2 RQ2: Speech-based AI vs Text-based AI

In the comparison between the speech-based emotion recognition model, Hume AI, and the text-based emotion recognition model NLP Cloud, the system overall seemed to show some levels of agreement for certain emotions. Using both descriptive statistics and visual analyses to calculate the differences, an overall comparison of both ai systems showed that the mean emotion scores differ across the two models. The average difference in the emotion scores showed values indicating that Hume AI obtained higher scores for the emotions anger and fear, while NLP Cloud proved to show higher scores for joy, sadness and surprise. Despite these findings, the score for sadness and surprise were sufficiently low, suggesting that the models were substantially aligned on specifically those emotions. Joy being highly scored by NLP Cloud indicates that joy may not have been as easily identified in speech-based emotion analysis, while the textual context may have conveyed a more positive tone from the text than appeared in the voice. In contrast, anger and fear appeared to have been more effectively captured by the speech-based emotion detection, possibly suggesting that someone might sound angry or fearful even though they may not be experiencing these emotions in the moment. Based on the Pearson correlation analysis showing the association between the text-based and speech-based emotion recognitions, the strongest alignments were shown for Joy ($r = 0.521$) and anger ($r = 0.468$). Joy and anger also showed statistically significant p-values, where joy had a p-value of 0.0069, and anger had a p-value of 0.0022. No further strong correlations or statistically significant p-values were found in the other emotions. Several factors may account for this result. For example, joy and anger are distinct emotions, while sadness, fear and surprise may likely involve more subtle cues and contextual factors. Being more complex to detect may have contributed to the lower consistency across the two models for these specific emotions.

For the full dataset, paired t-tests showed no significant differences for the mean score of the emotions across the dataset for all emotions except fear. Although the correlation between Hume AI and NLP Cloud showed non-significant scores for fear, the t-test indicated that while the systems do not align on detecting patterns for fear, Hume AI consistently rates the fear higher than NLP Cloud. Possible explanations for this result may reflect the differences in how emotions are conveyed and detected in the different models, whereas Hume AI possibly could have captured the more subtle vocal indicators that might not have been as easily expressed or detected in text.

Examining the t-tests for the positive oriented interviews in comparison to the negative oriented interviews, notable findings emerged. For the positive interviews, significant differences between Hume AI and NLP Cloud were found for all emotions with the exception of surprise, where Hume consistently detected higher levels of sadness and fear and NLP on the other hand overestimated joy and anger in comparison to Hume. This may be explained by the complexity of emotions and emotional expression. In the positive interviews, the participants discussed joyful topics, and while this may have been detected for the text-based emotion recognition, the vocal tone could reveal more subtle cues in the tone, rhythm and pitch. For a positive interview, the participant may have a lower and more neutral tone and pitch than an actor acting out happiness, which could be one explanation for this result. For the negative interviews, significant differences were only identified for joy and sadness, where Hume rated joy with a higher score, and NLP rated sadness higher. This indicates a better alignment for the different models for the analyses made for the negatively oriented interviews.

Possible explanations for these results are that people participating in the interviews may have used overly positive language out of politeness, even if the content of the words may have been negative.

Sentiment-Based Analysis RQ2

In comparing Hume AI and NLP Cloud, the sentiment-based analysis presented a distinct pattern in how the different AI models interpreted the positively oriented interviews versus the negatively oriented interviews, where some emotions were consistently rated higher than others.

NLP Cloud showed patterns of consistently rating joy higher than Hume AI, while Hume rated higher for negative emotions such as anger, sadness and fear in the positive interviews.

These results indicate that the subtle features such as pitch, loudness and more may have been interpreted by Hume as negative emotions even in positive conversations. Surprise remained a challenging emotion to detect, and NLP Cloud showed higher values of anger and sadness in the negative interviews, likely due to being able to better capture the negative context of the interviews through the text-based analysis. Hume rated joy unexpectedly high in the negative interviews, where a possible reason could be nervous laughter or other emotions that could have been misclassified. Overall, the results underscore that the two AI models differ in the job of emotion detection, possibly due to the vocal recordings involving subtly expressed emotions or possible irony or laughter that could have been incorrectly categorized as joy.

5.3 Result Discussion RQ3

For the third and final research question, the objective was to assess how the AI generated emotion labels obtained through the speech-based and text-based emotion recognition would compare to the self-reported emotions provided by the interviewees.

5.3.1 RQ3: AI vs Self-Assessed Emotions

For the third and final research question the alignment with the speech-based emotion labels from Hume AI and the text-based emotion labels from NLP Cloud in combination with self-assessed emotion scores were examined. Insightful findings revealed some levels of alignment dependent on both the model and emotion. With an analysis showing an average of the emotion scores across the entire dataset of interviews, joy emerged as the emotion with the highest average scores. Fear and surprise showed the lowest scores out of the emotions. A visualization of the average emotion scores across all three channels shows NLP Cloud overestimating joy

excessively, while Hume on the other hand overestimated anger to a certain degree. The rest of the emotions are relatively close in scores across the models and self-assessment scores, where sadness, fear and surprise were rated low for all channels. The low scores overall for sadness, fear and surprise could be explained by the spontaneous interview format in a calm setting, which may not encourage expressively conveying these emotions.

For the case of joy having a substantially higher score for NLP Cloud, the model may have interpreted language as joyful even though the tone was more neutral, also possibly missing out on cues such as irony. Hume estimated anger higher than NLP Cloud and the interviewees themselves, which may be due to misinterpreted signs of anger for example from pitch and intensity.

Hume AI and NLP Cloud vs Self-Reported Emotions

The correlations between Hume AI and the self-reported emotions indicated very weak to modest correlations for all emotions, where only anger showed a modest statistic significant correlation with a Pearson value of $r = 0.359$ and $p\text{-value} = 0.043$. This may be due to the nature of anger which often produces a distinct vocal change typically involving increased loudness and change of pitch, while emotions like fear and surprise may often be expressed with more subtle vocal expressions which may not have been captured as successfully. Once again, the calm setting in an interview environment might also have affected the results, further muting emotional expressions. This suggest that although Hume moderately detect some emotions based on voice, a multimodal approach with further analysis might be necessary for more complex emotion recognition.

The results for NLP Cloud presented statistically significant correlations for all emotions except surprise. This indicate a high degree of alignment between the self-reported emotions and the text-based emotion detection NLP Cloud, where the strongest correlations were presented in joy ($r = 0.863$, $p = 0.0000$), anger ($r = 0.739$, $p = 0.0000$), sadness ($r = 0.710$, $p = 0.0001$), and lastly fear ($r = 0.669$, $p = 0.0003$). This strongly indicates the effectiveness of NLP Cloud in capturing emotional content through text in combination with alignment with self-reported emotions gathered from the interviews.

Surprise being the only emotion not to show a strong correlation, emphasizes a consistent challenge observed over both models used in the study. During the self-evaluation segment of the interviews, multiple participants expressed certain confusion regarding the assessment of the emotion surprise. A large part of the interviews consisted of describing past emotional experiences which may have reduced the intensity of surprise. Typically, surprise is expressed as an immediate reaction to unexpected events and its unlikely that the interviewees are able to genuinely experience the same surprise felt in the original moment of the memory. This provides a possible explanation for why both AI models overall detected low levels of surprise, while an acted dataset could present higher correlations for this emotion. The statistical analysis made evident that both Hume AI and NLP Cloud showed partial alignment with the self-assessed values for some emotions. Hume AI showed a larger deviation for anger and surprise, whereas NLP Cloud deviated more for fear and joy. This suggest that the effectiveness of the AI models in emotion detection is not consistent across the emotions, although the challenging nature of self-assessing emotions, especially fear and surprise, retrospectively possibly complicates this process, effecting the results as well.

Sentiment-Based Analysis RQ3

In the sentiment-based analysis comparing Hume AI, NLP Cloud, and self-reported emotions insight was gained into how the AI models align with the personal perceptions of emotions.

In the positively oriented interviews, NLP showed higher rating of joy compared to the self-assessed scores, while Hume AI consistently rated joy lower than the self-assessed scores in combination with detecting higher levels of all negative emotions (anger, sadness and fear).

These results suggest that subtle vocal markers were captured by Hume that may not have matched the content. For surprise, NLP Cloud closely matched the scores of the self-assessment whereas Hume AI detected lower levels. This may reflect all challenges previously mentioned in regard to the emotion surprise.

In the negatively oriented interviews, both fear and surprise were relatively evenly rated across all three sources, with the self-assessed being the highest rated in both emotions.

The ratings for anger were high for all sources as well and fairly evenly matched between Hume and the self-assessed scores, while NLP Cloud rated anger higher. The higher rating by NLP Cloud was likely due to the context of negative wording in the negative interviews. Hume AI rated sadness low, while NLP Cloud was relatively close in score compared to the self-scores, suggesting the text-based model might have been better at capturing sad emotions from the vocal recordings. Hume may not have picked up the cues for sadness in the same capacity, likely due to the low expressions of sadness during the interviews. Further analysis showed Hume AI rating joy higher than both NLP Cloud and the self-reported emotions in the negative interviews, likely due to misclassifications of certain vocal cues that may have been subtle or complex, for example irony which could be difficult for a speech-based AI to recognize. Overall, the text-based AI NLP Cloud seems to align closer with the self-assessed scores rated by the participants of the interviews, possibly capturing the context for each interview more effectively. This underscores the limitations of relying exclusively on either speech-based or text-based emotion recognition.

5.4 Conclusions

5.4.1 Conclusion for RQ1

The result for the first research question has investigated if AI-based emotion recognition models align with existing research on vocal markers with a focus on the Swedish language. Exploring vocal markers correlation with Hume AI emotion labels as well as the correlation between vocal markers and Praat, the overall result demonstrated limited strength. Overall the results showed weak or moderate correlations, although some relevant patterns aligning with the existing research on Swedish vocal markers (Ekberg et al., 2023) was found. While the values from the Hume correlations appeared moderate at best, the Praat correlations overall showed weak correlations except some misleading numbers suggesting an over-reliance on the vocal markers pitch and HNR.

The segment level analysis gave important insight into the fluctuations in the emotions throughout entire clips compared to an average value of the different emotions.

The result indicates that while there are certain vocal features that remain relevant as indicators of emotional states, spontaneous speech presents challenges in emotion recognition. In comparison to acted datasets, emotions are more subtle with more variety and contextual dependence in spontaneous speech. Though there are challenges with spontaneous speech, the result suggest that AI-based emotion recognition systems such as Hume AI showed promise, demonstrating some flexibility and context-awareness.

Future work could benefit from incorporating a wider range of vocal features, emotions and a more dynamic approaches to capture the complexity of emotional expression.

5.4.2 Conclusion for RQ2

In answering RQ2 and RQ3, this study explored effectiveness of speech-based emotion recognition, Hume AI, text-based emotion recognition, NLP Cloud. These were later examined for potential alignments with self-assessed scores for emotions. For RQ2, partial agreements were found between the two AI models for some particular emotions such as joy and anger, though notable disagreements were present for fear and surprise.

This brought attention to the challenges of detecting the emotional cues for more complex emotions, where they might have more subtle cues for detection. In comparisons between the models, certain differences in how each model captured emotions were found.

5.4.3 Conclusion for RQ3

For RQ3, the comparison of the models with the self-assessed emotion scores indicated a stronger consistency for NLP Cloud than it did for Hume AI. Overall, surprise was presented as an emotion consistently challenging to detect across the models, possibly due to the complexity of the emotion combined with the nature of the interviews where this emotion may have been expressed the least.

The results for this study suggest that a multimodal approach integrating multiple sources could enhance the precision and reliability of the detection systems for emotions, as relying on only one type of analysis may not be enough for the complexity of emotions.

5.5 Method Discussion

5.5.1 RQ1 Methodological Considerations

To answer RQ1, the methodological approach involved analysis of emotional expression for vocal markers in Swedish speech in comparison to AI based emotion recognition models. The idea was to analyze emotions in a clip in its entirety and find correlations, which had some differing results, but it proved to be a notable strength to execute the analysis on a segment level to capture emotional fluctuations in a more dynamic way. While this offers another perspective, this approach introduced challenges of its own in having some inconsistent emotions not aligning completely across the segments. Therefore the methodological approach was partially fulfilled for answering RQ1 by identifying some emotional fluctuations, while also revealing challenges in both the analyses for segment-leveled clips and full clips.

One of the studies chosen to compare the results with, being the existing Swedish emotion research by Ekberg (Ekberg et al., 2023) proved some similarities and patterns which provided valuable information to this study. However, the research used pre-defined sentences, repeated by actors, which may have given a more consistent result than the dataset used in this research which consisted of interviews capturing spontaneous speech. With 16 participants, the dataset resulted in a total of 32 recordings across a diverse group of participants consisting of men and women with ages ranging from the 23-78. The spontaneous speech and large variety of interview questions combined with dataset size may indicate some limitations for the result. While RQ1 was addressed, a larger and more controlled dataset with acted emotions along with repeated sentences, could possibly have validated some observed patterns, ensuring more consistent emotions throughout the recordings.

Hume AI was one of the models used and provided some advantages such as avoiding manual labeling and being pre-trained, although the Hume AI emotion scores had to be normalized and the emotions were filtered to use only the specific five emotions necessary for the comparisons in this research, which may have had some limitations on the model's capacity. Along with

working well for the research's purpose, the model has some downsides. For example, there is limited publicly available information about functions of the model, making it difficult to fully assess possible limitations and biases.

Despite these limitations, Hume AI contributed with valuable insights in answering RQ1.

A set of basic vocal features which consisted of pitch, intensity, harmonic-to-noise ratio (HNR), jitter, shimmer, was extracted through Praat. These features are well established indicators of emotional expression but proved to be somewhat of a limitation which possibly could have been avoided by incorporating additional vocal features. Given that the dataset for this research consisted of interviews capturing spontaneous speech, a broader range of vocal features might have contributed to the detection of the complex vocal patterns and given a more nuanced understanding of the correlations for emotion recognition in the Swedish language.

While the selected vocal markings chosen for this study gave some insight into addressing RQ1, expanding the set of features could have helped address RQ1 more comprehensively.

5.5.2 RQ2 and RQ3 Methodological Considerations

The methodological approach to address RQ2 combines analysis of transcribed text in emotion recognition using NLP Cloud in order to assess the emotional content of speech transcripts in relation to speech-based AI models.

In addressing RQ3, the approach was to compare self-reported emotions with the AI-generated labels from both speech and text-based models to analyze potential alignments. This is a multi-modal approach with several methodological considerations, but also some methodological strengths.

The vocal recordings were transcribed and analyzed with the text-based emotion recognition tool NLP Cloud and self-assessments of emotions were collected after each interview, allowing a comparison between speech-based AI, text-based AI and self-assessment scores given by the participants in the interviews. The use of three different methods resulted in triangulation, which increased the flexibility and credibility in the findings. In addition to this, the usage of pre-trained AI models ensured consistent processing. However, some information loss was expected for the transcript text analyzed in NLP Cloud. When the model takes in what is said rather than how it is said, many important emotional cues such as intensity or pitch get lost. This could possibly have led to some emotions being misinterpreted or not catching the full complexity of the emotions expressed, based on only the text-based analysis.

While NLP Cloud contributed in addressing RQ2, some limitations in loss of prosodic information may have reduced the full emotional understanding. The self-reported emotions introduced a valuable reference point for this research. Some agreement was found between the AI models and self-reported emotions, but some of the self-assessed scores may also have been slightly exaggerated. The emotional memories and personal interpretations of emotions by participants can have influenced the self-assessed emotion scores. While the self-reported emotion scores have some limitations which likely contributed to some variability in the analyses, they helped valuably address RQ3.

The complexity of emotion detection across different modalities is highlighted by the AI models being able to capture some emotions in a quite robust way while struggling more with others. The few emotional categories used for this research may have limited the emotion recognition, where an implementation of more emotions and features possibly could have captured the emotions in a better way.

All methods used in answering RQ2 and RQ3 provided valuable information and findings, however the research could have benefited from an expansion of the emotion categories to help identify emotions in a more accurate and complex way.

5.5.3 Summary of Methodological Considerations

To address all research questions, this study utilized a multi-method approach combining speech-based and text-based emotion recognition with self-reported emotion scores of the participants from the interviews.

Although all methods contributed with important findings and significant insights into emotional expression in Swedish speech, a number of limitations emerged.

While the triangulation of speech, text and self-assessment scores contributed to the strength and credibility of the findings, size of dataset, model transparency and other limitations such as variabilities and inconsistency in having spontaneous interviews may have impacted the effectiveness of the findings. Although highlighting some areas for improvement for future studies, the methods chosen for this research overall contributed to answering the research questions in a comprehensive manner.

5.6 Limitations

5.6.1 Limitations RQ1

There are several limitations which should be taken into consideration when interpreting the results presented for the first research question.

Firstly, the dataset conducted and used for this study consist of spontaneous conversational semi-structured interviews. This likely has resulted in more subtle emotions than an acted dataset would contain with weaker emotional expressions reducing the detectability of different vocal markers. This study has also been restricted to a small set of acoustic features for the analysis. There are possibilities other acoustic features would have contributed with relevant aspects for the analyses if included. This study does not involve all vocal features that are included in the research by Ekberg (Ekberg et al., 2023) which potentially could have shown relevant information and possible alignments, however, due to the timeframe and scope of this study only five vocal markers were chosen.

Another important limitation to consider is the static averages of the vocal features. Despite effort where some clips have been visualized in its entirety, showing spikes in one specific emotion combined with one specific vocal feature from time to time, most clips involved in the process of answering the first research question are the result of an average per clip. This may have obscured some dynamic fluctuations of all emotions over time.

A further constraint worth noting is that the Hume AI emotion scores are not perfect estimates of true emotional states themselves and should not be considered ground truth in this study, only a way of comparison and finding potential similarities. Additionally, although Hume showed varied and appropriate emotion outputs, it is important to note that the Hume AI emotion outputs are based on soft scoring. This often produces mixed emotions rather than single emotional states, making comparisons more difficult.

Finally, the usage of the Swedish vocal data collected from semi-structured interviews. These interviews consist of spontaneous and conversational speech which likely do not involve as strong emotional expressions as acted datasets. Conversational speech tends to be more subtle and may have a reduced level of clear vocal markers.

Worth noting is that the interviews conducted for this study were context-dependent and influenced by topics selected by the participants of the interviews themselves, introducing potential additional variability from interview to interview. Contrary, the Swedish research by Ekberg consisted of 14 repeated sentences.

5.6.2 Limitations RQ2 and RQ3

This study has presented several important insights in emotion detection using AI models, although there are several limitations that should be noted. The spontaneous nature of the interviews remains a limitation through RQ2 and RQ3, as these vocal recordings may have given more subtle and muted emotional expressions compared to an acted dataset would. The calm setting of the interviews may also be an explanation to why fear and surprise especially was not detected to a high degree.

The self-reported emotions unavoidably involve subjective biases, which possibly could have resulted in some variety across interviews. The dataset size and the limited emotions remains a limitation, where a larger dataset and more emotions possibly could lead to broader findings and correlations.

Conclusion

6.1 Summary of Key Findings and Answering Research Questions

Based on the findings presented in chapter 4, the following key findings and answers to the research questions were found.

In the first research question, the aim was to compare the vocal features from speech recordings gathered from interviews to see whether or not they correlate with AI-based emotion detection from the model Hume AI. Existing research on vocal markers in Swedish speech was used for better comparison. Here, the results showed some alignments between AI-based emotion recognition models with the existing research on vocal markers in Swedish language. Many correlations were weak or moderate and the results showed some limitations, for example. the nature of spontaneous speech involved in interviews presented challenges for some of the emotions, while other emotions proved to be more relevant indicators of emotional states. The results revealed that Hume AI showed promise with the contextual awareness and the flexibility of the model. In this research question, important insight was also given through the segment level analyses which presented the fluctuations in the emotions better than average clip analyses.

For the second research question, where the purpose was to investigate if emotions from textual content of speech, transcribed and analysed with the AI model NLP Cloud could be understood through comparisons made with the speech based AI model, Hume AI. Here, some agreements were found between the AI models Hume AI and NLP Cloud, for some emotions more than others. Joy and anger showed better alignment whereas fear and surprise did not align as much. The overall cause for this most likely being the complex nature of some emotions where the cues may not be as pronounced.

For the third and last research question, the comparisons between the AI models Hume AI and NLP Cloud and self-assessed emotion scores from the interviews were investigated. These results showed there were some alignments between the models and self-assessed scores, though these alignments were stronger with NLP Cloud than with Hume AI. This gave insight into the fact that although some alignments were stronger with NLP Cloud than others, a multimodal approach that integrates several sources results in better detection of emotions, where only relying on either Hume or NLP Cloud would not give as much insight in the results.

Overall, the findings in this research provided answers to all research questions, contributing to a deeper understanding of emotion recognition in the Swedish language.

6.2 Contribution to the Field

This thesis contributes to the vocal and linguistic aspect of the growing fields affective computing and natural language processing (NLP).

This study utilized a multimodal approach, using both text and audio. The results of this suggested that using more than one modality improves the accuracy when classifying emotions.

While there is a lot of existing research focusing on vocal emotion recognition in affective computing and NLP, there was a noticeable gap for research specifically in the Swedish language. Much of the existing research also used acted datasets, which was not used in this study. Instead, this study consisted of semi-structured interviews with a set of questions for the interviewees to choose from. With the spontaneous nature of the datasets from speech recorded from interviews, this offers valuable insight into how emotions can be recognized in a setting more similar to real life. This can also contribute to the development of more emotionally aware AI models and systems through the insights of the more subtle cues for different emotions in these settings. Possible areas where this can be applied is in human-computer interactions, virtual assistance, and mental health monitoring or similar.

6.3 Limitations of the Study

There were several limitations of significance throughout this thesis. The limited size of the dataset and the chosen emotions, as the acoustic features, may have impacted the results. A bigger dataset with more emotions and acoustic features could have given important insight and clearer results if included.

Another important limitation to mention is the spontaneous nature of the speech gathered from the interviews. Using an acted dataset may have resulted in emotions being more accurately and stronger identified, whereas the interview setting may have led to some emotions being left undetected as they may have been expressed too subtly. This may be due to conversational speech being more subtle than acted speech, reducing levels of clear emotions and vocal markers. Worth noting is also that the Hume AI emotion scores used for comparisons, should not be considered ground truth, but more as a something to compare against.

This study also used self-reported emotion scores as a comparison, which may have given a large variety across the dataset due to the participants own self-perception.

6.4 Future Research

There are several ways to build on this study for future research. Key areas for further advancement are a bigger dataset and a bigger set of emotions and vocal features. Using more emotions and vocal features can contribute to a more nuanced understanding of complex emotions such as confusion or boredom, or even blended emotions. This, combined with a bigger and more diverse dataset from real word scenarios, could give robust data and insight in how emotions are recognized in linguistic data, important for future contributions for this field in the future such as in mental health monitoring etc.

A possible valuable direction for future research is to apply and try out emotion-recognition systems in real time settings and environments to provide scientifically justified insights, though this should be done with considerations for ethical implications and responsibility, specifically in sensitive areas involving mental health or similar.

6.5 Final Conclusion

This thesis investigated AI based emotion recognition in textual and vocal content in the Swedish language. The results showed some similarities with existing research on vocal markers in Swedish language and strongly demonstrated the importance of using more than one modality to improve the accuracy of the emotion recognition, especially for more complex emotions.

All of the findings gathered from the research in this thesis help contribute to the growing fields of affective computing and natural language processing. With technology becoming an increasingly larger part of society, understanding human emotions is a step towards AI becoming more effective and empathetic, which can help mold the future of education, health care and human-computer interactions.

Bibliography

- Abbaschian, B. J., Sierra-Sosa, D., & Elmaghriby, A. (2021). Deep learning techniques for speech emotion recognition, from databases to models. *Sensors Basel, Switzerland*, 21, 1–27. <https://doi.org/10.3390/s21041249>
- Adebisi, M. O., Adeliyi, T. T., Olaniyan, D., & Olaniyan, J. (2024). Advancements in accurate speech emotion recognition through the integration of cnn-am model. *Telkomnika*, 22, 606–618. <https://doi.org/10.12928/TELKOMNIKA.v22i3.25708>
- Ahammed, M., Sheikh, R., Hossain, F., Liza, S. M., Rahman, M. A., Mahmud, M., Brown, D. J., Ahmed, M. R., Ben-Abdallah, H., Kaiser, M. S., & Zhong, N. (2024). Speech emotion recognition: An empirical analysis of machine learning algorithms across diverse data sets. In *Applied intelligence and informatics* (pp. 32–46). Springer. https://doi.org/10.1007/978-3-031-68639-9_3
- AI, H. (n.d.-a). Prosody. <https://www.hume.ai/products/speech-prosody-model>
- AI, H. (n.d.-b). Vocal expression. <https://www.hume.ai/products/vocal-expression-model>
- Alroobaea, R. (2024). Cross-corpus speech emotion recognition with transformers: Leveraging handcrafted features and data augmentation. *Computers in biology and medicine*, 179, 108841. <https://doi.org/10.1016/j.combiomed.2024.108841>
- Areshey, A., & Mathkour, H. (2024). Exploring transformer models for sentiment classification: A comparison of bert, roberta, albert, distilbert, and xlnet. *Expert systems*, 41. <https://doi.org/10.1111/exsy.13701>
- Auphonic. (n.d.). Features. <https://auphonic.com/features>
- Babu, P. A., Nagaraju, V. S., & Vallabhuni, R. R. (2021). Speech emotion recognition system with librosa. *10th IEEE International Conference on Communication Systems and Network Technologies CSNT*, 421–424. <https://doi.org/10.1109/CSNT51715.2021.9509714>
- Baird, A., Tzirakis, P., Brooks, J. A., Gregory, C. B., Schuller, B., Batliner, A., Keltner, D., & Cowen, A. (2022). The acii 2022 affective vocal bursts workshop & competition. *10th International Conference on Affective Computing and Intelligent Interaction Workshops and Demos ACIIW*, 1–5. <https://doi.org/10.1109/ACIIW57231.2022.10086002>
- Barbon, R. S., & Akabane, A. T. (2022). Towards transfer learning techniques—bert, distilbert, bertimbau, and distilbertimbau for automatic text classification from different languages: A case study. *Sensors (Basel, Switzerland)*, 22, 8184. <https://doi.org/10.3390/s22218184>
- Brooks, J. A., Tzirakis, P., Baird, A., Kim, L., Opara, M., Fang, X., Keltner, D., Monroy, M., Corona, R., Metrick, J., & Cowen, A. S. (2023). Deep learning reveals what vocal bursts express in different cultures. *Nature human behaviour*, 7, 240–250. <https://doi.org/10.1038/s41562-022-01489-2>
- Bruce, P., & Bruce, A. (2017). *Practical statistics for data scientists*. O'Reilly.
- Bryman, A., Bell, E., Reck, J., & Fields, J. (2022). *Social research methods*. Oxford University Press.
- Cai, Y., Li, X., & Li, J. (2023). Emotion recognition using different sensors, emotion models, methods and datasets: A comprehensive review. *sensors*. <https://doi.org/10.3390/s23052455>
- Cloud, N. (n.d.). Advanced ai platform. <https://nlpcloud.com/>

- Cohen, J. (1977). *Statistical power analysis for the behavioral sciences (revised edition)*. Academic Press.
- Cowen, A. S., Laukka, P., Elfenbein, H. A., Liu, R., & Keltner, D. (2019). The primacy of categories in the recognition of 12 emotions in speech prosody across two cultures. *Nature human behaviour*, 3, 369–382. <https://doi.org/10.1038/s41562-019-0533-6>
- Creswell, J. W., & Creswell, J. D. (2023). *Research design : Qualitative, quantitative, and mixed methods approaches* (Fifth). SAGE.
- Demszky, D., D, M.-A., Ko, J., Cowen, A., Nemade, G., & Ravi, S. (2020). Goemotions: A dataset of fine-grained emotions. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 4040–4054. <https://doi.org/10.48550/arxiv.2005.00547>
- DeSouza, D. D., Robin, J., Gumus, M., & Yeung, A. (2021). Natural language processing as an emerging tool to detect late-life depression. *Frontiers in psychiatry*, 12, 719125. <https://doi.org/10.3389/fpsyg.2021.719125>
- Drougkas, G., Bakker, E. M., & Spruit, M. (2024). Multimodal machine learning for language and speech markers identification in mental health. *BMC medical informatics and decision making*, 24, 320–354. <https://doi.org/10.1186/s12911-024-02772-0>
- Ekberg, M., Stavrinos, G., Andin, J., Stenfelt, S., & Dahlström, Ö. (2023). Acoustic features distinguishing emotions in swedish speech. *Journal of voice*. <https://doi.org/10.1016/j.jvoice.2023.03.010>.
- Ekman, P. (2016). What scientists who study emotion agree about. *Perspectives on psychological science*, 11, 31–34. <https://doi.org/10.1177/1745691615596992>
- Ekman, P., & Cordaro, D. (2011). What is meant by calling emotions basic. *Emotion Review*, 3, 364–370. <https://doi.org/10.1177/1754073911410740>
- Ermakova, T., Fabian, B., Golimblevskaia, E., & Henke, M. (2023). A comparison of commercial sentiment analysis services. *SN computer science*, 4, 477–. <https://doi.org/10.1007/s42979-023-01886-y>
- Esfahani, S. H. N., & Adda, M. (2024). Classical machine learning and large models for text-based emotion recognition. *Procedia Computer Science*, 241, 77–84. <https://doi.org/10.1016/j.procs.2024.08.013>
- Fröhholz, S., & Belin, P. (2019). *The oxford handbook of voice perception*. Oxford University Press.
- HappyPlanetIndex. (n.d.). What is the happy planet index? <https://happyplanetindex.org/learn-about-the-happy-planet-index/>
- Hume, A. (n.d.-a). About hume. <https://www.hume.ai/about>
- Hume, A. (n.d.-b). About the science. <https://dev.hume.ai/docs/resources/science>
- Jadoul, Y., de Boer, B., & Ravignani, A. (2024). Parselmouth for bioacoustics: Automated acoustic analysis in python. *Bioacoustics Berkhamsted*, 33, 1–19. <https://doi.org/10.1080/09524622.2023.2259327>
- Jadoul, Y., Thompson, B., & de Boer, B. (2018). Introducing parselmouth: A python interface to praat. *Journal of phonetics*, 71, 1–15. <https://doi.org/10.1016/j.wocn.2018.07.001>
- Juslin, P. N., Laukka, P., & Bänziger, T. (2018). The mirror to our soul? comparisons of spontaneous and posed vocal expression of emotion. *Journal of nonverbal behavior*, 42, 1–40. <https://doi.org/10.1007/s10919-017-0268-x>
- Kansara, D., Sawant, V., Shekokar, N., Vasudevan, H., Narvekar, M., & Michalas, A. (2020). Comparison of traditional machine learning and deep learning approaches for sentiment analysis. In *Advanced computing technologies and applications* (pp. 365–377). Springer. https://doi.org/10.1007/978-981-15-3242-9_35

- Khalil, R. A., Jones, E., Babar, M. I., Jan, T., Zafar, M. H., & Alhussain, T. (2019). Speech emotion recognition using deep learning techniques: A review. *IEEE access*, 7, 117327–117345. <https://doi.org/10.1109/ACCESS.2019.2936124>
- Kumar, S., & Singh, S. (2024). Fine-tuning llama 3 for sentiment analysis: Leveraging aws cloud for enhanced performance. *SN computer science*, 5, 1161. <https://doi.org/10.1007/s42979-024-03473-1>
- Kusal, S., Patil, S., Choudrie, J., Kotecha, K., Vora, D., & Pappas, I. (2023). A systematic review of applications of natural language processing and future challenges with special emphasis in text-based emotion detection. *The Artificial intelligence review*, 56, 15129–15215. <https://doi.org/10.1007/s10462-023-10509-0>
- Kusal, S. D., Patil, S. G., Choudrie, J., & Kotecha, K. V. (2024). Understanding the performance of ai algorithms in text-based emotion detection for conversational agents. *ACM transactions on Asian and low-resource language information processing*, 23, 1–26. <https://doi.org/10.1145/3643133>
- Lee, C. M., & Narayanan, S. S. (2005). Toward detecting emotions in spoken dialogs. *IEEE transactions on speech and audio processing*, 13, 293–303. <https://doi.org/10.1109/TSA.2004.838534>
- Lee, S. J., Lim, J., Paas, L., & Ahn, H. S. (2023). Transformer transfer learning emotion detection model: Synchronizing socially agreed and self-reported emotions in big data. *Neural computing & applications*, 35, 10945–10956. <https://doi.org/10.1007/s00521-023-08276-8>
- Lian, H., Lu, C., Li, S., Zhao, Y., Tang, C., & Zong, Y. (2023). A survey of deep learning-based multimodal emotion recognition: Speech, text, and face. *Entropy (Basel, Switzerland)*, 25, 1440–. <https://doi.org/10.3390/e25101440>
- Madhuri, S., & Lakshmi, V. (2021). Detecting emotion from natural language text using hybrid and nlp pre-trained models. *Turkish journal of computer and mathematics education*, 12, 4095–4103.
- Maruf, A. A., Khanam, F., Haque, M. M., Jiyad, Z. M., Mridha, M. F., & Aung, Z. (2024). Challenges and opportunities of text-based emotion detection: A survey. *IEEE access*, 12, 18416–18450. <https://doi.org/10.1109/ACCESS.2024.3356357>
- Montasem, A., Brown, S. L., & Harris, R. (2013). Do core self-evaluations and trait emotional intelligence predict subjective well-being in dental students? *Journal of Applied Social Psychology*, 43, 1097–1103. <https://doi.org/10.1111/jasp.12074>
- Núñez, A. Á., del C Santiago Díaz, M., Vázquez, A. C. Z., Marcial, J. P., & Linares, G. T. R. (2024). Emotion detection using natural language processing. *International Journal of Combinatorial Optimization Problems and Informatics*, 15, 108–114. <https://doi.org/10.61467/2007.1558.2024.v15i5.564>
- Oatley, K., Keltner, D., & Jenkins, J. M. (2019). *Understanding emotions fourth edition*. Blackwell.
- OpenAI. (2022, September). Introducing whisper. <https://openai.com/index/whisper/>
- Praseetha, V. M., & Joby, P. P. (2022). Speech emotion recognition using data augmentation. *International journal of speech technology*, 25, 783–792. <https://doi.org/10.1007/s10772-021-09883-3>
- Radford, A., Kim, J. W., Xu, T., Brockman, G., McLeavey, C., & Sutskever, I. (2022). Robust speech recognition via large-scale weak supervision. *Robust Speech Recognition via Large-Scale Weak Supervision*. <https://doi.org/10.48550/arxiv.2212.04356>
- Rahman, M. M., Hossain, M. A., Hasan, T., Ahmed, M. K., Sultana, R., & Islam, M. S. (2024). Emotionnet: Pioneering deep learning fusion for real-time speech emotion recognition with convolutional neural networks. *2024 6th International Conference on Electrical*

- Engineering and Information & Communication Technology ICEEICT*, 592–597. <https://doi.org/https://doi.org/10.1109/ICEEICT62016.2024.10534404>
- Rathi, T., & Tripathy, M. (2024). Analyzing the influence of different speech data corpora and speech features on speech emotion recognition: A review. *Speech communication*, 162, 103102–. <https://doi.org/https://doi.org/10.1016/j.specom.2024.103102>
- Repede, S. E., & Brad, R. (2024). Llama 3 vs. state-of-the-art large language models: Performance in detecting nuanced fake news. *Computers (Basel)*, 13, 292. <https://doi.org/https://doi.org/10.3390/computers13110292>
- Ri, F. A. D., Ciardi, F. C., & Conci, N. (2023). Speech emotion recognition and deep learning: An extensive validation using convolutional neural networks. *IEEE Access*, 11, 1. <https://doi.org/https://doi.org/10.1109/ACCESS.2023.3326071>
- Safari, F., & Chalechale, A. (2023). Emotion and personality analysis and detection using natural language processing, advances, challenges and future scope. *The Artificial intelligence review*, 56, 3273–3297. <https://doi.org/https://doi.org/10.1007/s10462-023-10603-3>
- Sahoo, C., Wankhade, M., & Singh, B. K. (2023). Sentiment analysis using deep learning techniques: A comprehensive review. *International journal of multimedia information retrieval*, 12, 41–. <https://doi.org/https://doi.org/10.1007/s13735-023-00308-2>
- Scherer, K. R., Frühholz, S., & Belin, P. (2018). *Acoustic patterning of emotion vocalizations*. Oxford University Press. <https://doi.org/10.1093/oxfordhb/9780198743187.013.4>
- Shelke, N., Chaudhury, S., Chakrabarti, S., Bangare, S. L., Yogapriya, G., & Pandey, P. (2022). An efficient way of text-based emotion analysis from social media using lra-dnn. *Neuroscience informatics*, 2, 100048. <https://doi.org/https://doi.org/10.1016/j.neuri.2022.100048>
- Siedlecka, E., & Denson, T. F. (2019). Experimental methods for inducing basic emotions: A qualitative review. *Emotion Review*, 11, 87–97. <https://doi.org/https://doi.org/10.1177/1754073917749016>
- Simcock, G., McLoughlin, L. T., Regt, T. D., Broadhouse, K. M., Beaudequin, D., Lagopoulos, J., & Hermens, D. F. (2020). Associations between facial emotion recognition and mental health in early adolescence. *International Journal of Environmental Research and Public Health*, 17, 330. <https://doi.org/10.3390/ijerph17010330>
- Singh, S. (2023). Emotion recognition for mental health prediction using ai techniques: An overview. *International Journal of Advanced Research in Computer Science*, 14, 87–107. <https://doi.org/10.26483/ijarcs.v14i3.6975>
- Sönmez, Y. Ü., & Varol, A. (2024). In-depth investigation of speech emotion recognition studies from past to present. the importance of emotion recognition from speech signal for ai. *Intelligent systems with applications*, 200351–. <https://doi.org/https://doi.org/10.1016/j.iswa.2024.200351>
- Thompson, W. F., Schellenberg, E. G., & Husain, G. (2004). Decoding speech prosody: Do music lessons help? *Emotion*, 4, 46–64. <https://doi.org/https://doi.org/10.1037/1528-3542.4.1.46>
- Tian, L., Oviatt, S., Muszyński, M., Chamberlain, B. C., Healey, J., & Sano, A. (2022). *Applied affective computing*. Association for Computing Machinery.
- Tomasello, R., Grisoni, L., Boux, I., Sammler, D., & Pulvermüller, F. (2022). Instantaneous neural processing of communicative functions conveyed by speech prosody. *Cerebral cortex*, 32, 4885–4901. <https://doi.org/https://doi.org/10.1093/cercor/bhab522>
- TwinWord. (n.d.). Just the best keywords. <https://www.twinword.com/ideas/>
- Tyagi, S., & Szénási, S. (2024). Semantic speech analysis using machine learning and deep learning techniques: A comprehensive review. *Multimedia tools and applications*, 83, 73427–73456. <https://doi.org/https://doi.org/10.1007/s11042-023-17769-6>

- Zhang, J., & M, Z. (2024). Is llama 3 good at identifying emotion? a comprehensive study. *Proceedings of the 2024 7th International Conference on Machine Learning and Machine Intelligence MLMI*, 128–132. <https://doi.org/10.1145/3696271.3696292>
- Zhang, L., Wang, S., & Liu, B. (2018). Deep learning for sentiment analysis: A survey. *Data mining and knowledge discovery*, 8, 1253–n/a. <https://doi.org/10.1002/widm.1253>
- Zhang, S., Tao, X., Chuang, Y., & Zhao, X. (2021). Learning deep multimodal affective features for spontaneous speech emotion recognition. *Speech communication*, 127, 73–81. <https://doi.org/10.1016/j.specom.2020.12.009>